

# Computer Architecture & Organization

## CS 520

Timothy N. Miller

(Course materials from Kanad Ghose)

# About the Instructor

- Timothy Normand Miller
- 16 years of industry experience
  - Chip design, embedded systems, software
- PhD from Ohio State, 2012
  - Specialization: Computer Architecture

# Prerequisites

- Coding skills
- Undergraduate logic design
- Assembly language programming
- Virtual memory and memory protection
- Undergraduate operating systems

# Prerequisites

- Computer organization
- Instruction set design
  - Opcode encoding, addressing modes
  - RISC vs. CISC, load-store
  - Data representation
  - Computer arithmetic

# Course Goals

- Thorough understanding of modern microprocessor design, primarily uniprocessor
  - Performance, design tradeoffs, implementation
  - Interaction/interface with compilers
  - Influence on OS design

# Course Materials

- Syllabus
- Homework assignments, project
- No text book
  - Online notes, papers, etc.
- Optional reading material:
  - *Computer Architecture, a Quantitative Approach*, Hennessy and Patterson
- Blackboard

# Document password

C0rEKO0L16!

# Homework

- 4 or so major assignments
- Design and coding assignments
- Some exam and quiz questions based on homework problems

# Exams & Quizzes

- Pop quizzes to review specific topics
- Midterm
- Final

# Class Project

- Coding assignment
- Two-phase
- Due around midterm and after final
- Typically involving some CPU simulator
  - Implementing a CPU yourself or
  - Making algorithm changes to msim

# Class Participation

- Subjective
- Make sure I know you!
- Ask questions in class
- Show up for class
- Pop Quizzes
- Visiting during office hours

# Grading

- Consider the grader's ability to understand your work
- Instructor will not override grader
- No give-away grade changes
- The grader's decision is final

# Cheating

- Know the Academic Honesty Policy
- **Cheating will be reported to the school.**
- “Germany's defence minister has given up his doctoral title for good, after allegations that he had plagiarised sections of his thesis.” -- BBC, 21 February 2011
  - <http://www.bbc.co.uk/news/world-europe-12532877>
- **DO YOUR OWN WORK**

# What is cheating?

- Plagiarism: Representing someone else's work as your own.
- Improper sharing: Allowing another to represent your work as theirs.

# Examples of cheating

- Copying answers from a classmate.
- Copying answers from the Internet or another course.
- Copying an algorithm.
- Sharing your answers with a classmate.
- Any case where you did not do your own work.
- All cheating is of the same magnitude.

# Not cheating

- Discussing homework problems in class, with instructor, or with TA.
- Discussing homework *questions* with classmates (i.e. what do they mean?) — but doing the work on your own.

# Why is cheating wrong?

- Grader cannot assess your performance.
- Penalizes those who do not cheat.
- Your unrecognized incompetence could cost millions of dollars or even kill people.
- HOWEVER:
- We are not looking for an excuse to fail you.

# Contact Time

- Class:
  - TR 10:05AM to 11:30AM
  - Science Library 210
- Office hours:
  - TR 1PM to 2PM and By appointment
  - Engineering (EB) Q01
- Email: [millerti@binghamton.edu](mailto:millerti@binghamton.edu)
  - Short questions, arrange appointment

# **Part I: Basic Concepts**

# Outline

- Costs, economics, markets
- Architecture basics
  - Machine cycle
  - Instruction Set Architecture (ISA)
- Processor design & implementation
- Fabrication
- Circuits, technology trends, power
- Performance
  - RISC vs. CISC
  - Locality of reference
  - Pipelining
  - The “APEX” archetypical CPU
  - Advanced Pipelining

# **Costs, Economics, Markets**

# Economics of Processor Design

- Cost of CPU =
  - High-level design (this class)
  - ***Testing***
  - Gate/transistor level implementation
  - ***Testing***
  - Masks, fabrication, packaging
  - ***Testing***
  - Business overhead & profit margin

# Cost to design new high-end CPU

- ~ 600 man years (150 designers for 4 years)
- At \$300K per person per year:
  - \$180 million

# CPU cost drivers beyond basic design

- Functional testing
- VLSI/Chip-level testing
- These increase with chip complexity

# Economy of scale

- 10K units                    \$18K / unit
- 1M units                    \$180 / unit
- Requires high volumes

# Market segments

- PCs are largest market in \$ sales
- Embedded are largest in units sold
- Workstations, high-end
- Servers

# Server processor market, 2006-2011

- 22 million server processor chips shipped
- 60% were x86-based
- \$500 to \$1000 retail per microprocessor

# Desktop & Workstation markets

- Commodity
- Highly competitive
  - Features: multi-core, memory latency, graphics, encryption, performance per watt
  - Price: \$200 in 2000, \$150 in 2005, \$105 in 2010
  - High-end chips, well over \$800 / chip

# Cell Phones & Tablets

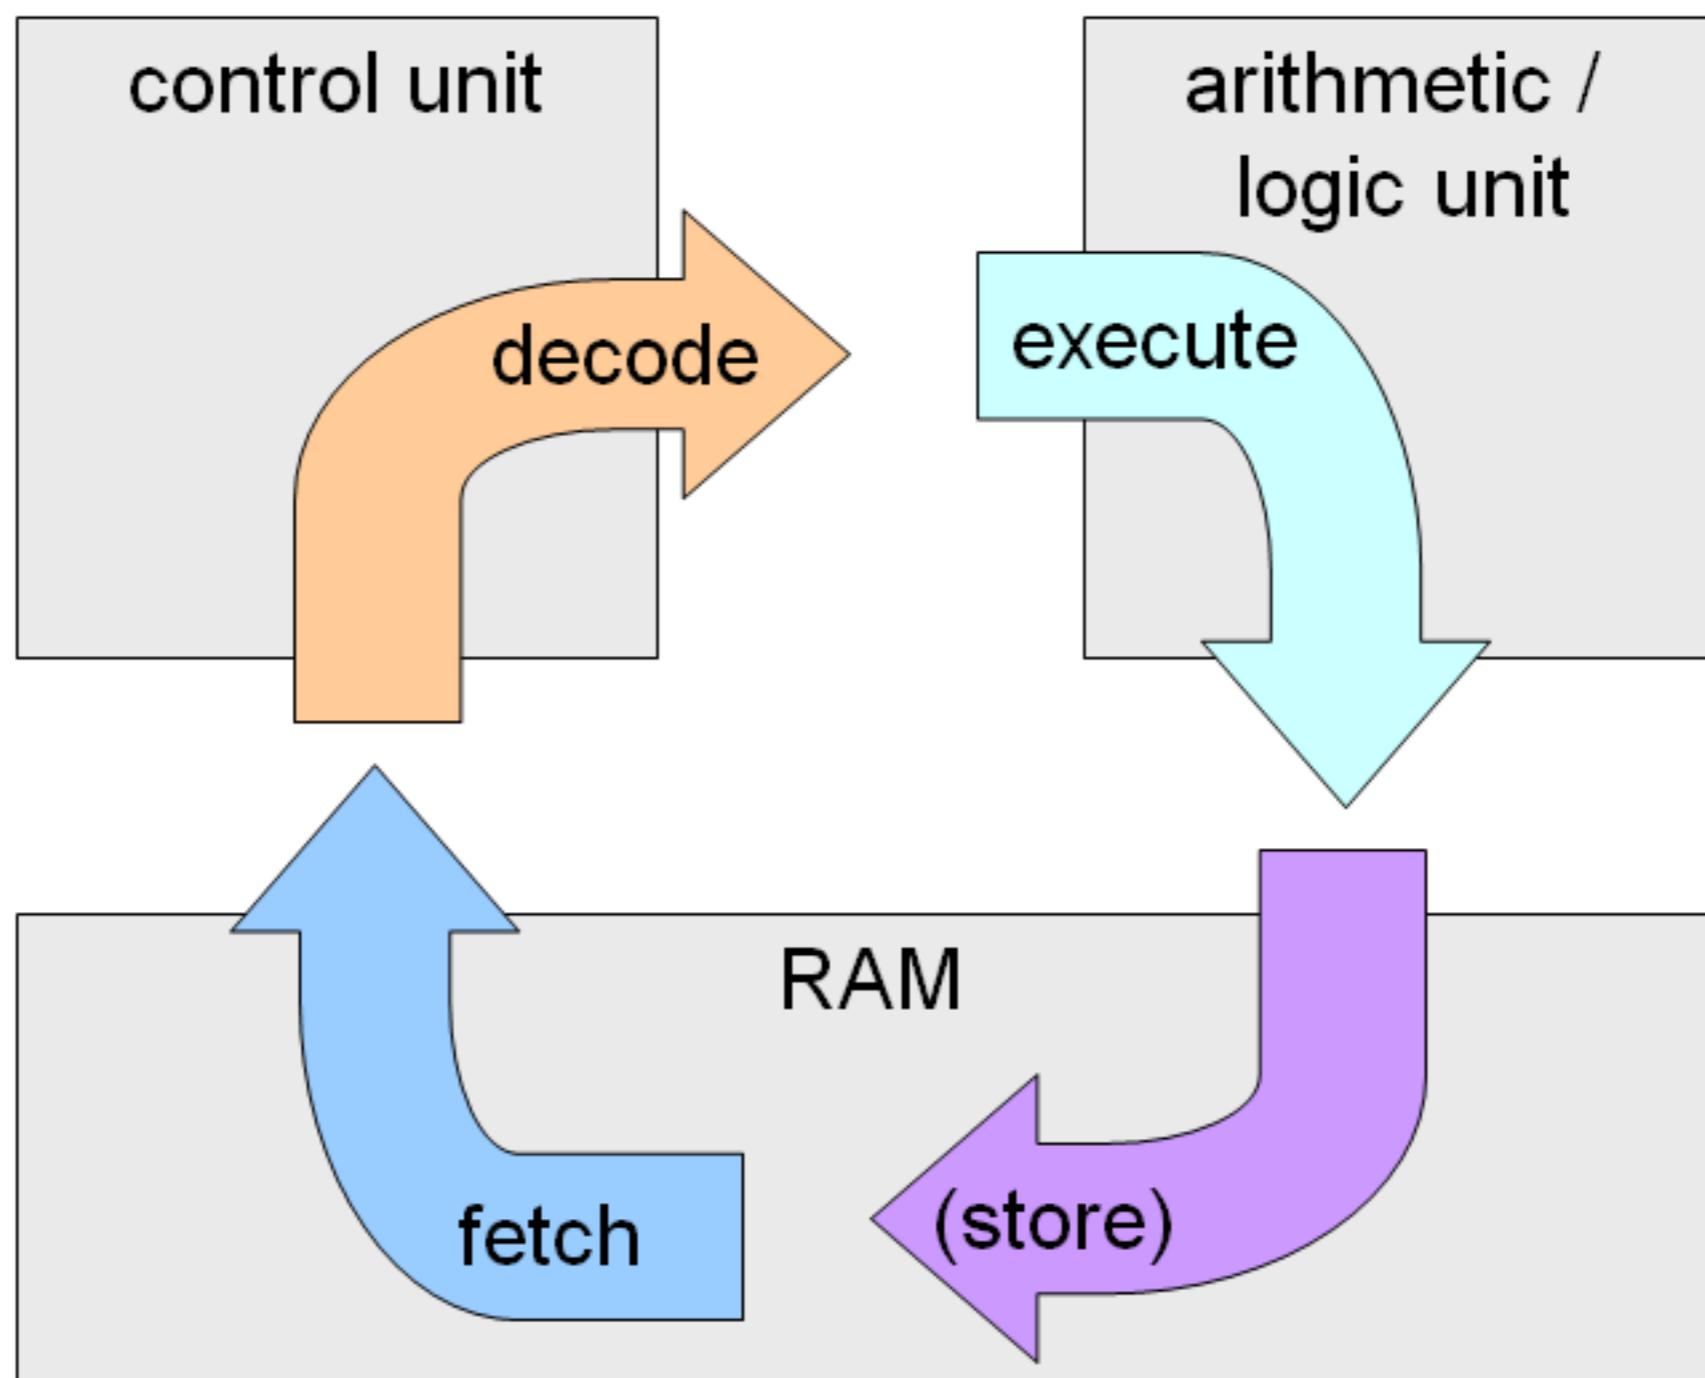
- Very competitive with many vendors and many sophisticated designs
- Most important: Performance, features, battery life
- Features: Graphics, video, encryption, speech recognition
- Evolution: Discrete chips toward highly-integrated multi-core at 1+ GHz clock
- Cost: Order of \$1 to \$10 per SoC

# Intel x86 server/ desktop market

- Revenue: \$26.5 B in 2001 to \$43.6 B in 2010
- Breakdown for 2010: 60% PC processors, 17% servers, 18% chipsets
- Intel owns 80% of PC CPU market, AMD owns most of the rest

# Architecture Basics

# Basic CPU Machine Cycle



# Instruction Set Architecture (ISA)

- Functional behavior of processor
- Characterizes processor family
- Can have multiple implementations
- Produced by compiler
- Implemented by ***microarchitecture***

# Characteristics of an ISA

- Specification provides:
  - Instruction types and semantics
  - Instruction formats
  - Addressing modes
  - Data formats (word sizes, endianness)

# Examples of ISAs

- x86                    Intel, AMD, Via, Transmeta
- x64                    Intel, AMD
- IA-64                 Intel, HP
- SPARC                 Sun, Fujitsu
- PowerPC              IBM, Motorola
- MIPS                    SGI, NEC

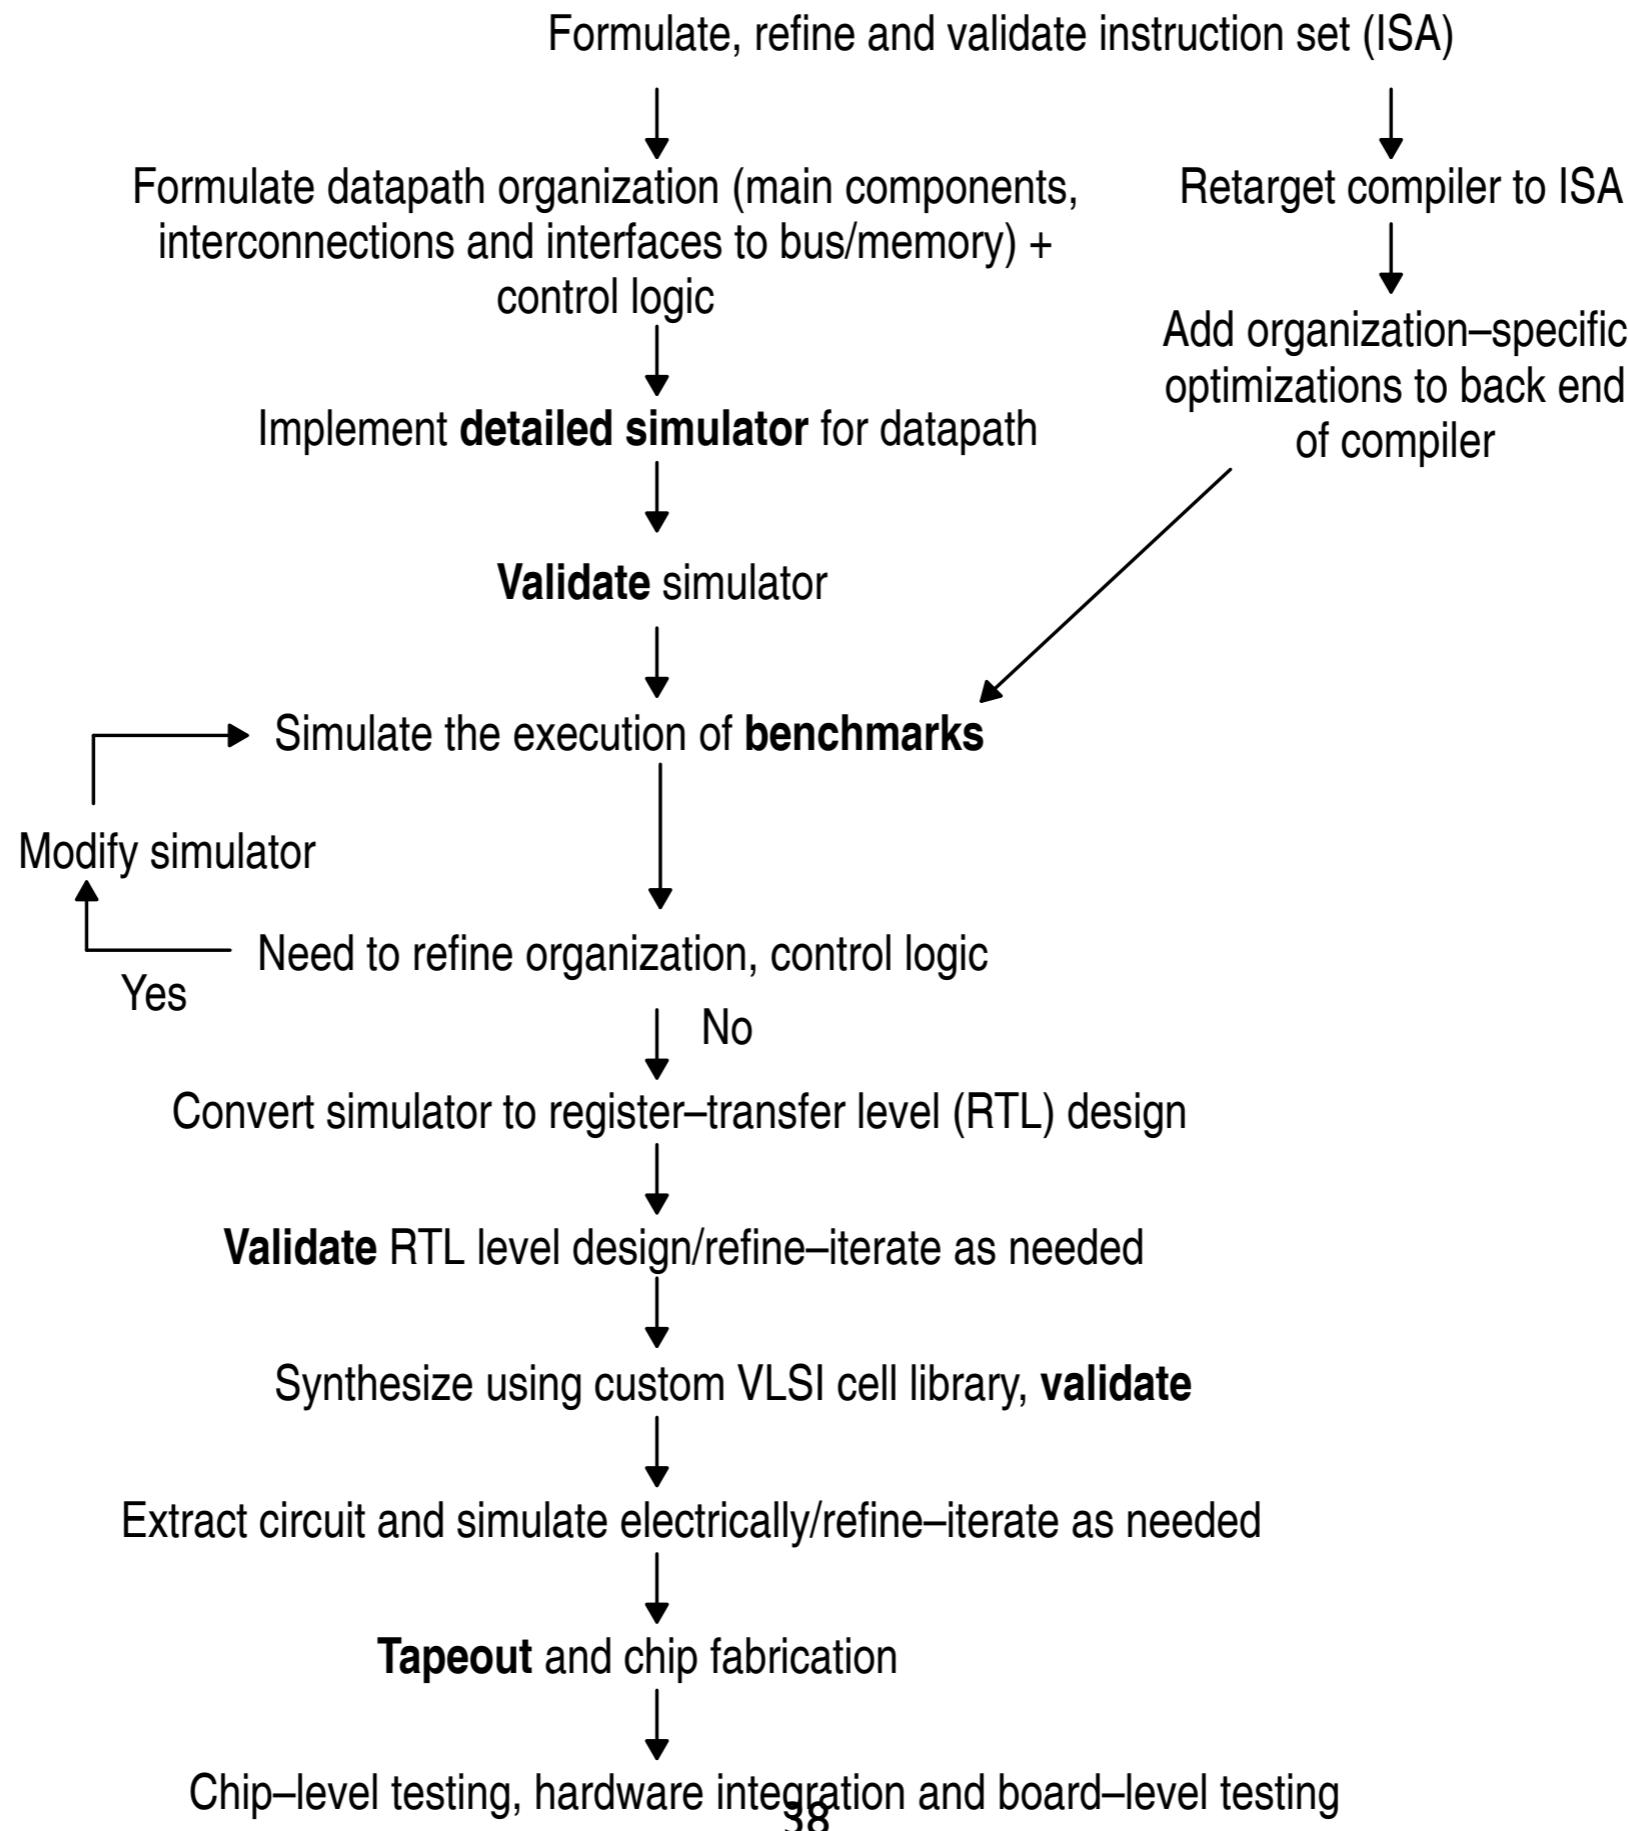
# New ISA is very expensive

- ISA design
- New hardware
- New compilers
- New software
- Must have adequate base of applications
- Would take 5+ years to recompile x86 software for PowerPC

# Designing and Implementing a Processor

## Steps Involved in Designing and Implementing a Modern Processor

---



# Typical Design & Implementation

- ISA already exists
- Design/validation team of 150 to over 200
- 2 validators for each designer
- Separate teams for interfaces, buses, chipsets
- Validate based on detailed simulation
- Discover bugs at every validation point

# Typical Design & Implementation

- Bugs discovered after release, often not serious
- Trend towards using HDL like Verilog for simulator
- Simpler CPUs go through same steps

# Fabrication

# Contemporary VLSI Technology

- Virtually all CPUs implemented completely within a single chip
- Circuit components: Transistors, wires, limited capacitors, limited resistors

# Microchip fabrication process

- Cut extremely pure silicon wafer
- Lithographically fabricate many identical copies of same design on wafer
- Series of steps that form layers of various materials:
  - Etching & deposition

**Prepare Wafer**

**Coat with Photoresist**

**Prebake**

**Align and Expose**

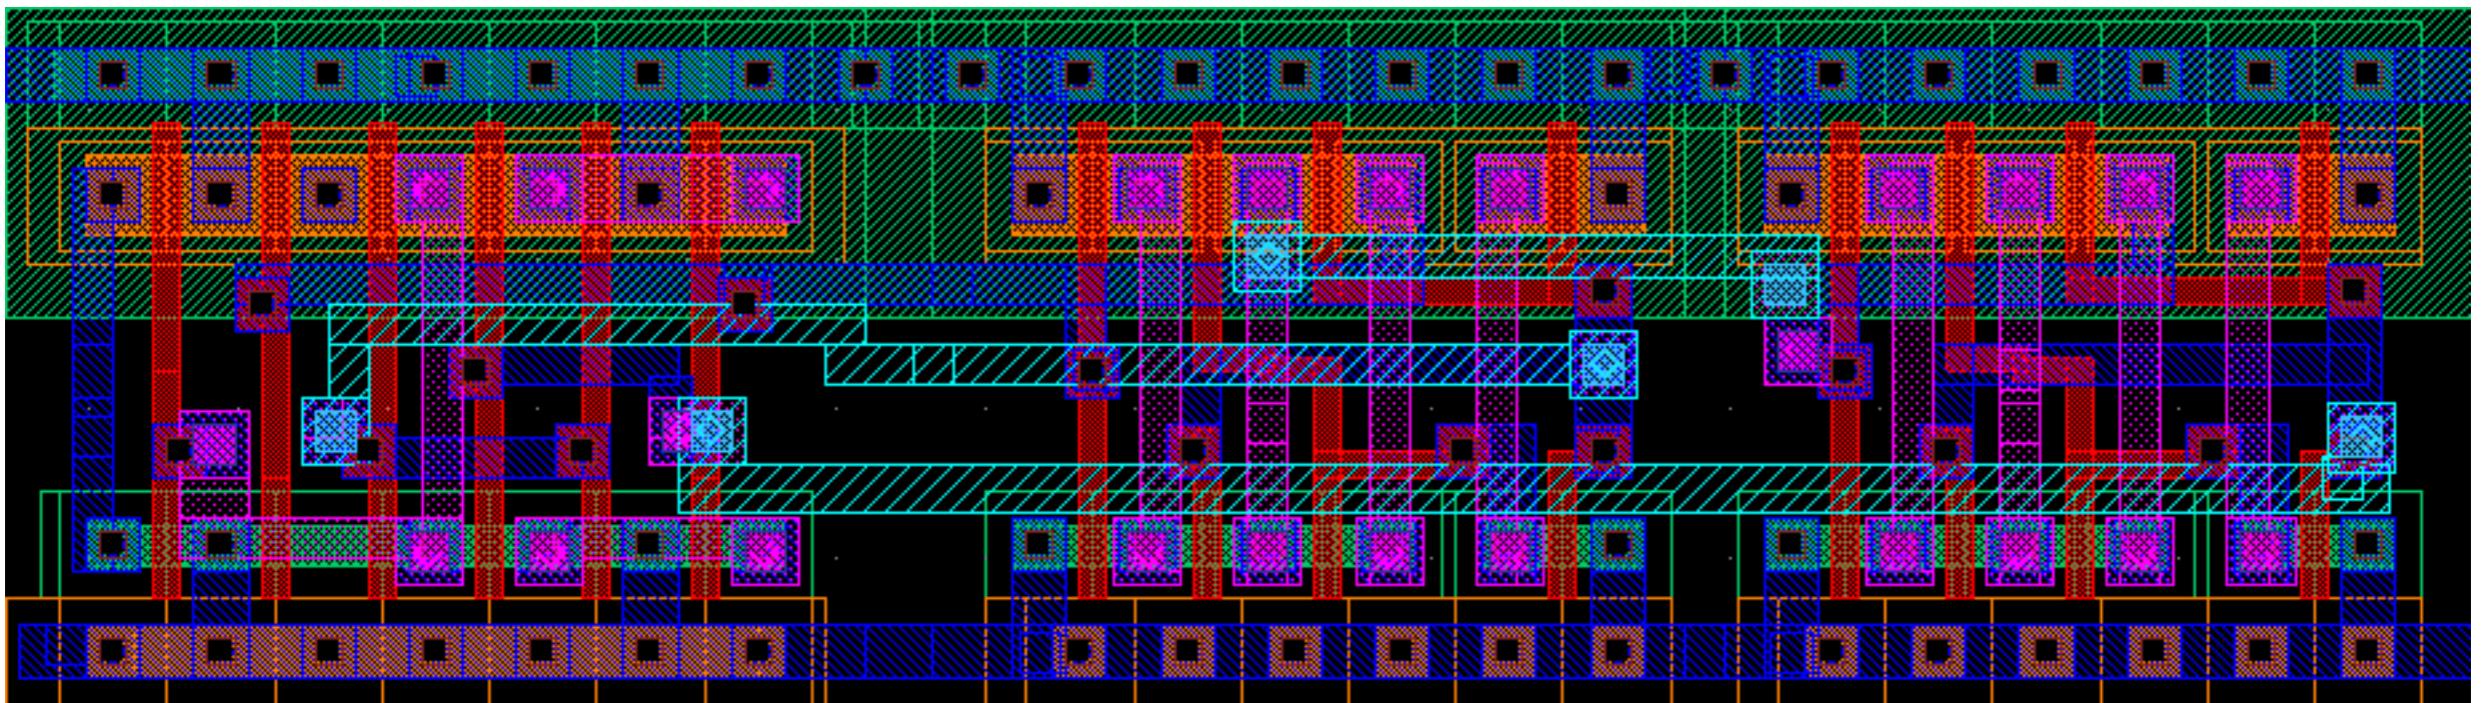
**Develop**

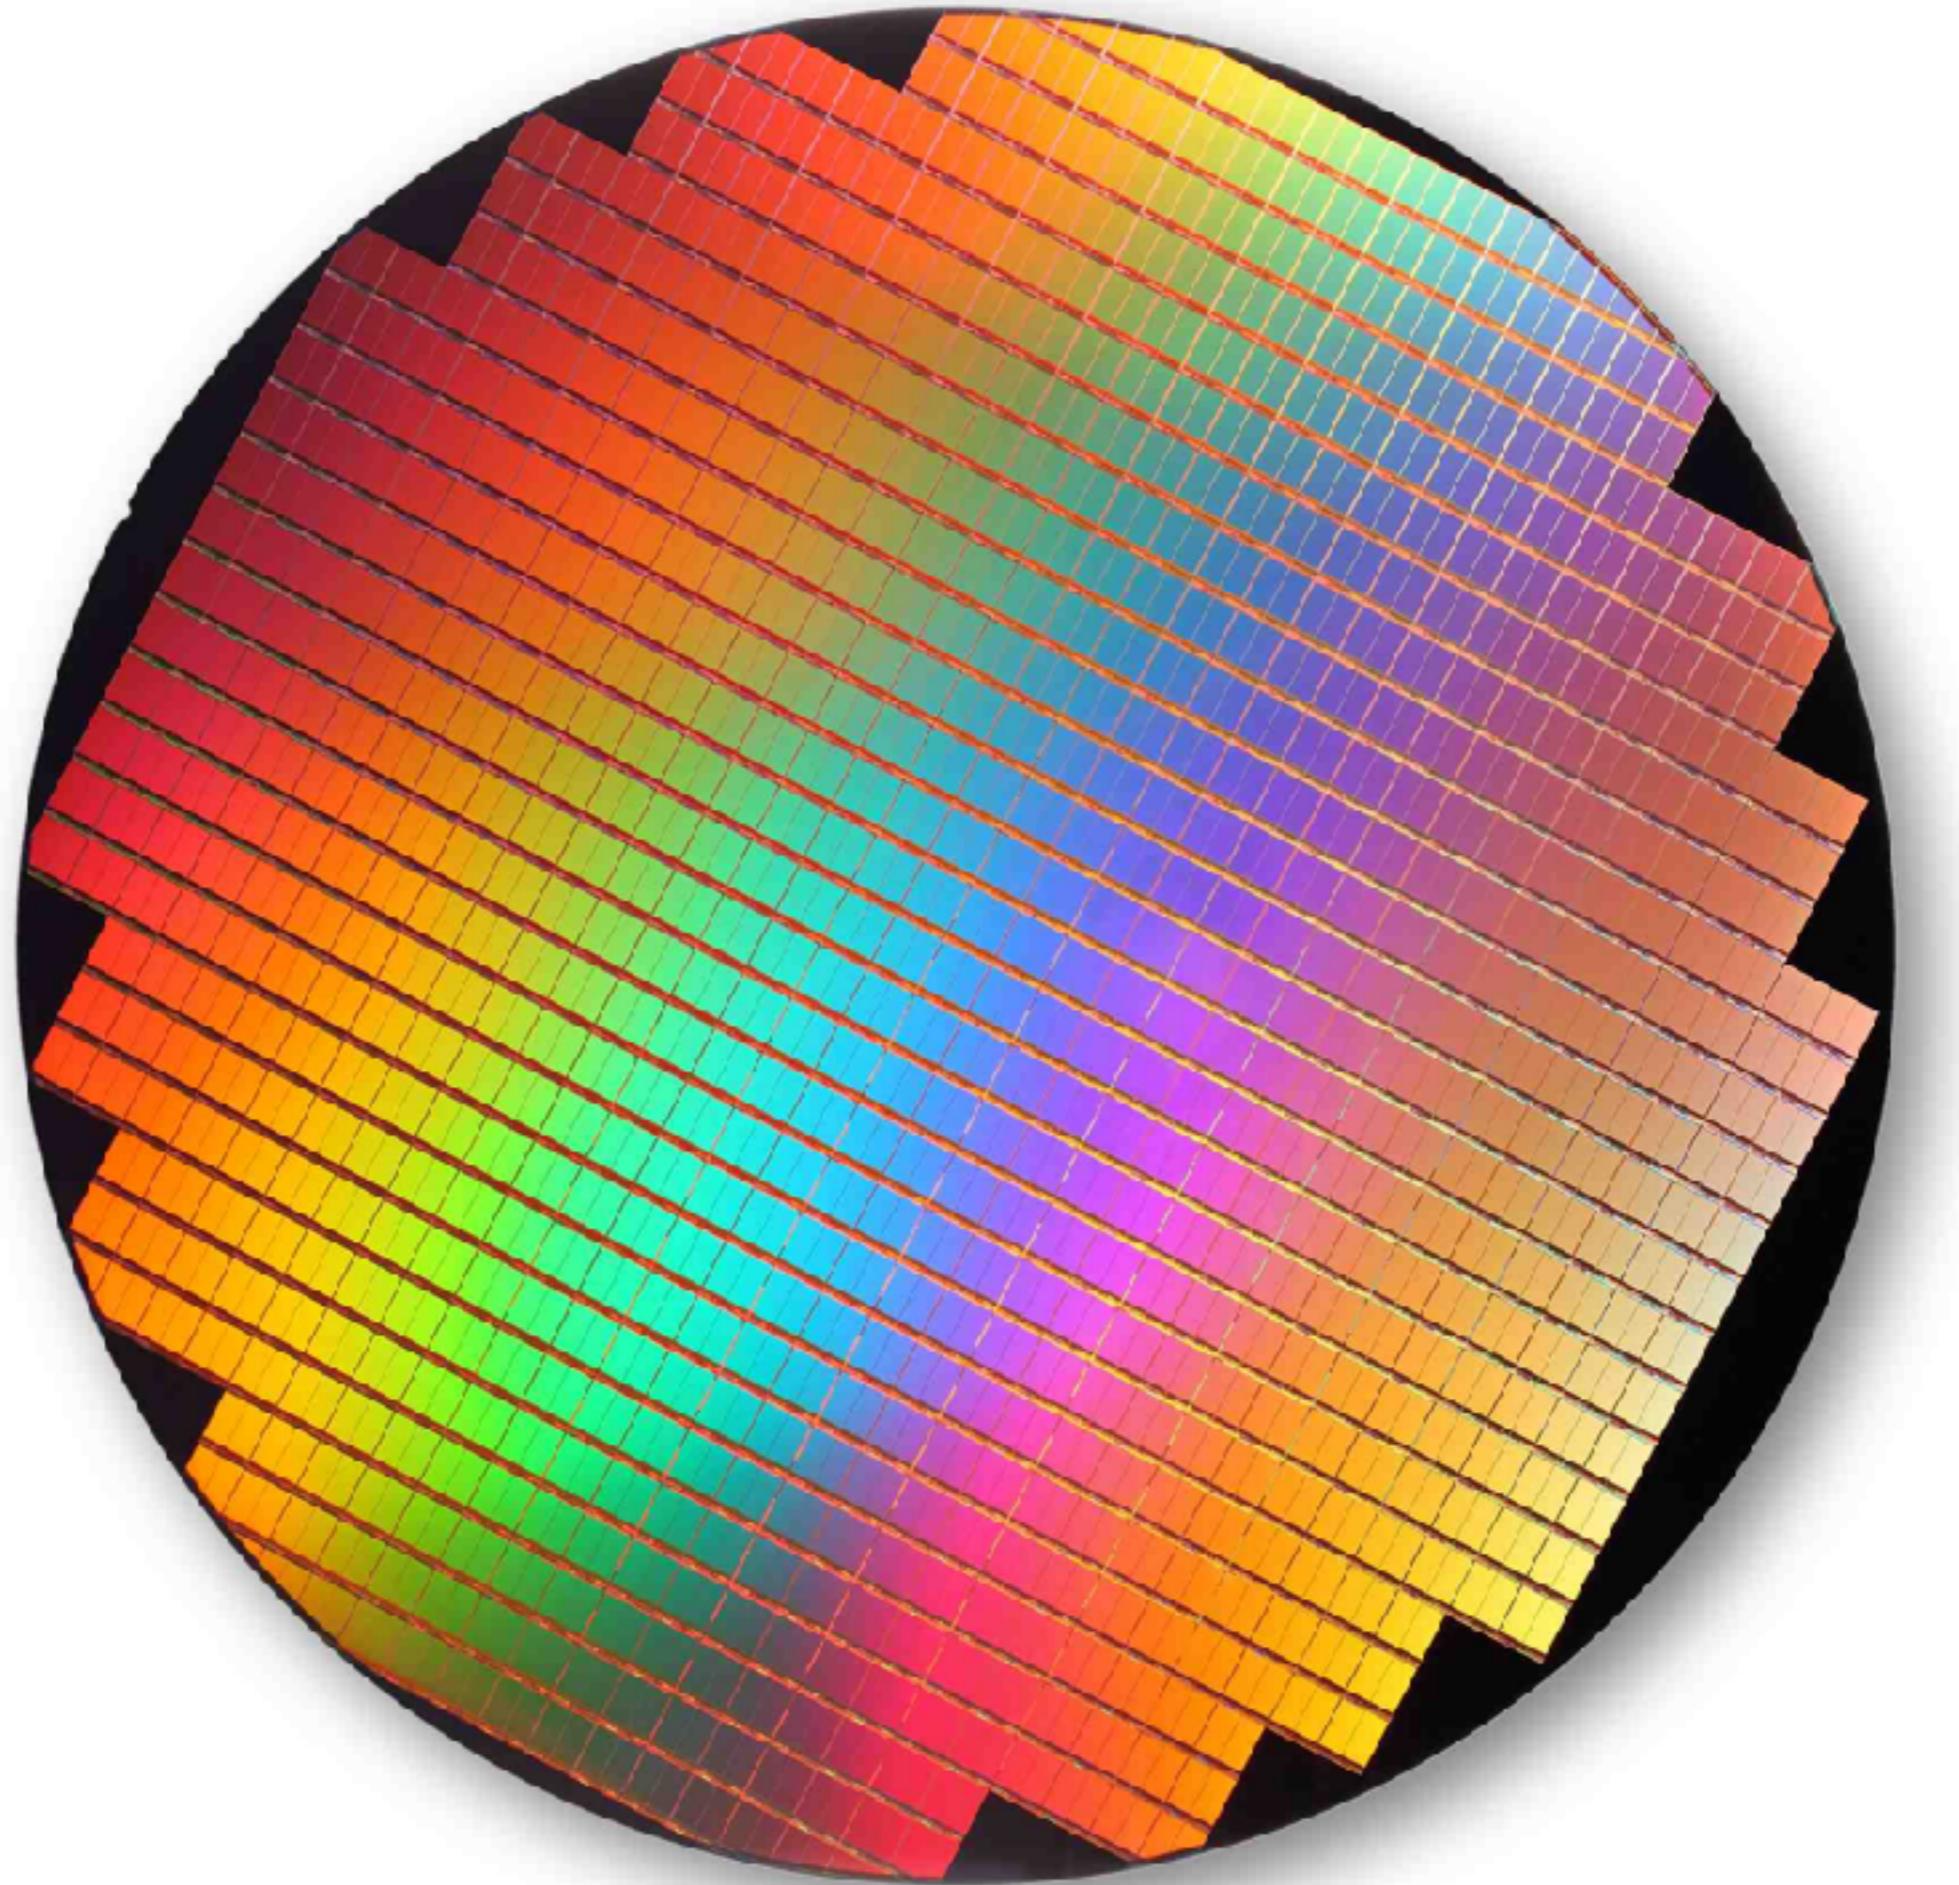
**Etch, Implant, etc.**

**Strip Resist**

# Types of microchip layers

- Semiconducting (doped silicon)
- Poor conductors (polysilicon)
- Insulators ( $\text{SiO}_2$ )
- Conductors / Interconnects (2 to 5 layers of metal, separated by insulators)
- Patterns specified by MASKs. Careful alignment required.
- 192 nm light, smaller features





# Fabrication process

- Lithographically lay out layers of doping, insulators, and interconnects
- Cut dice from wafers
- Test & reject
- Package

# Yield

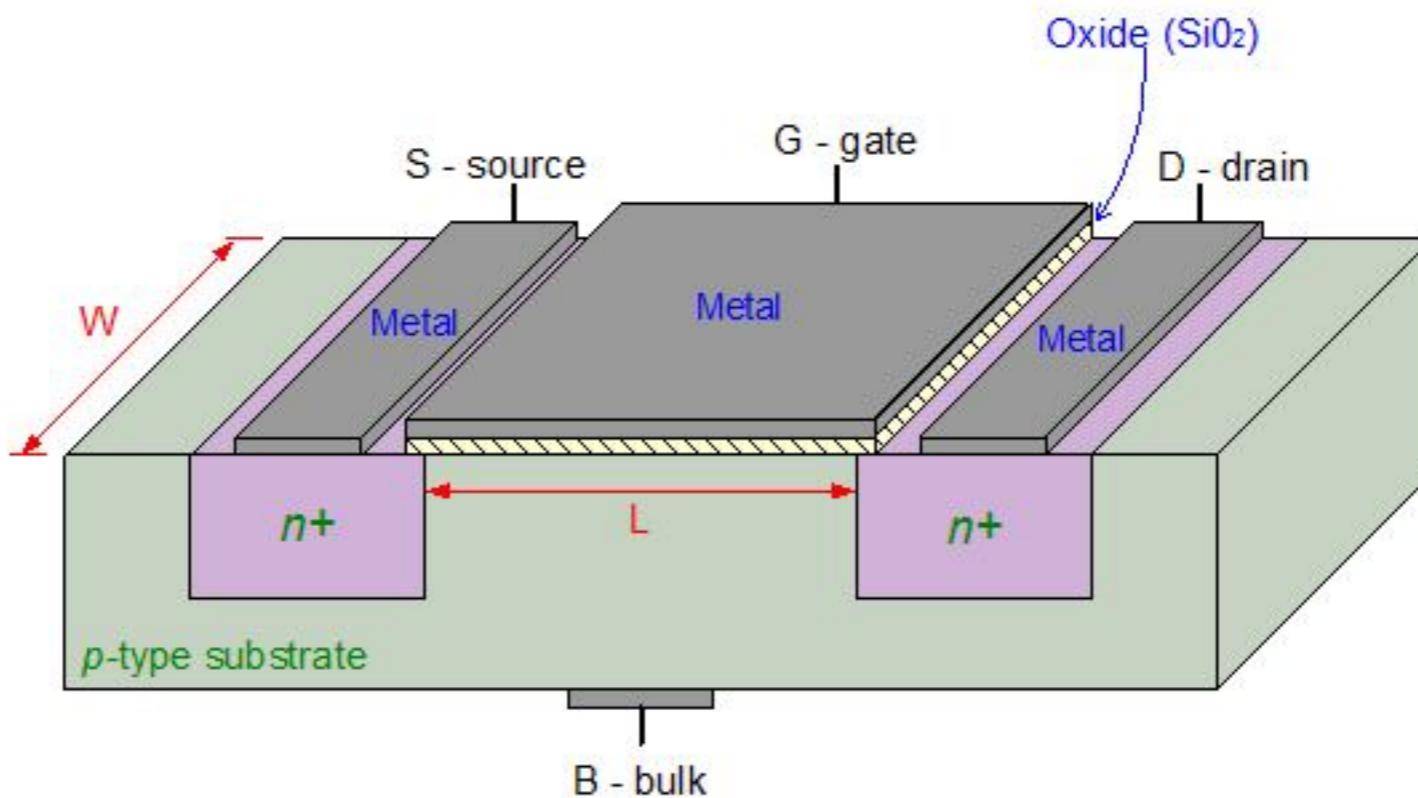
- Fraction of good dice to total
- Lower die area -- better yield
- More regularity, better mask alignment -- better yield

# **Circuitry, Tech Limitations, Power**

# Logic vs. DRAM

- CPU/logic/SRAM: Requires high speed, many layers of metal
- DRAM: Requires high density, few layers of metal, much cheaper

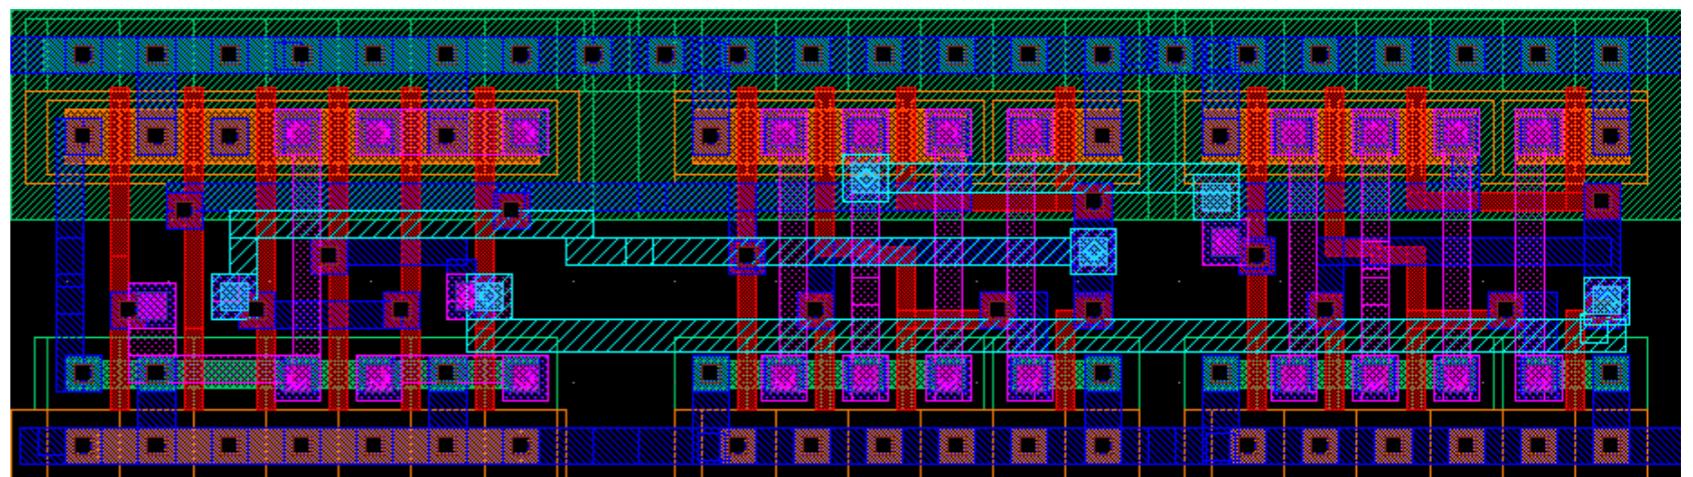
# Circuit area & speed

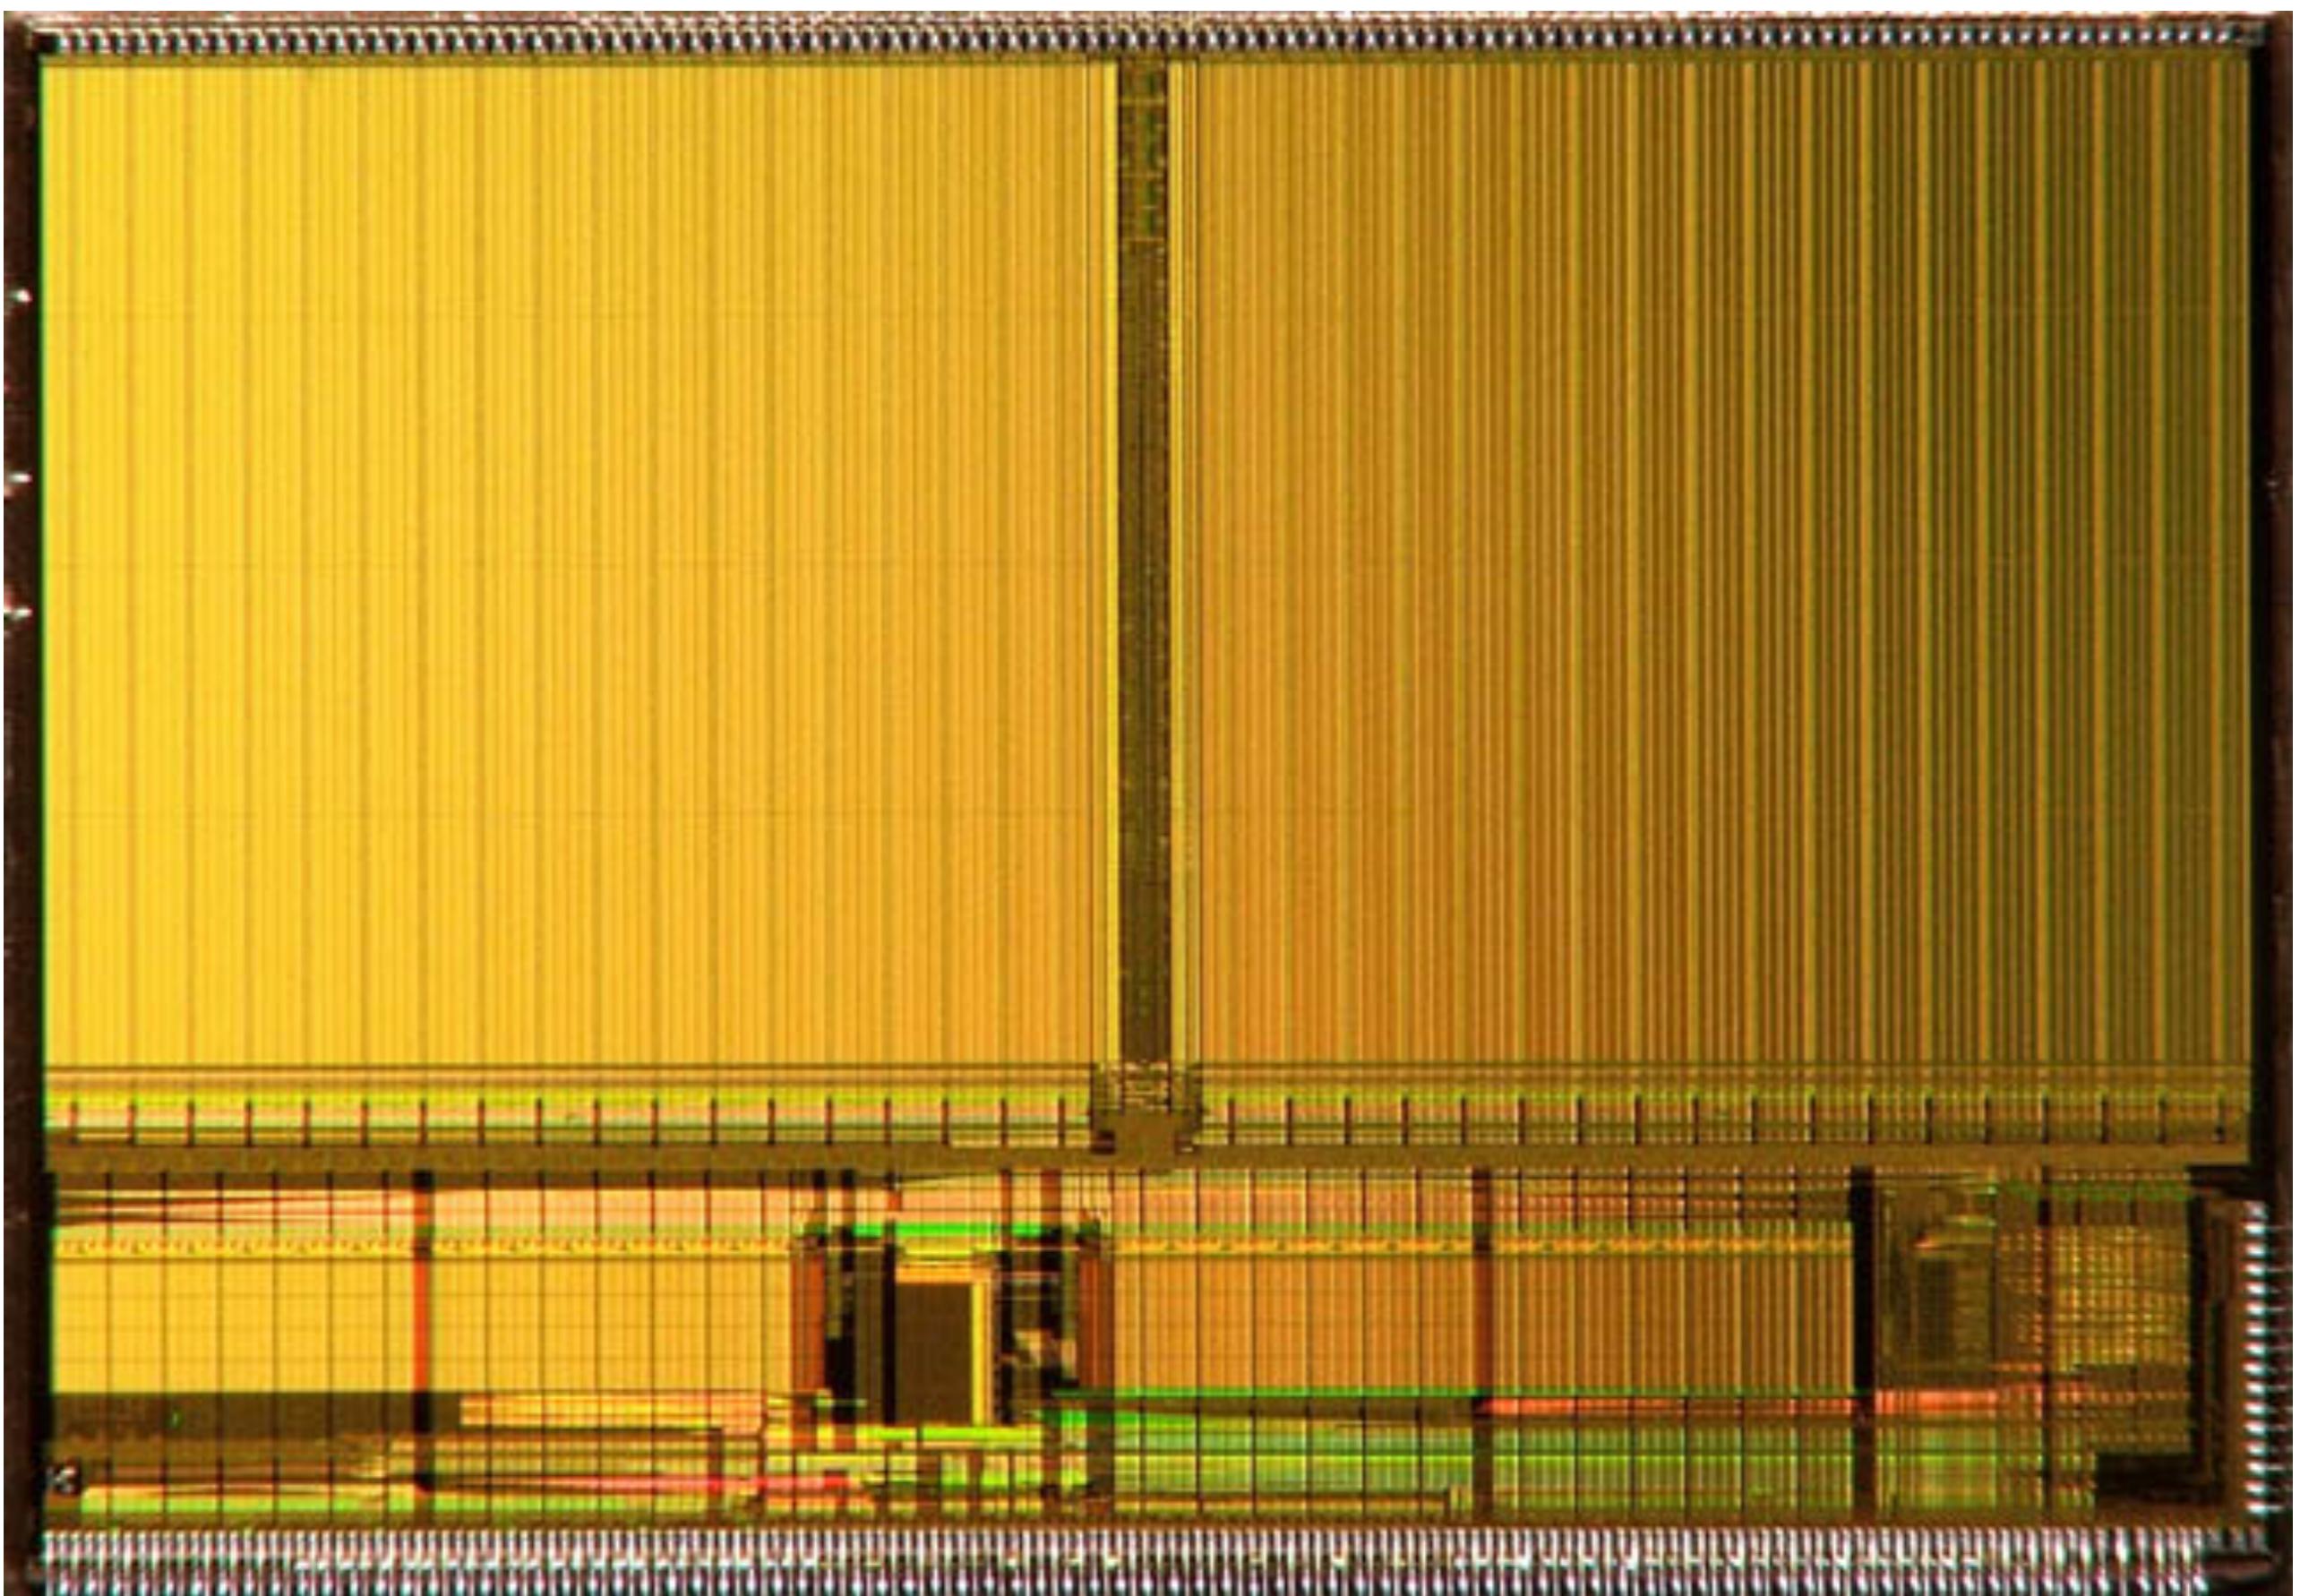


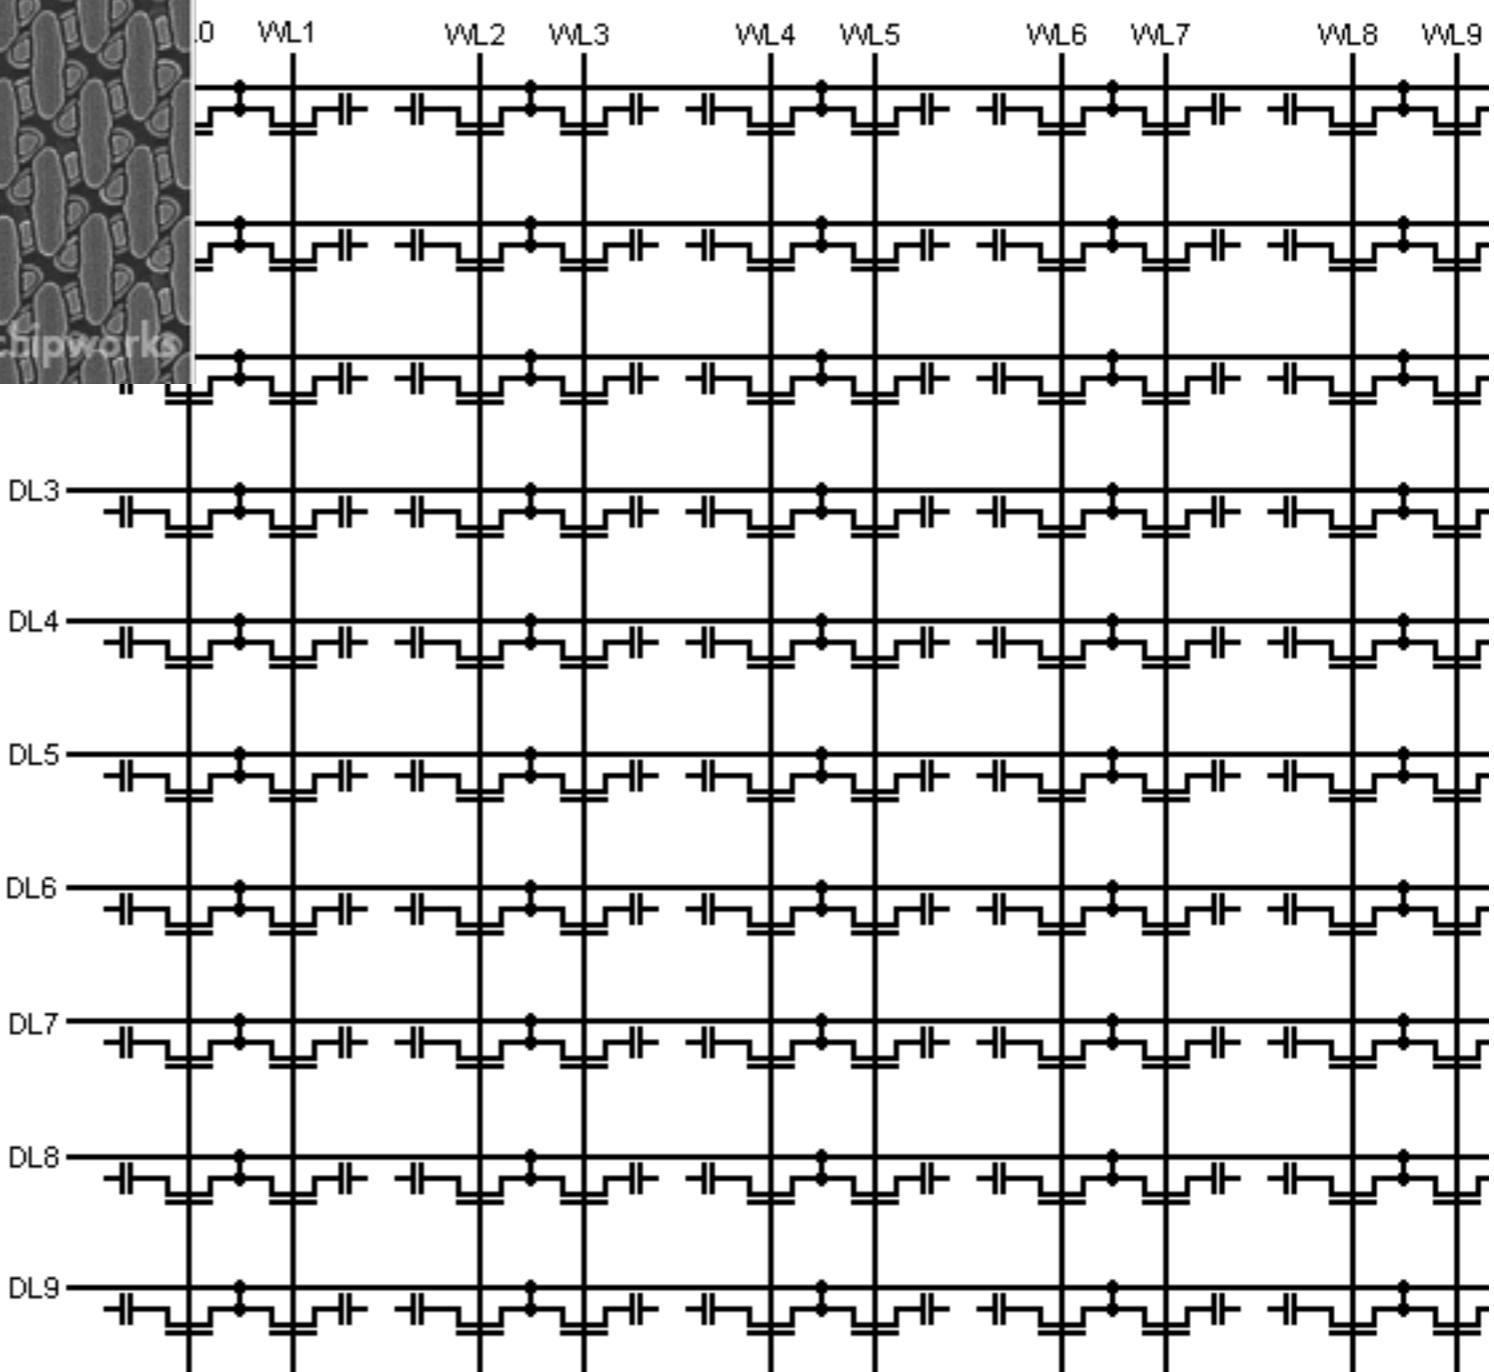
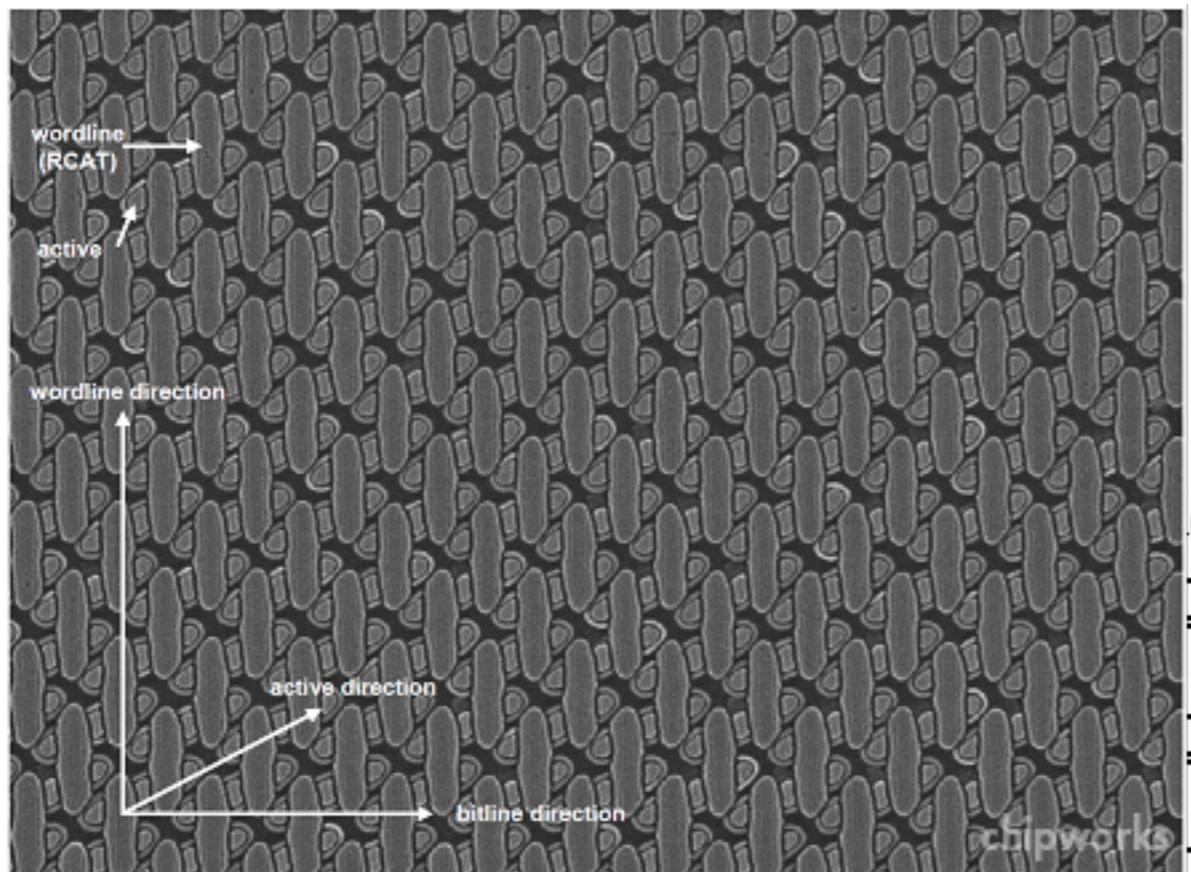
- Smaller transistors (channel length) are faster but use more energy
- Smaller circuits are faster, shorter wires
- Modern circuits dominated by wire delays

# Technological Constraints

- Design irregularities in logic limit maximum number of transistors per chip
- SRAMs and FPGAs more dense due to regularity
- DRAMs smallest at 1.5 B transistors per die







# Technological Constraints

- Off-chip I/O delay 10x to 50x on-chip
- Limited I/Os: Pin count max at around 2000; about half are pwr & ground.

# External I/O limitations

- Confine most activity to within chip
- Exploit spatial and temporal LOR
- Multiplexing of I/Os for multiple purposes
- Burst transfers, streaming, high-speed differential serial
- Clock multipliers

# Transistor count limitations

- Put only the most essential functions on a single chip.
- 1971 -- Simple 4-bit CPU, Intel 4004, first microprocessor
- Late 70's -- 16-bit general purpose CPU on single chip, no floating point, cache or virtual memory

# Transistor count limitations

- Early 80's -- 32-bit CPU, no cache or floating point. Some microcontrollers included limited RAM + communications. Some pipelined designs.
- Late 80's, early 90's -- 32-bit and 64-bit CPUs, on-chip floating point, single level of cache. Early superscalar.

# Transistor count limitations

- 1993 -- DEC 21164: 64-bit, on-chip floating point, two level cache, superscalar
- Late 90's -- multiprocessors, high-speed cluster networking, MCMs, VLIW, etc.
- 2000's -- SMT, multi-core, on-chip media processors, on-chip graphics, etc.
- Future -- Greater integration of CPU, RAM, chipset. Single-chip multiprocessors. Power-aware designs.

# Recent limitations

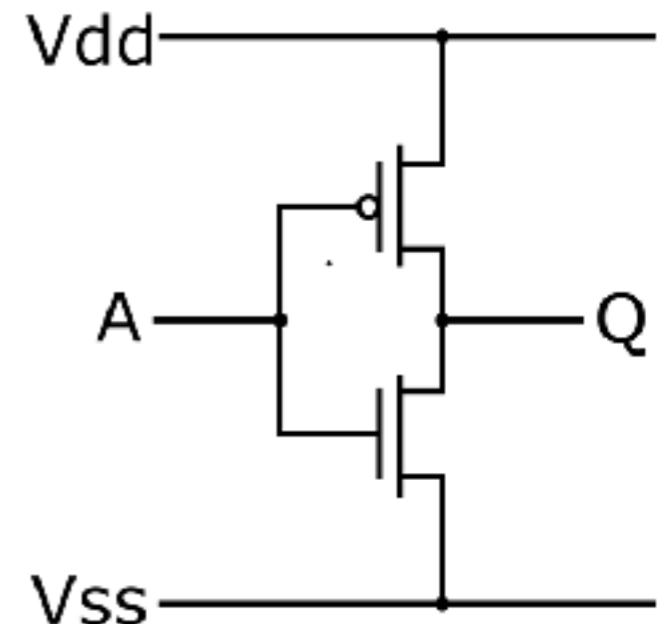
- Power wall -- performance limited by maximum heat dissipation
- Energy (charge) wall -- limited by battery life
- Dominance of interconnect delays, clock skew

# Peak power dissipation

<b>Processor</b>	<b>Peak Power (W)</b>	<b>Clock rate (MHz)</b>
Alpha 21264C	95	1001
AMD Athlon XP	67	1800
AMD Athlon 64 3800+	89	2400
IBM Power 4	135	1300
Intel Pentium 4 (dual MT)	115	3600
Intel Itanium 2	130	1000
Intel Xeon 5680 (6-core)	130	3333
AMD FX-8170 Bulldozer (8-core, Sept 2011)	125+	3900

# Power components

- Dynamic/switching power -- Energy is used every time a CMOS logic gate switches.
  - $P = \alpha f V^2$
- Static power -- Subthreshold leakage current
  - Increases as transistors shrink, increases with temperature, linear with voltage



# Problems with high power dissipation

- Decreased processor lifetime
- Hotspots
- Limits performance
- Cooling costs

# Trends in Fabrication and Processor Design

# CPU/Logic/SRAM trends

- Logic transistor counts increase 60% to 80% per year
  - Moore's law -- transistor count doubles every 18 months
- Switching speeds double every 3 years
  - Trend has slowed since 2001

# DRAM trends

- Densities increase 60% per year (4x in 3 years)
- Access times have increased by only about 30% in the last 10 years.

# External buses on PCB

- Speed trends: 25MHz in 1990's to 1GHz+ in 2000's, DDR.
- 500Mbits/sec on parallel bus in late 1990's to 4 Gbits/sec in 2000's using point-to-point differential signaling.

# Recent Trends high-end processors

- Multithreading (SMT)
- Virtualization
- Adaptations to power wall
  - Throttling, more simpler cores, slower clocks, clock/power gating
- Multi-core (CMP, SMP)

# Recent Trends Embedded processors

- Integration
  - Memory, DSPs, programmable logic
- Multi-core
- Array processors (video, signal processing, network, graphics)
- Micro LMS3S10I: \$1 at 10K unit pricing
- Dark silicon

# Performance Considerations

# Performance factors

$$T_{exec} = N \times \tau \times CPI$$

- $N$  = Instructions executed
- $\tau$  = Clock period
- CPI = Average clocks per instruction

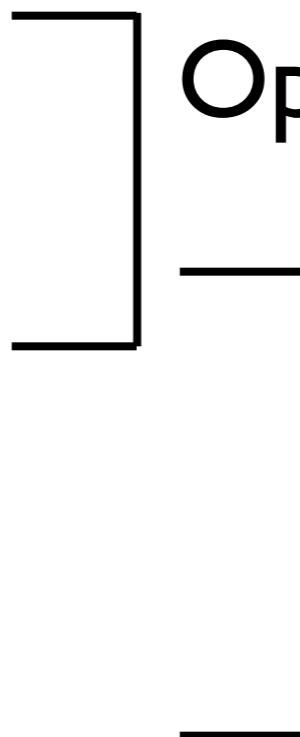
# Reduced Instruction Set Computer (RISC)

- Reduce Texec by
  - Reducing CPI
  - Reducing  $\tau$

# Complex Instruction Set Computer (CISC)

- Reduce Texec by:
  - Reducing N

# Characteristics of RISC ISA

- Load-store architecture
  - Simple instruction semantics
  - Few addressing modes
  - Few instruction formats
- 
- Op times are fast
- Quick fetch and decode

# Consequences of RISC

- Simpler hardware
- Smaller logic delays -- faster clock
- N goes up
- Ensure that growth in N is balanced by reduction in CPI and  $\tau$ .

# Characteristics of CISC CPU

- ALU instructions can reference memory
  - High-level semantics
  - Many addressing modes
  - Non-uniform instruction formats
- Op times are slow
- Large decode times
- Large fetch times --  
memory alignment
- 
- The diagram consists of four black circular bullet points arranged vertically. From the right side of the third point, a horizontal line extends upwards to the first two points. From the right side of the fourth point, a diagonal line extends upwards and to the left to the third point, while a vertical line extends downwards to the second point.

# Consequences of CISC

- Complex hardware
- Larger logic delays
- Higher CPI
- Philosophy: Amortize cost of fetch & decode by performing more work per instruction
- Challenge: Ensure increase in CPI and  $\tau$  is compensated by drop in  $N$

# Locality of Reference (LOR)

- Temporal LOR: Repeated access of same address
- Spatial LOR: Likely access of nearby addresses
- Basis for most performance optimization techniques

# Examples of LOR

• Mechanism	<i>Type of LOR Exploited</i>
• Registers	Temporal
• Ins. Cache	Temporal & Spatial
• Data Cache	Temporal & Spatial
• Pipelining	Spatial
• Interleaved mem	Spatial

# Pipelining and the “APEX” processor

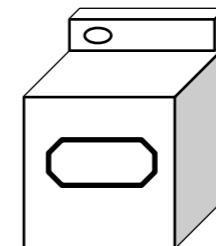
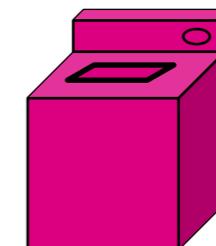
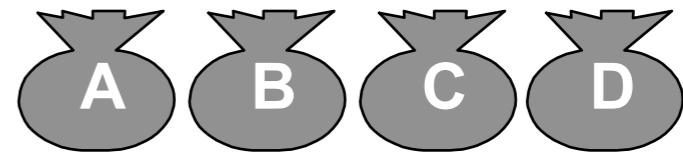
# Pipelining

- Divide labor among multiple hardware functional units
- Route work through chain of units
- Analogy: Assembly line or fluid flow through pipeline
- Up to K consecutive tasks can be in various states of processing simultaneously
- Logic for each step is called a **stage**.

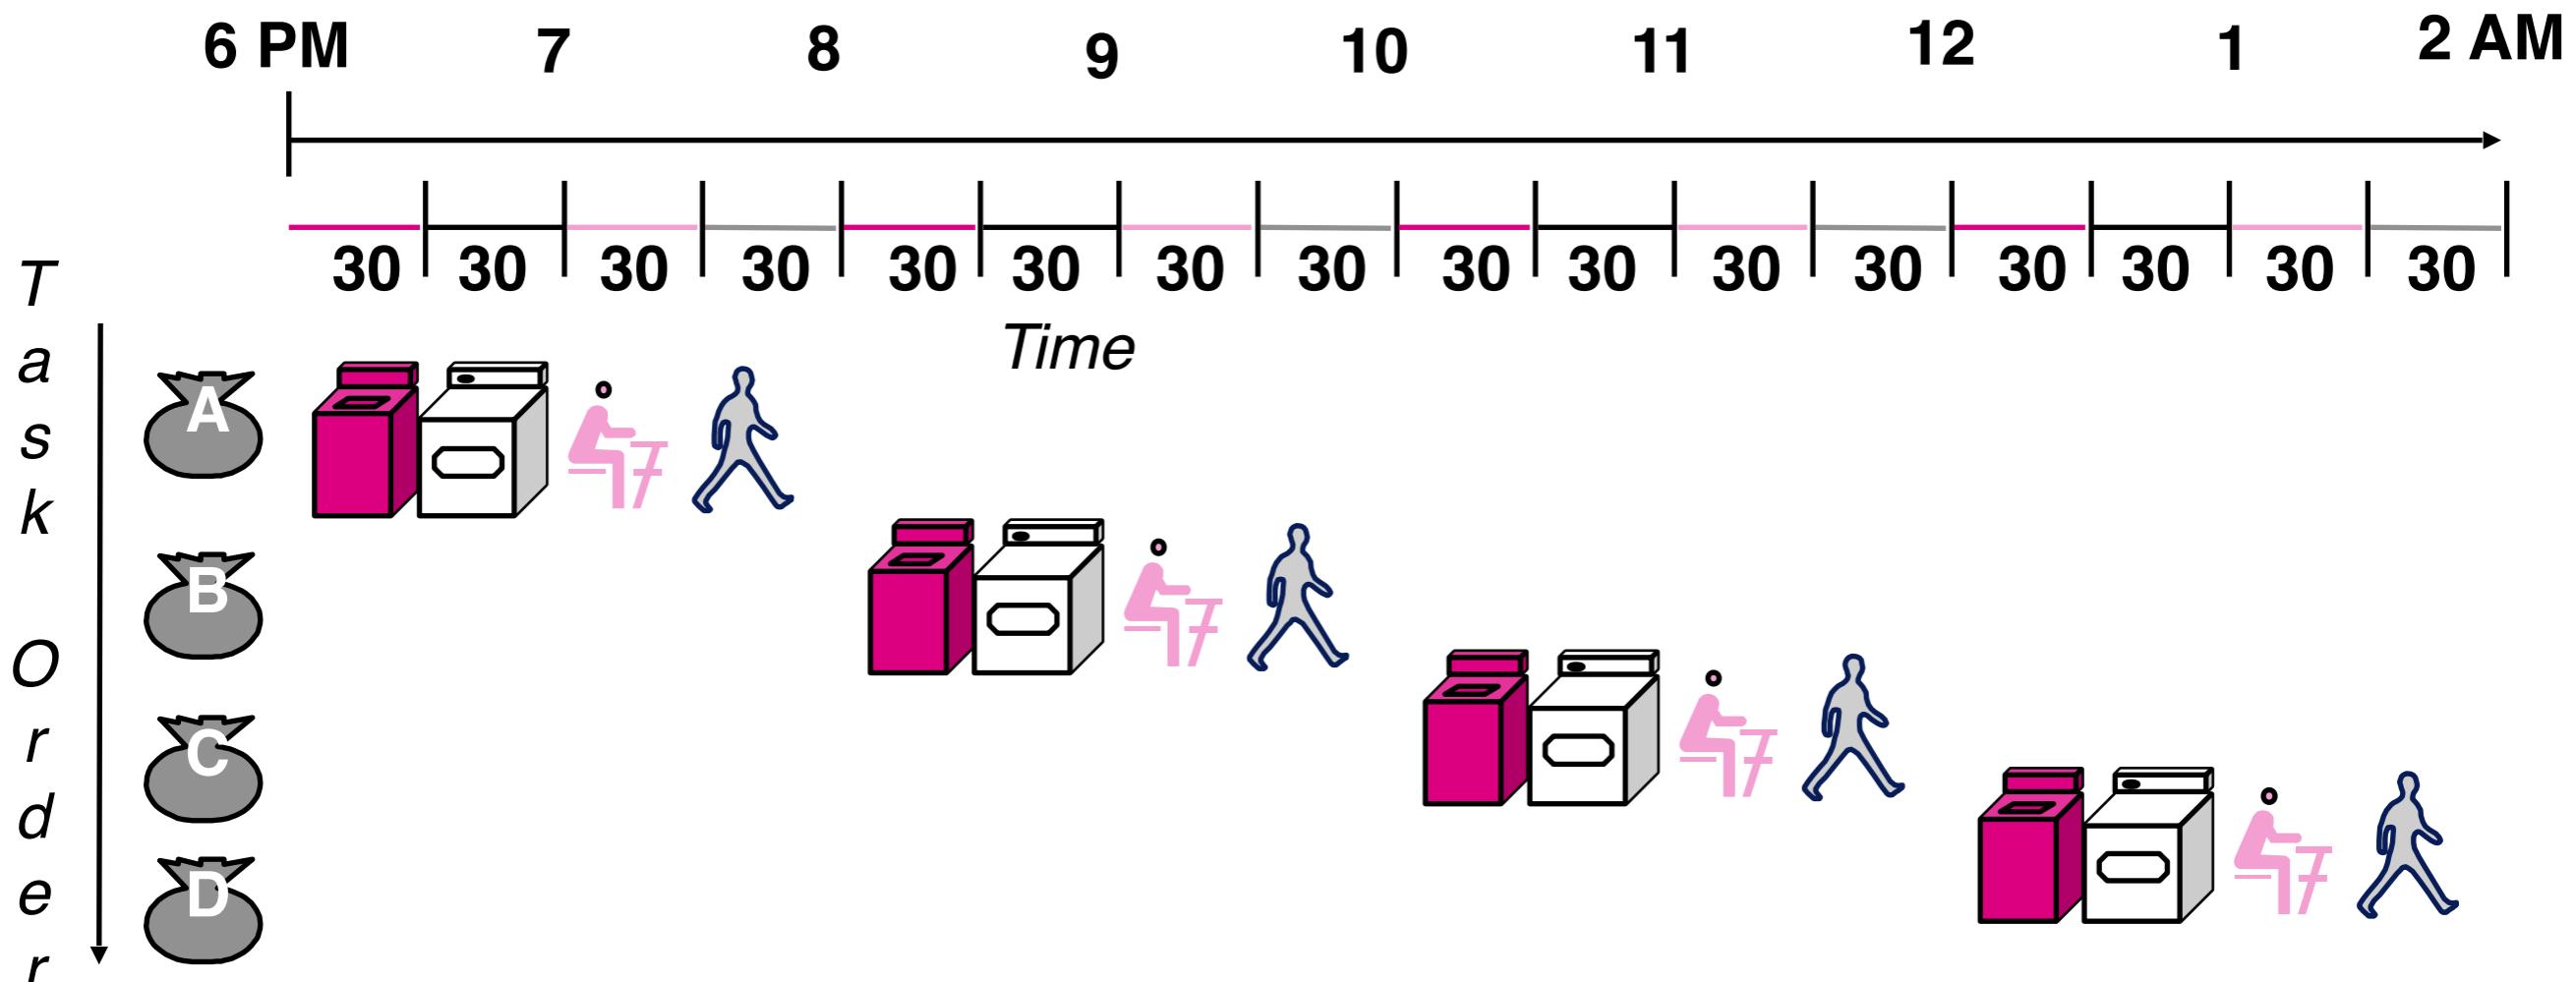
# Pipelining is Natural!

---

- Laundry Example
- Ann, Brian, Cathy, Dave each have one load of clothes to wash, dry, and fold
- Washer takes 30 minutes
- Dryer takes 30 minutes
- “Folder” takes 30 minutes
- “Stasher” takes 30 minutes to put clothes into drawers

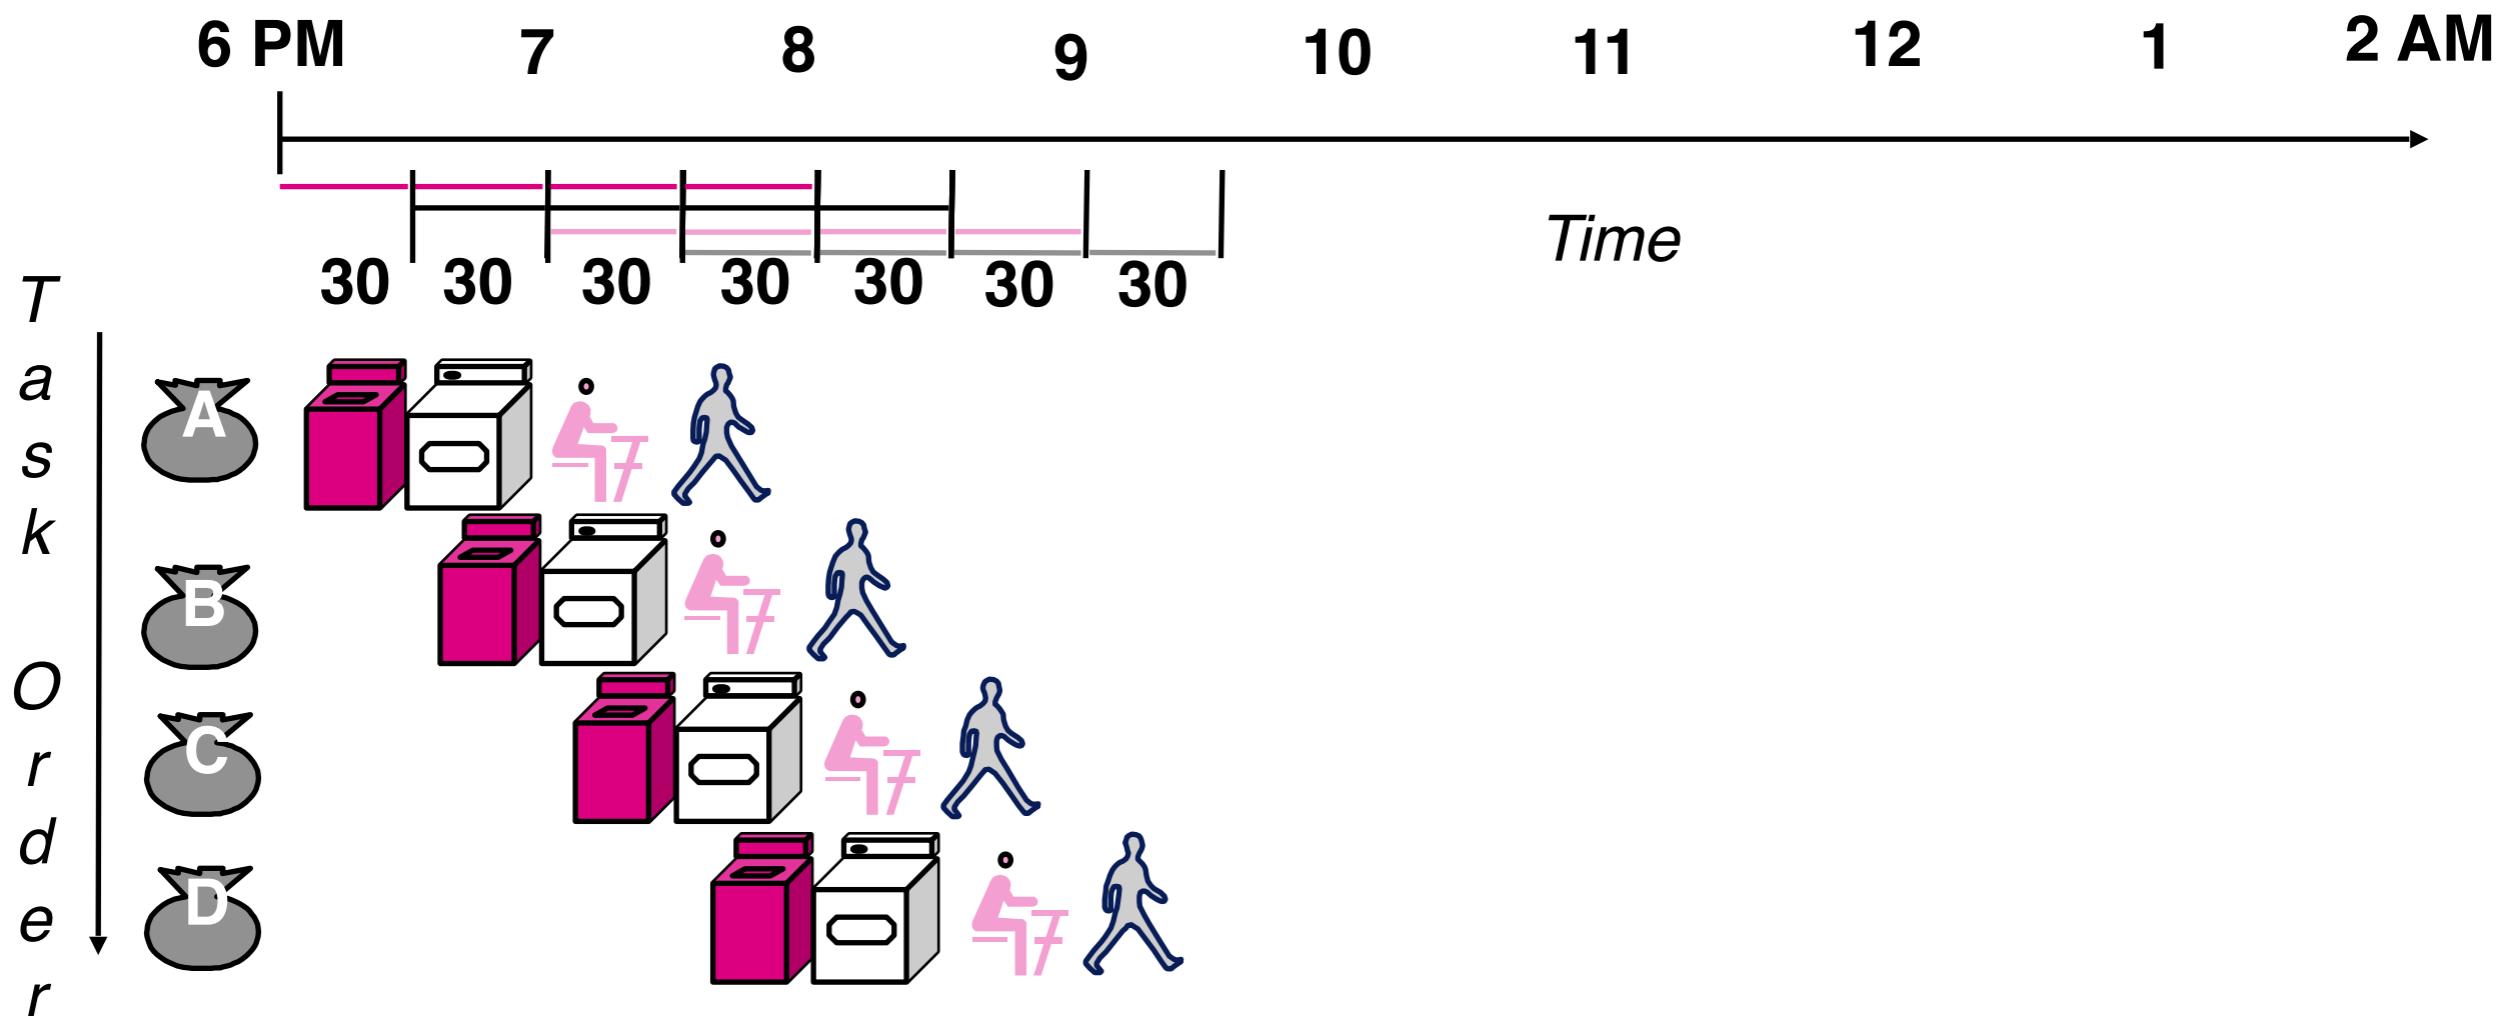


# Sequential Laundry



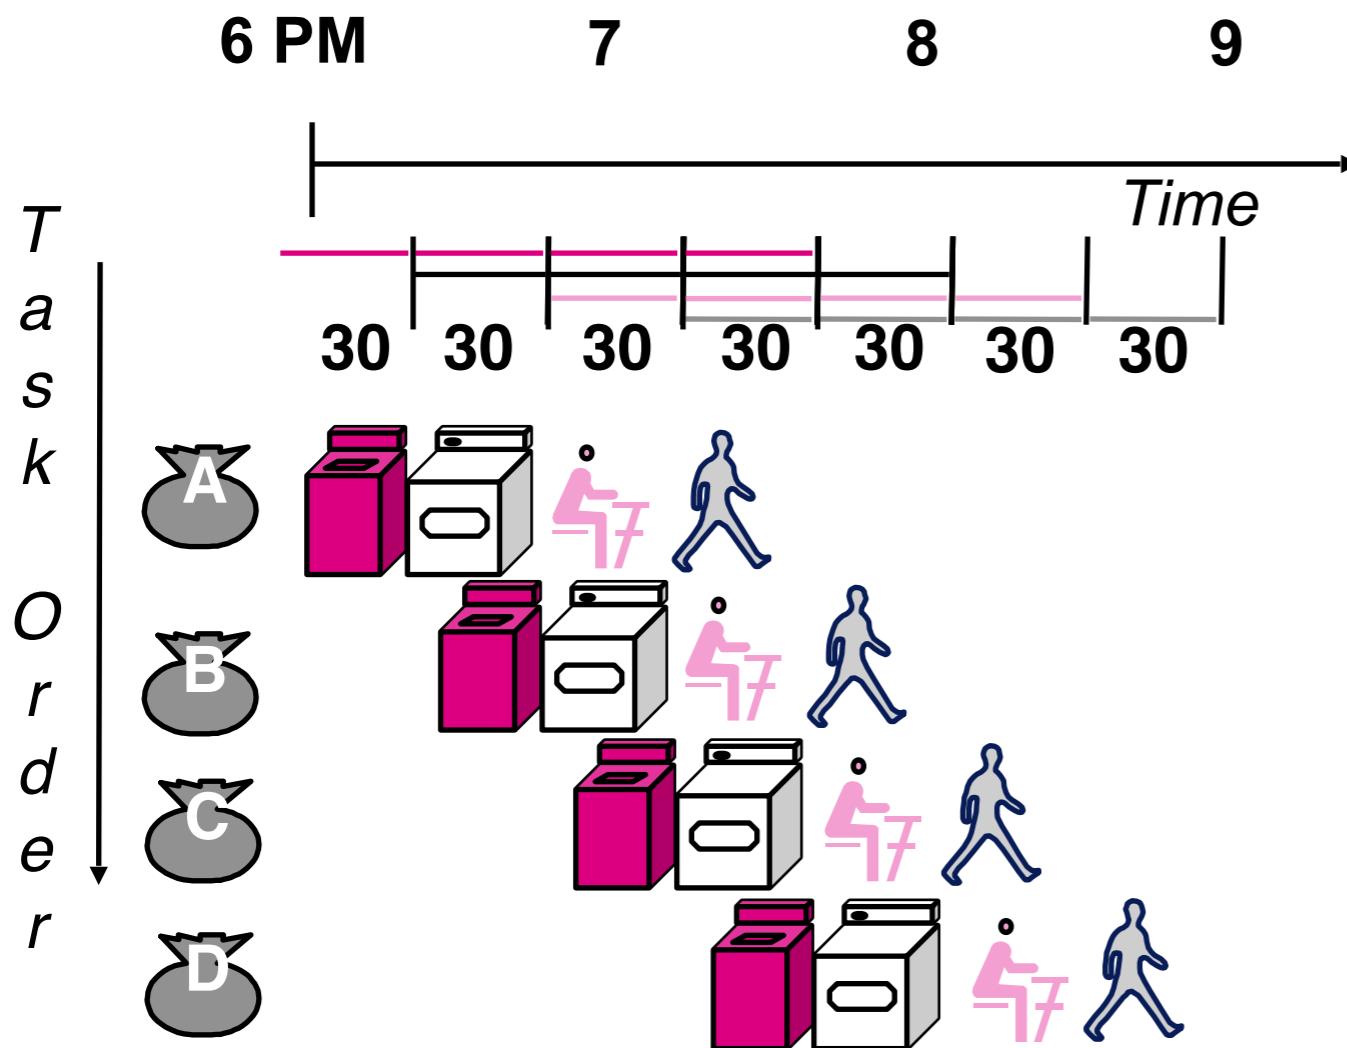
- ° Sequential laundry takes 8 hours for 4 loads
- ° If they learned pipelining, how long would laundry take?

# Pipelined Laundry: Start work ASAP



◦ Pipelined laundry takes 3.5 hours for 4 loads!

# Pipelining Lessons

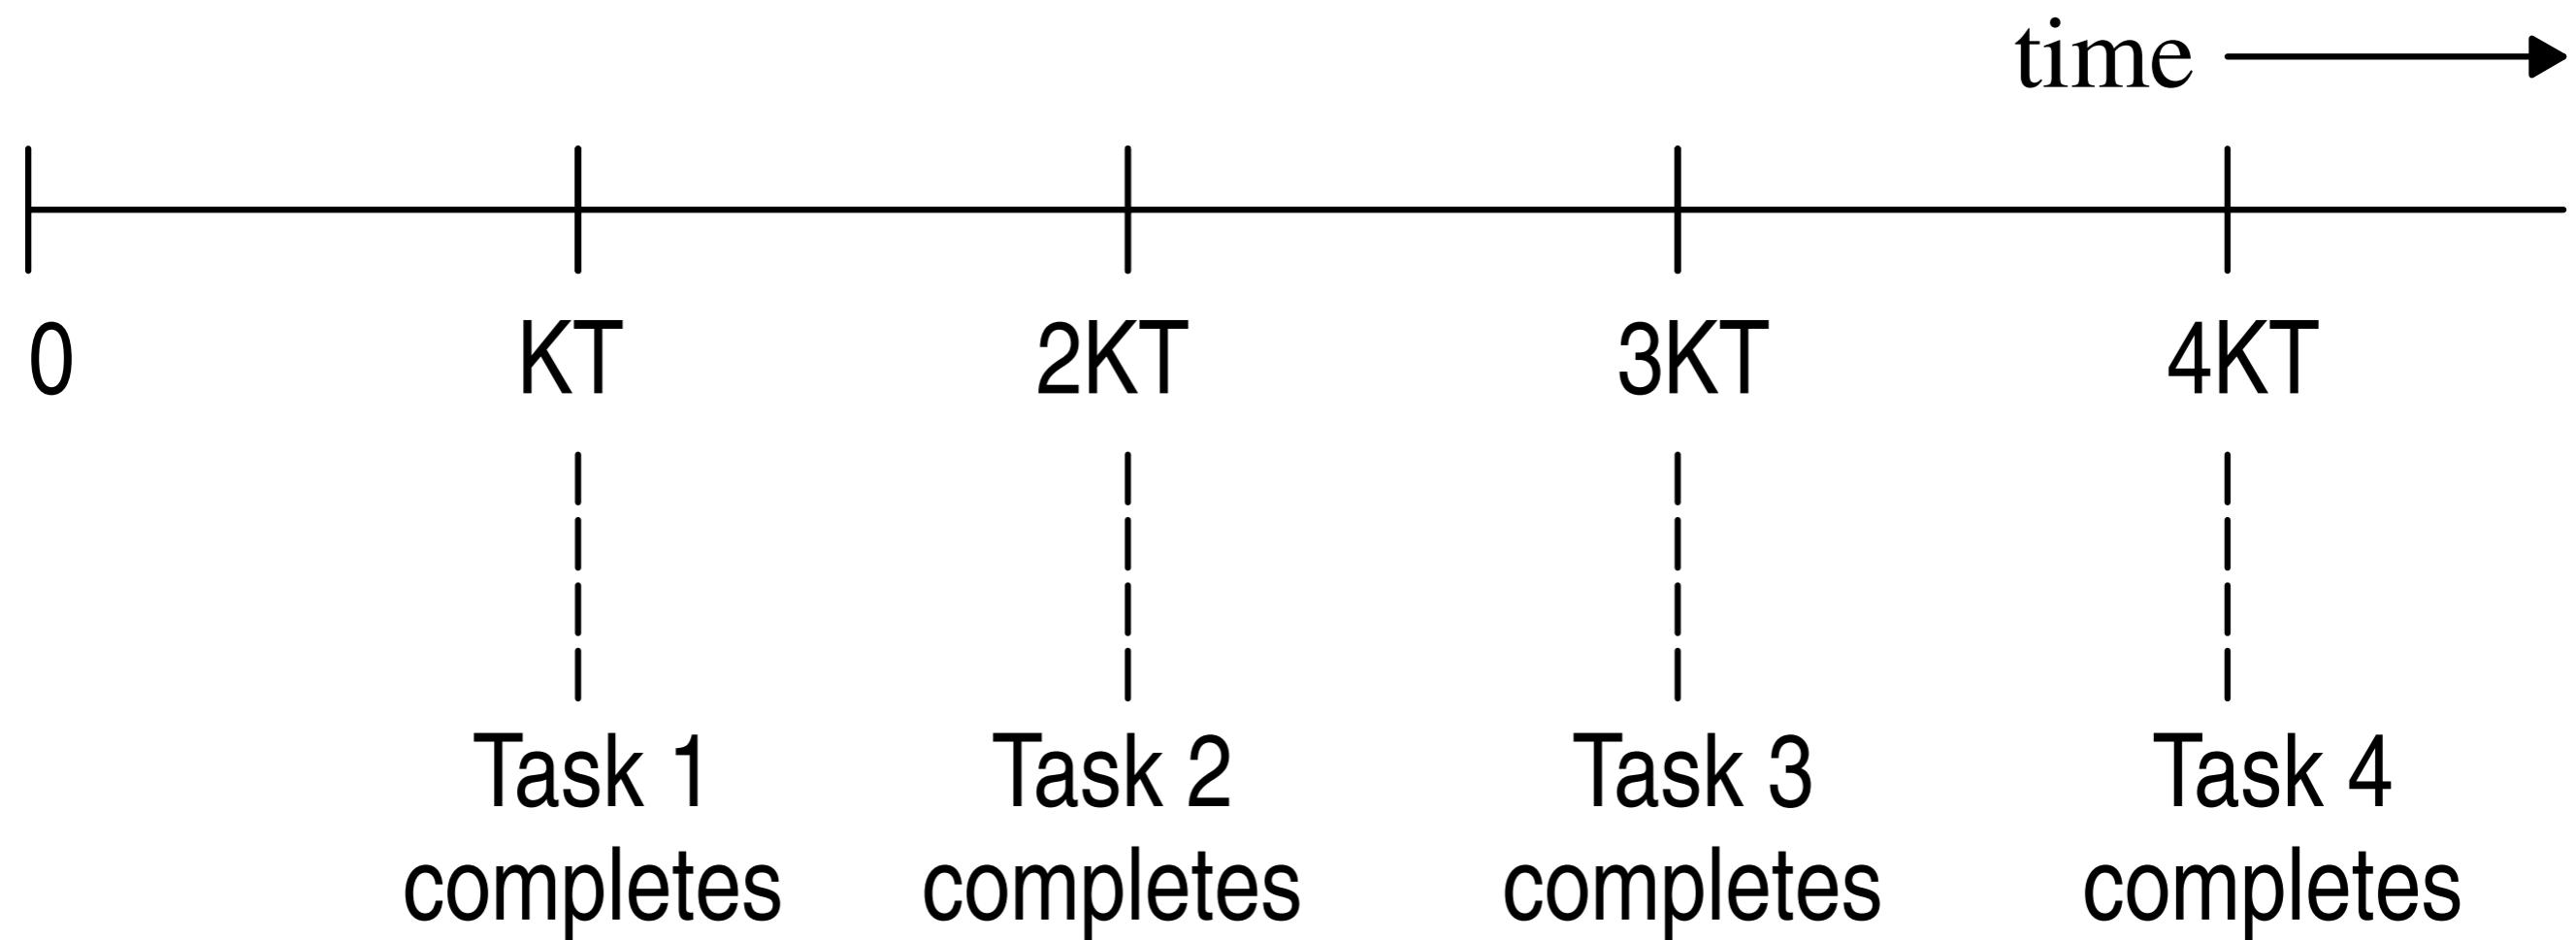


- Pipelining doesn't help **latency** of single task, it helps **throughput** of entire workload
- Multiple tasks operating simultaneously using different resources
- Potential speedup = **Number pipe stages**
- Pipeline rate limited by **slowest** pipeline stage
- Unbalanced lengths of pipe stages reduces speedup
- Time to “**fill**” pipeline and time to “**drain**” it reduces speedup
- Stall for Dependencies

# Improving throughput by Pipelining

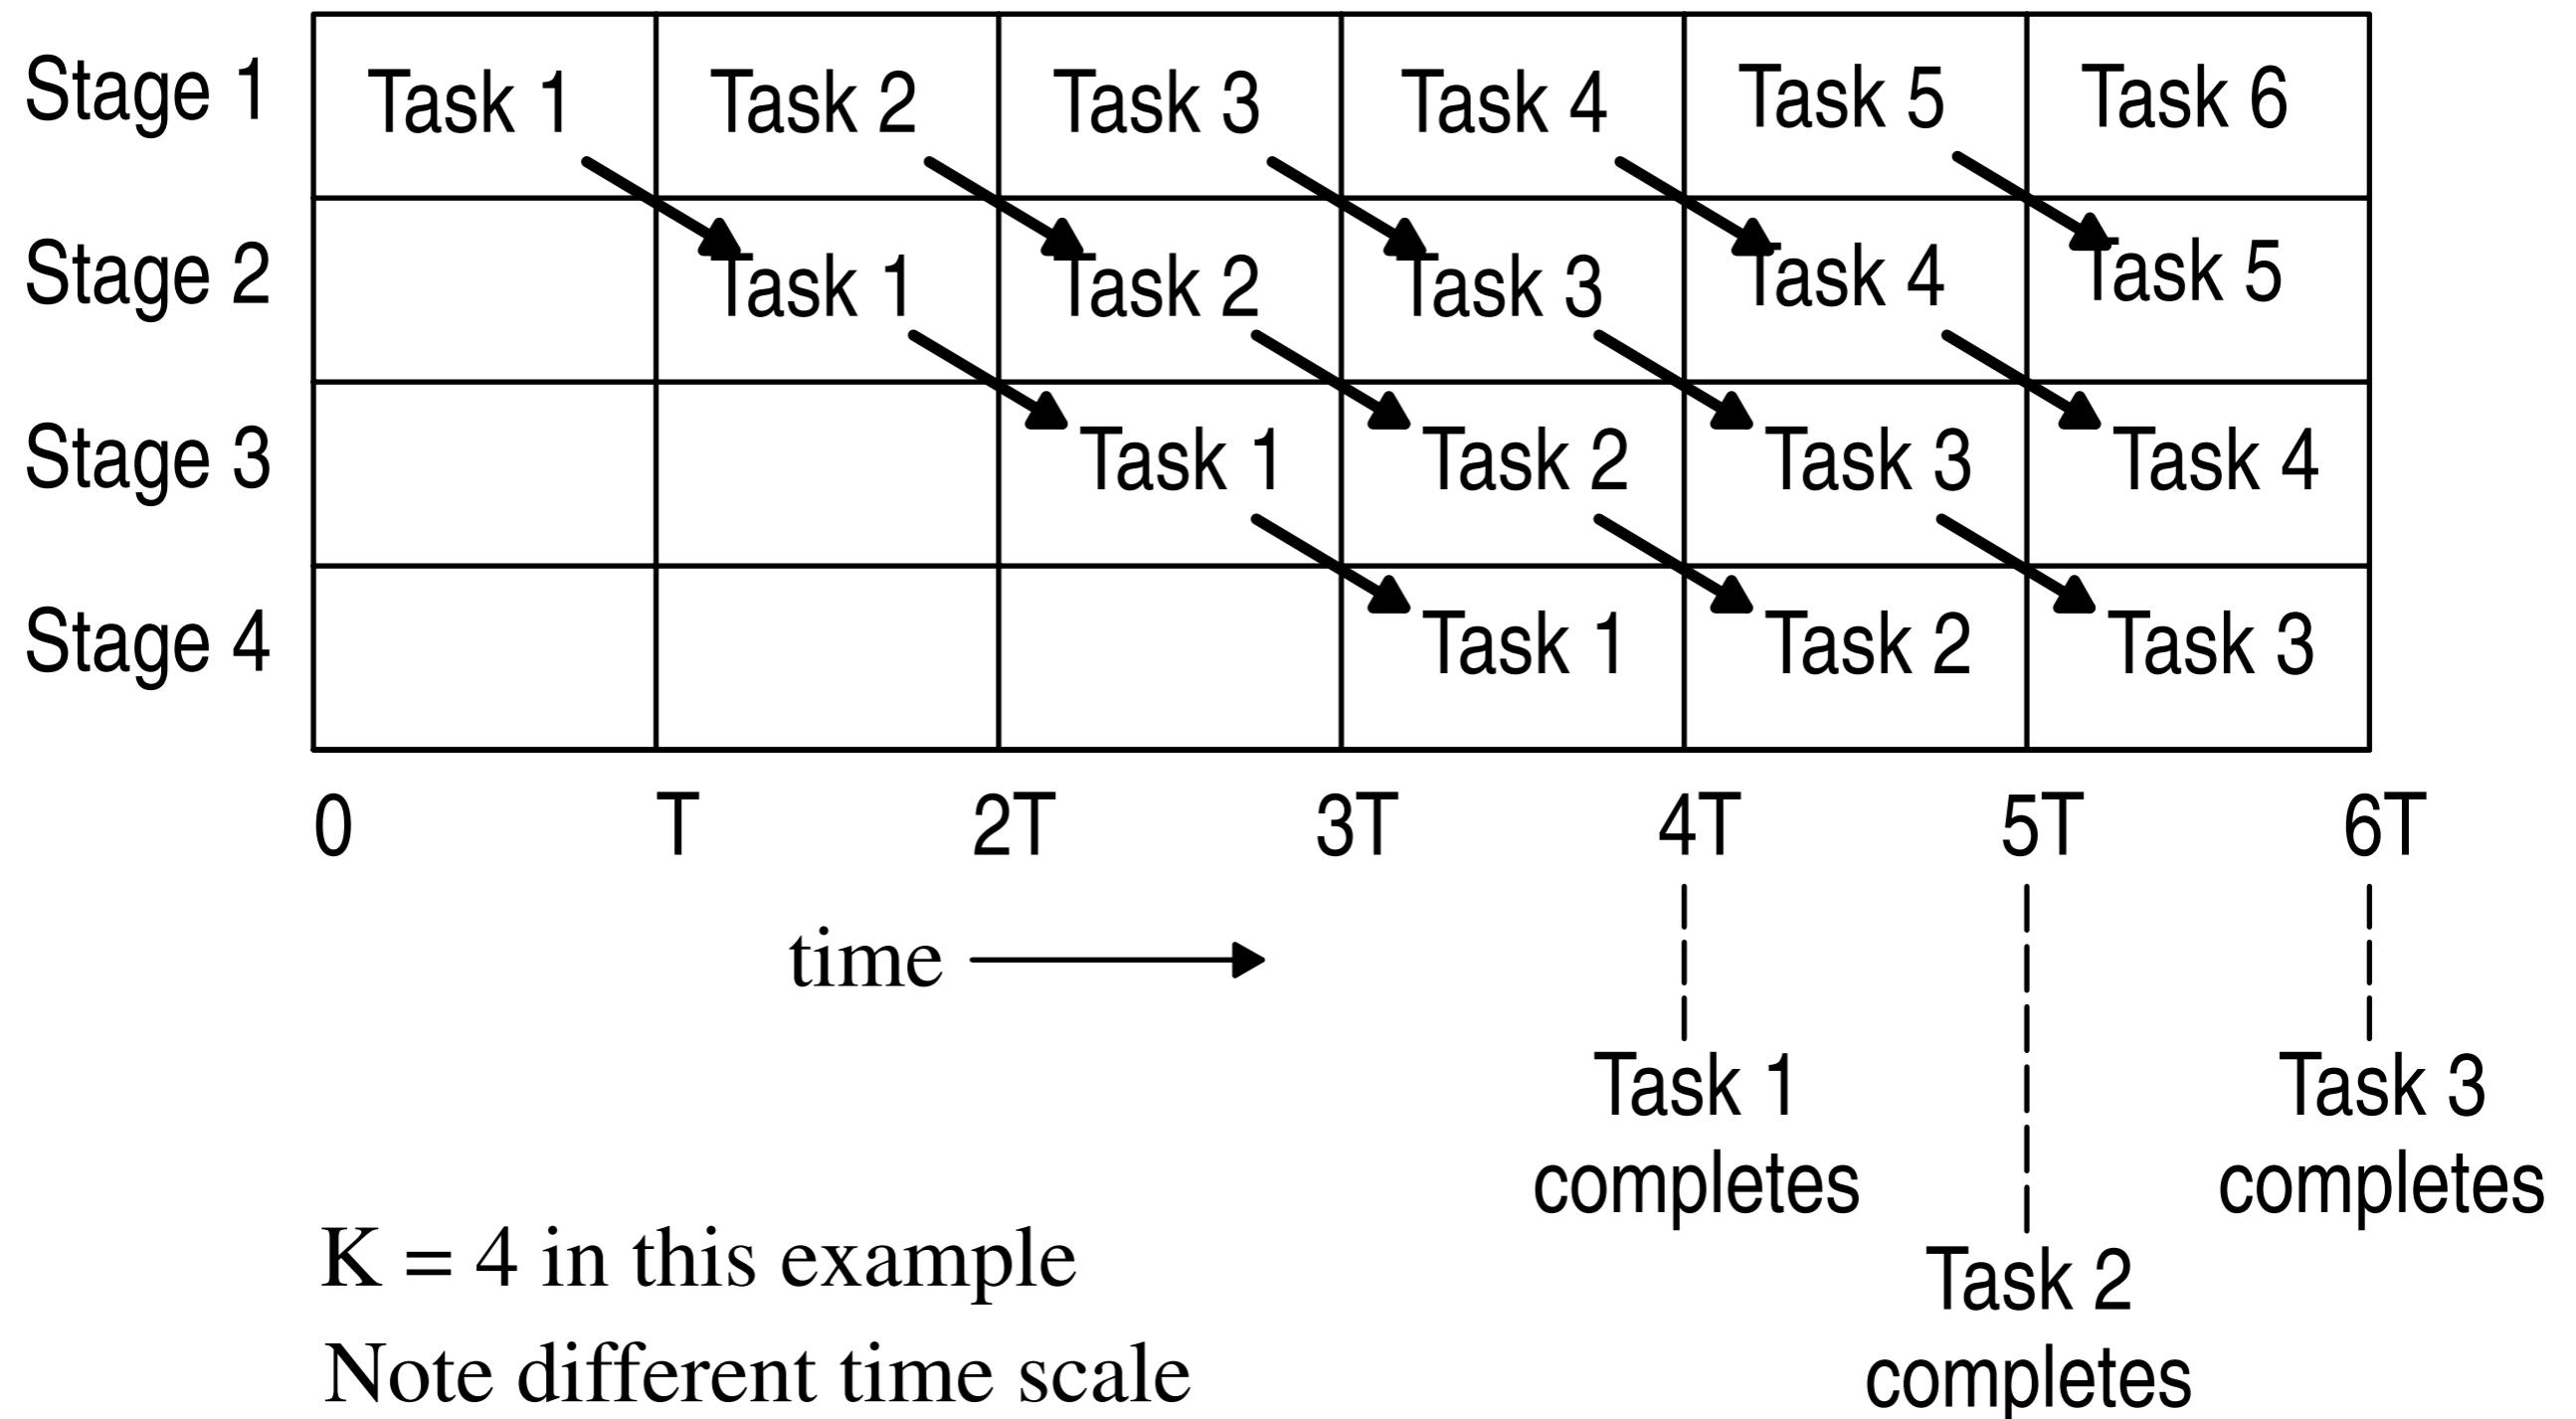
- N processing tasks
- K steps per task
- T duration each step
- Non-pipelined:  $N * (K * T)$

$$T_{np} = N * (K * T)$$



# Without pipelining

- No overlap in processing of consecutive tasks
- Throughput = completion rate = one task every  $K*T$  time units
- Processing latency per task = time between start and end =  $K*T$



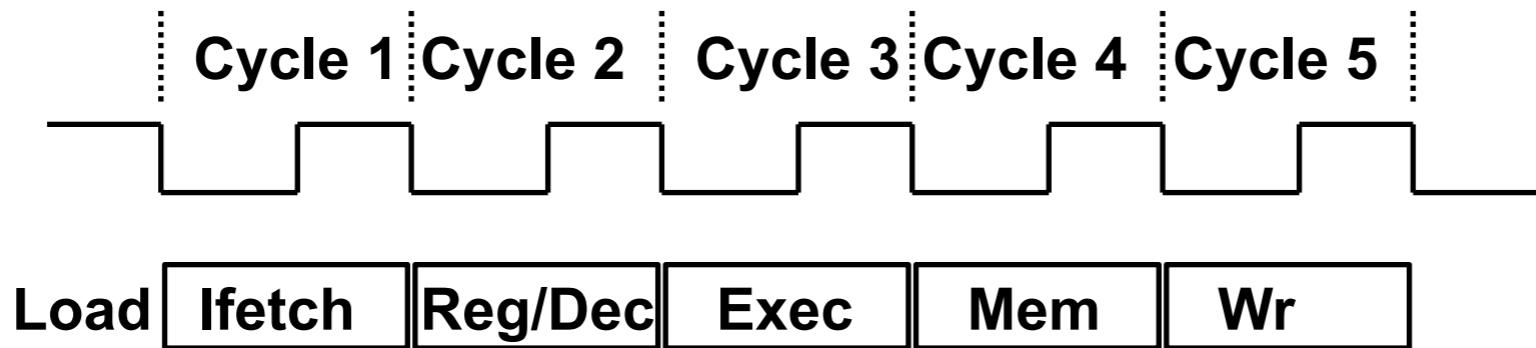
# Pipelining

- Total processing time
  - $T_p = K * T + (N-1)*T$
- First task takes  $K*T$ , subsequent complete every  $T$
- Throughput = one task per  $T$  time units, improved by a factor of  $K$
- Latency =  $K*T$  (unchanged)

# Instruction Pipelining

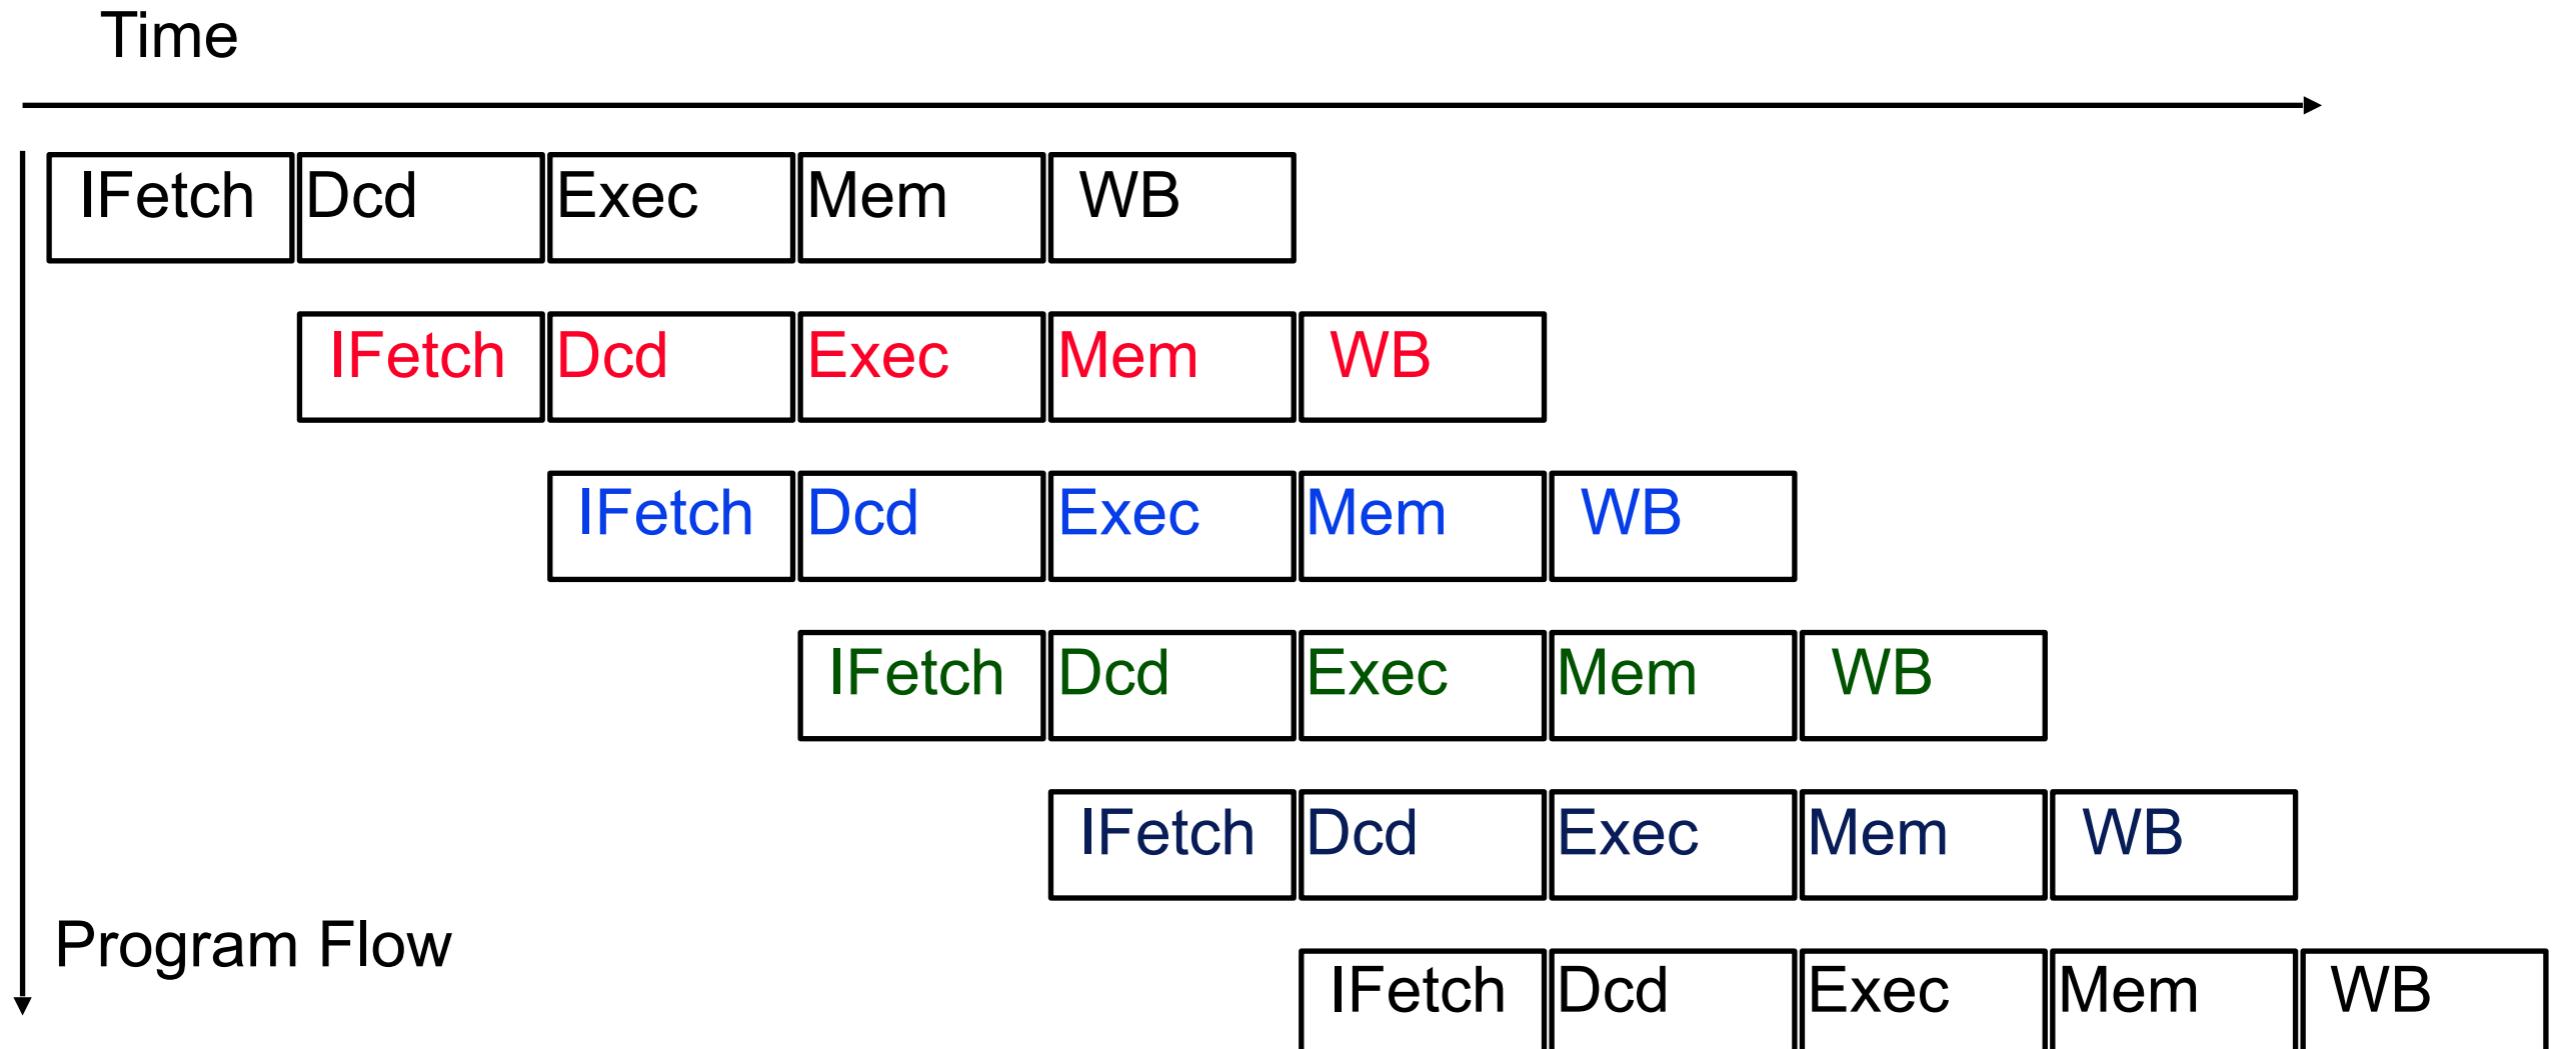
- Stages:
  - F: Instruction fetch
  - D/RF: Decode, register fetch
  - EX: Execute
  - MEM: Memory access
  - WB: Write-back
- Typical of load-store architecture

# The Five Stages of Load

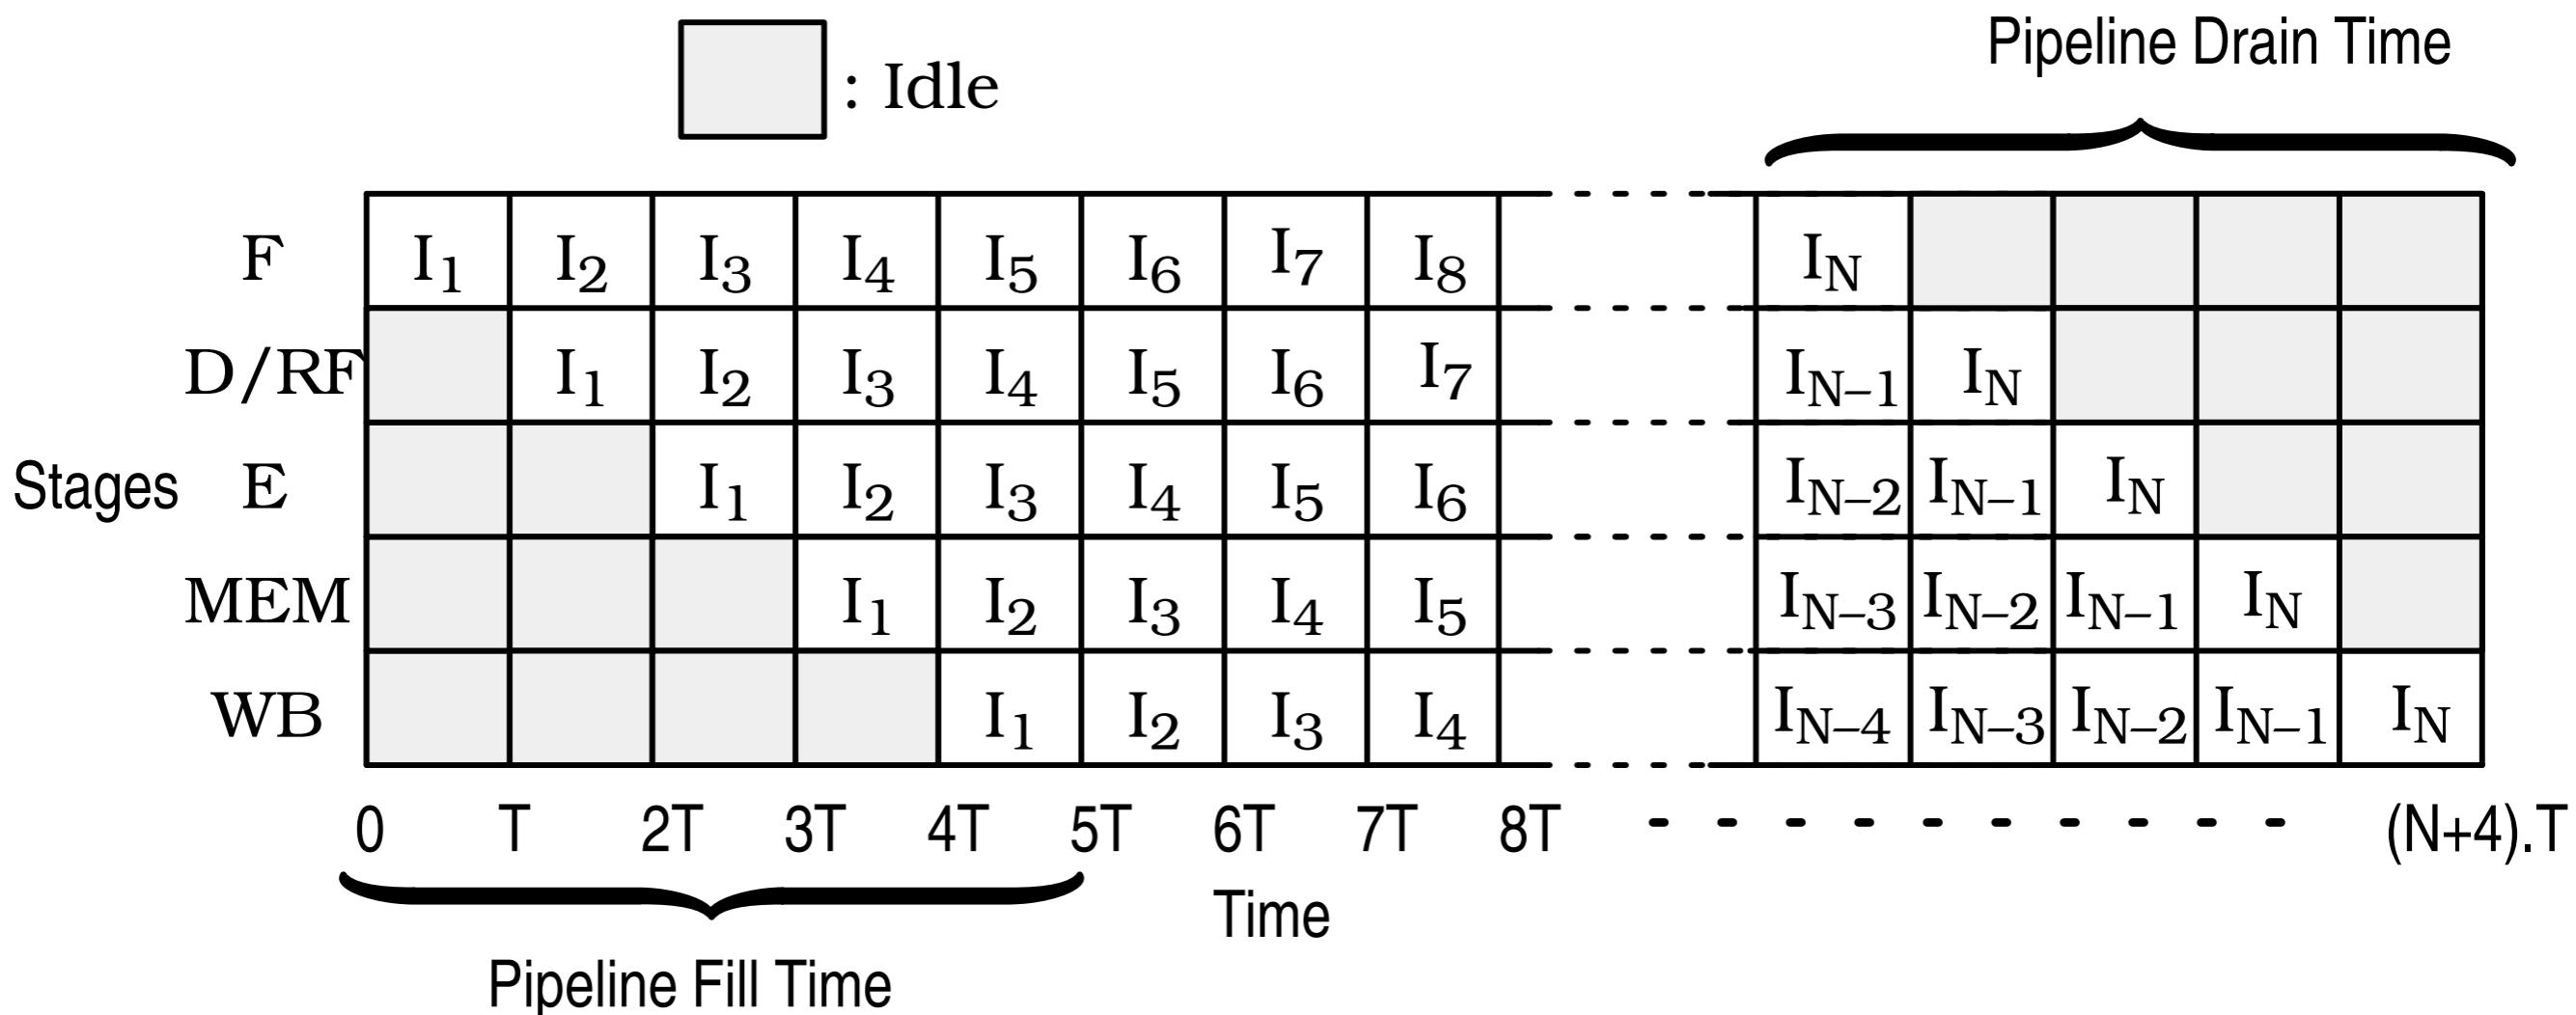


- ° **Ifetch: Instruction Fetch**
  - Fetch the instruction from the Instruction Memory
- ° **Reg/Dec: Registers Fetch and Instruction Decode**
- ° **Exec: Calculate the memory address**
- ° **Mem: Read the data from the Data Memory**
- ° **Wr: Write the data back to the register file**

# Conventional Pipelined Execution Representation

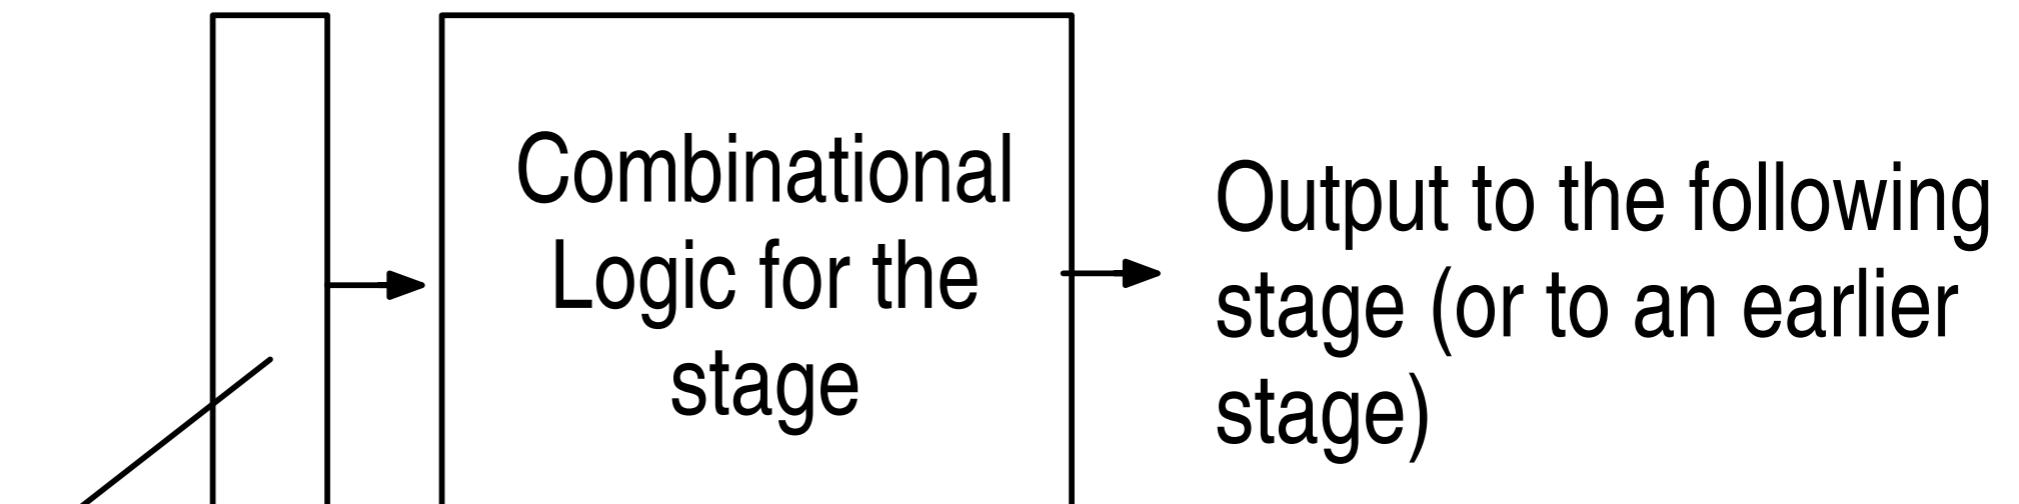


# $N$ instructions to be processed



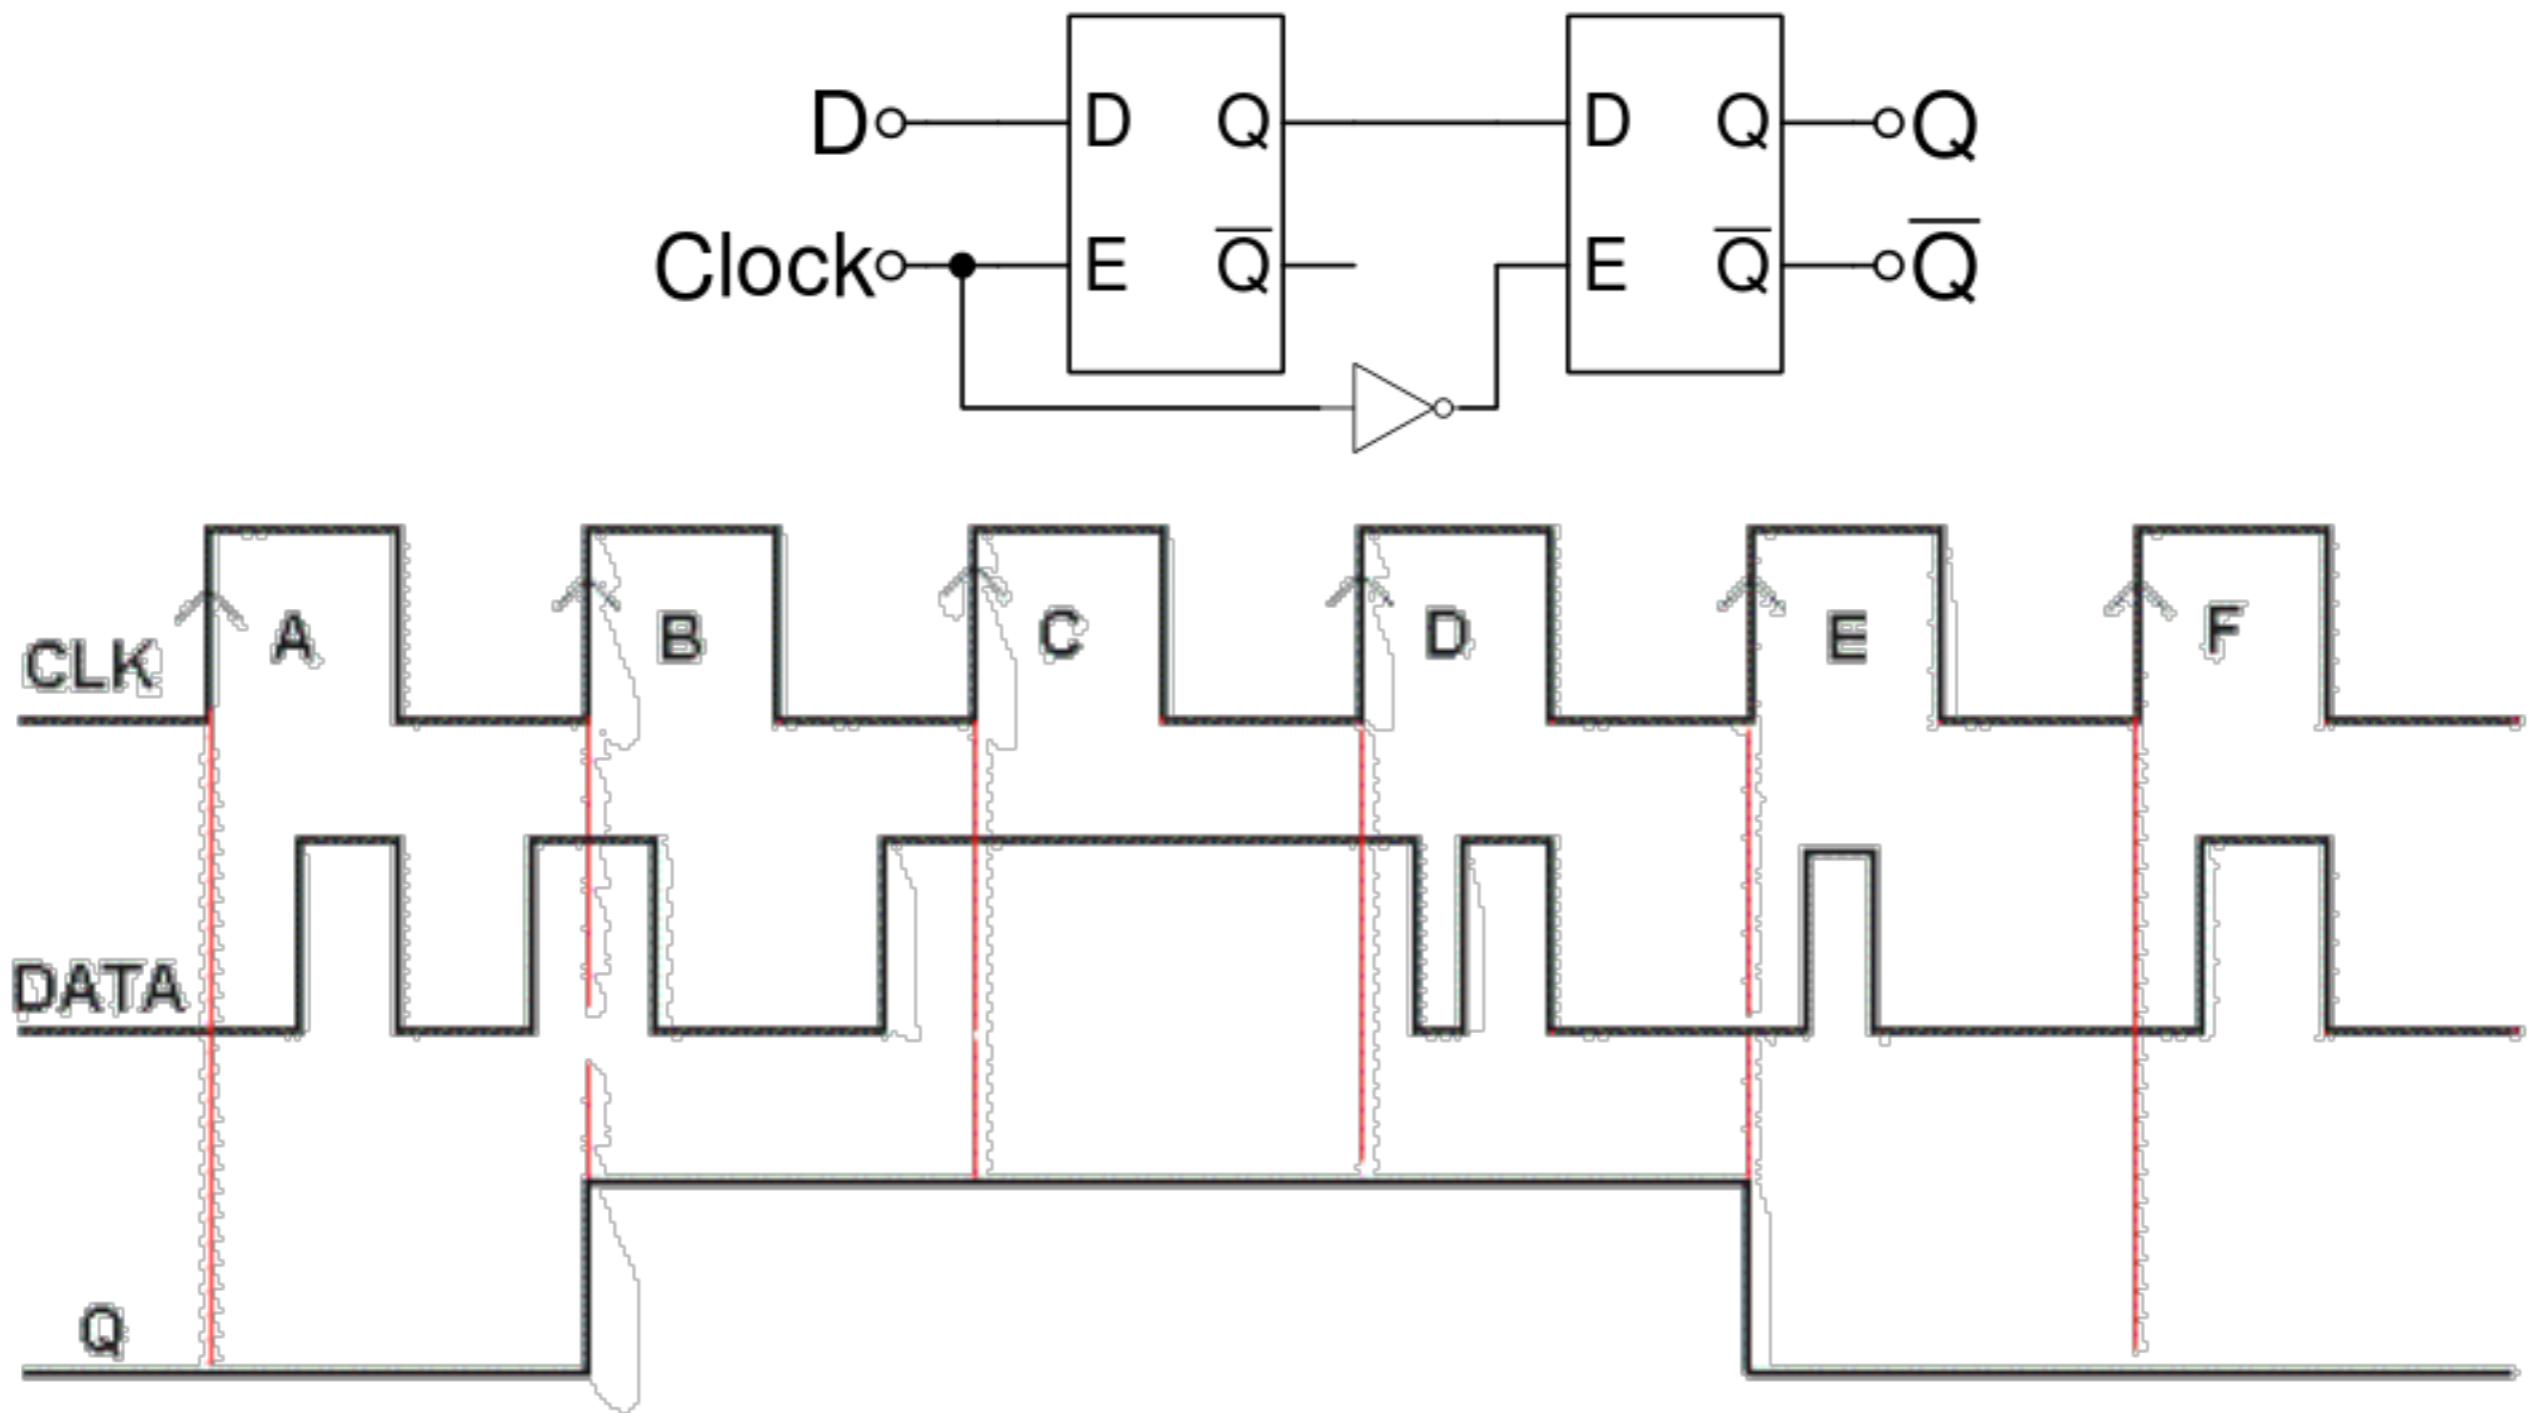
The Execution of a Sequence of  $N$  Instructions on APEX

# Single stage of a pipeline

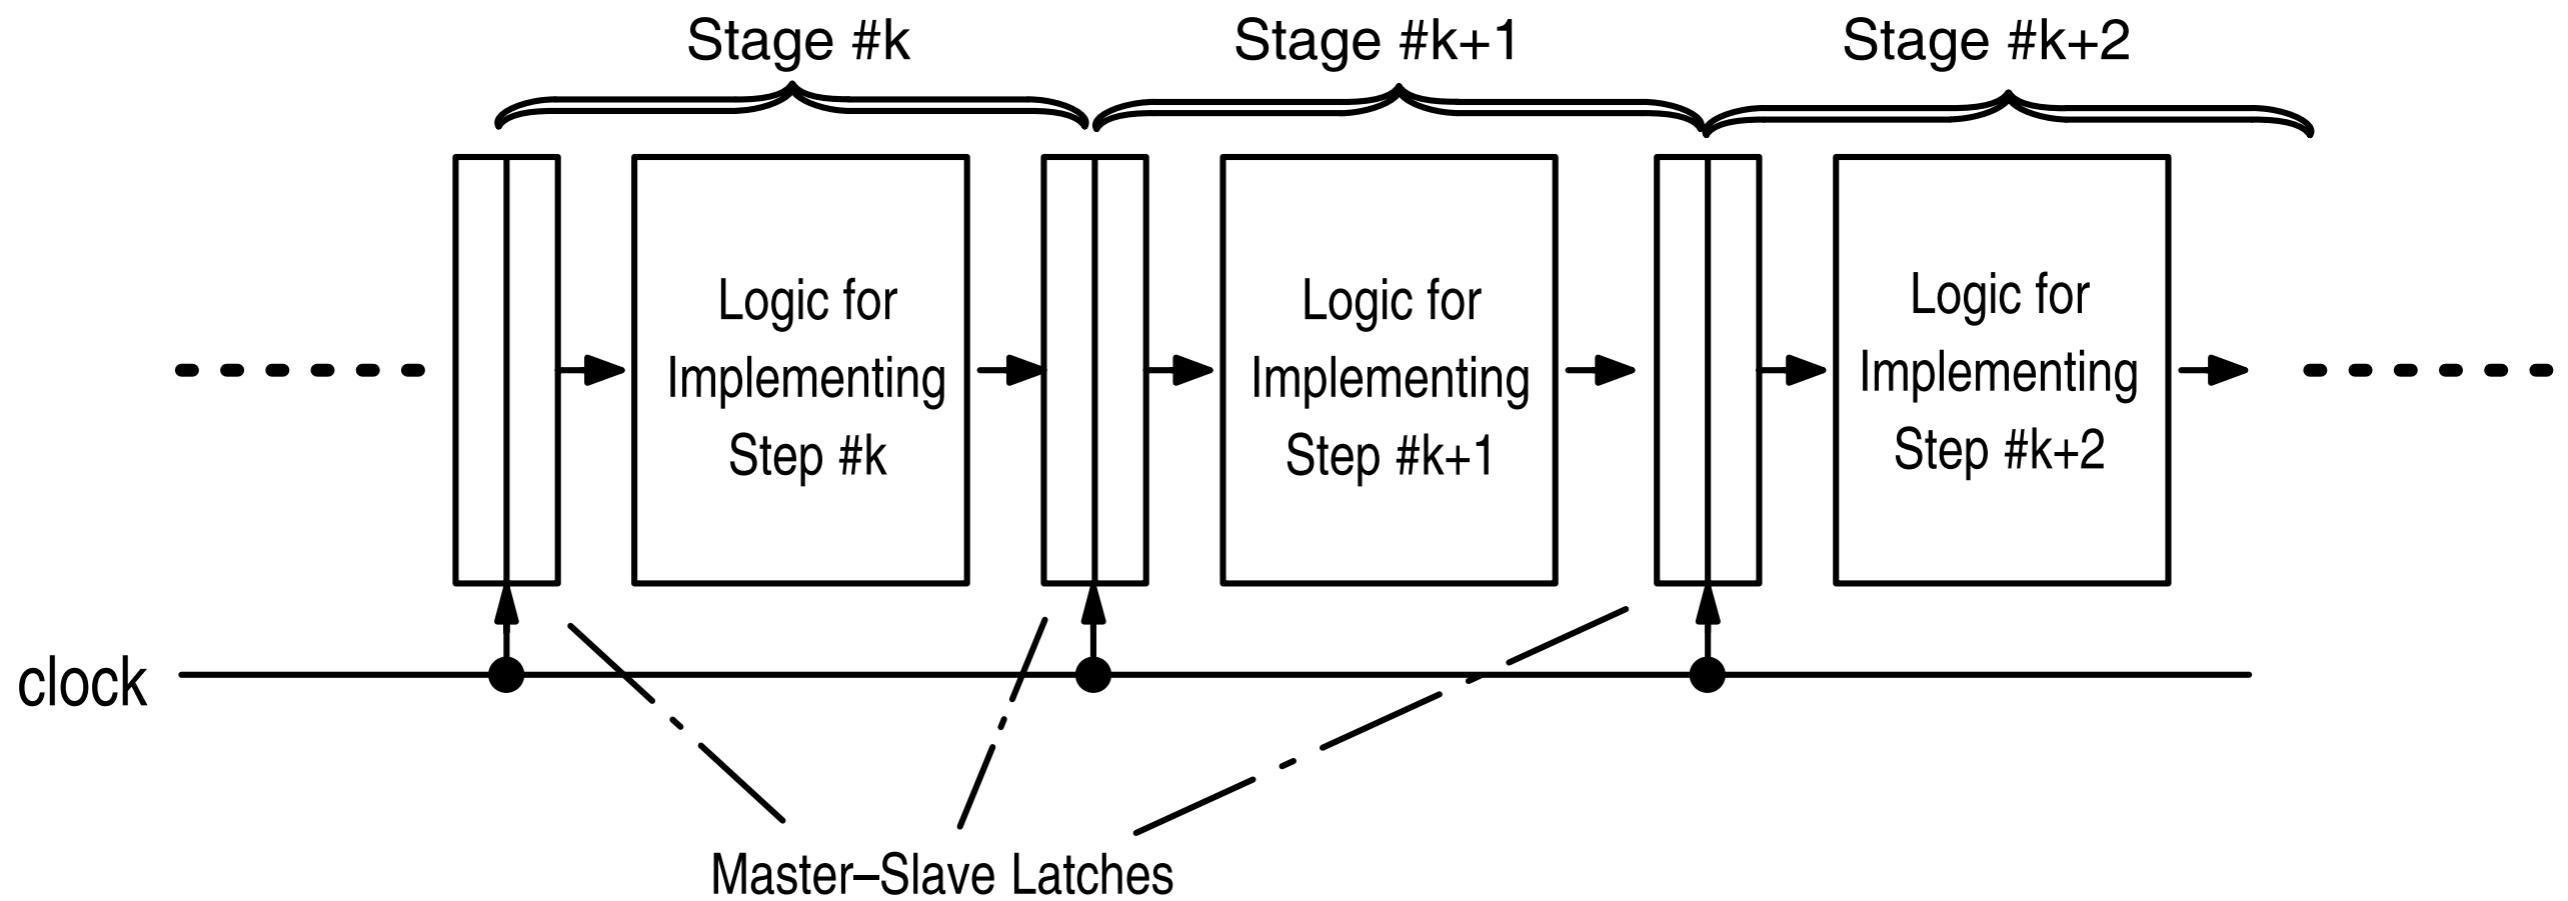


One or more input latches or registers

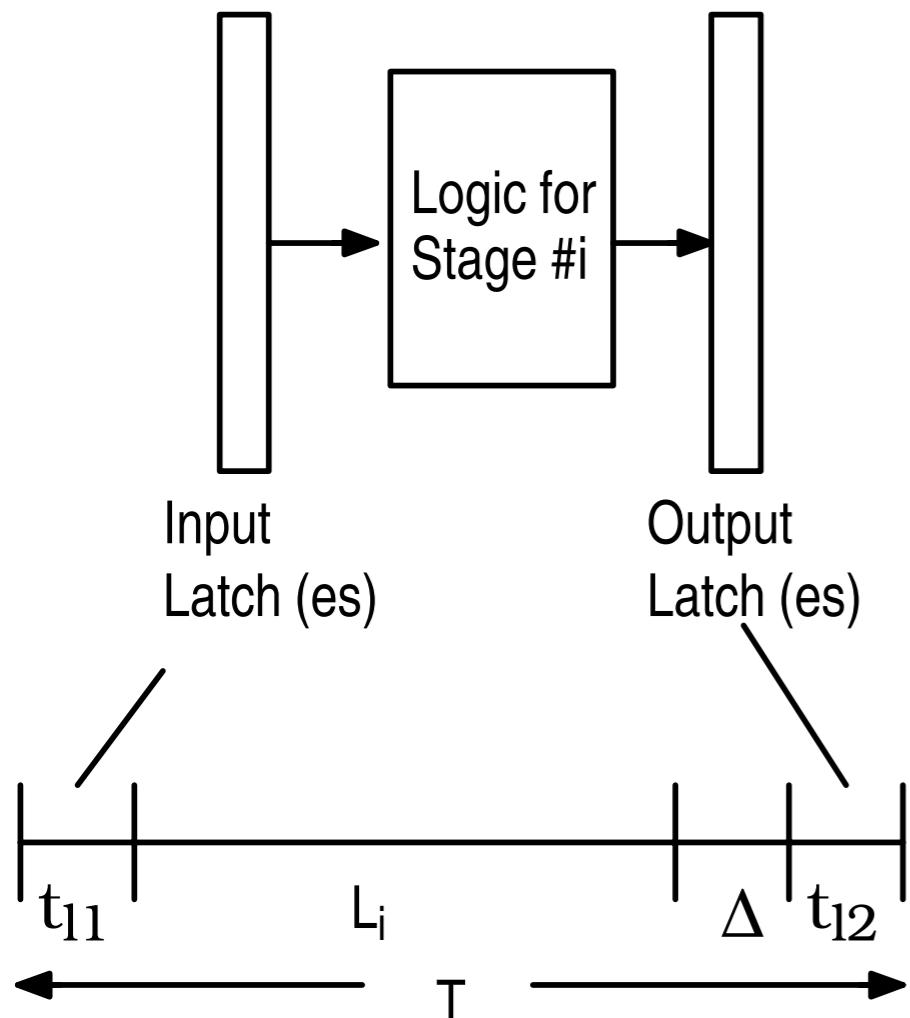
# Edge Triggered Flip-flop



# Pipeline as a whole



Synchronous pipeline,  
common clock



$t_{l1}$  = input latching delay (delay of slave latch)

$t_{l2}$  = output latching delay (delay of master latch)

$t_{\text{latch}} = t_{l1} + t_{l2}$  = overall latching delay: roughly the same for all stages

$L_i$  = processing delay for the logic of stage #i

$L_{\max}$  = processing delay for the slowest logic associated with any stage =  $\max \{L_i\}$

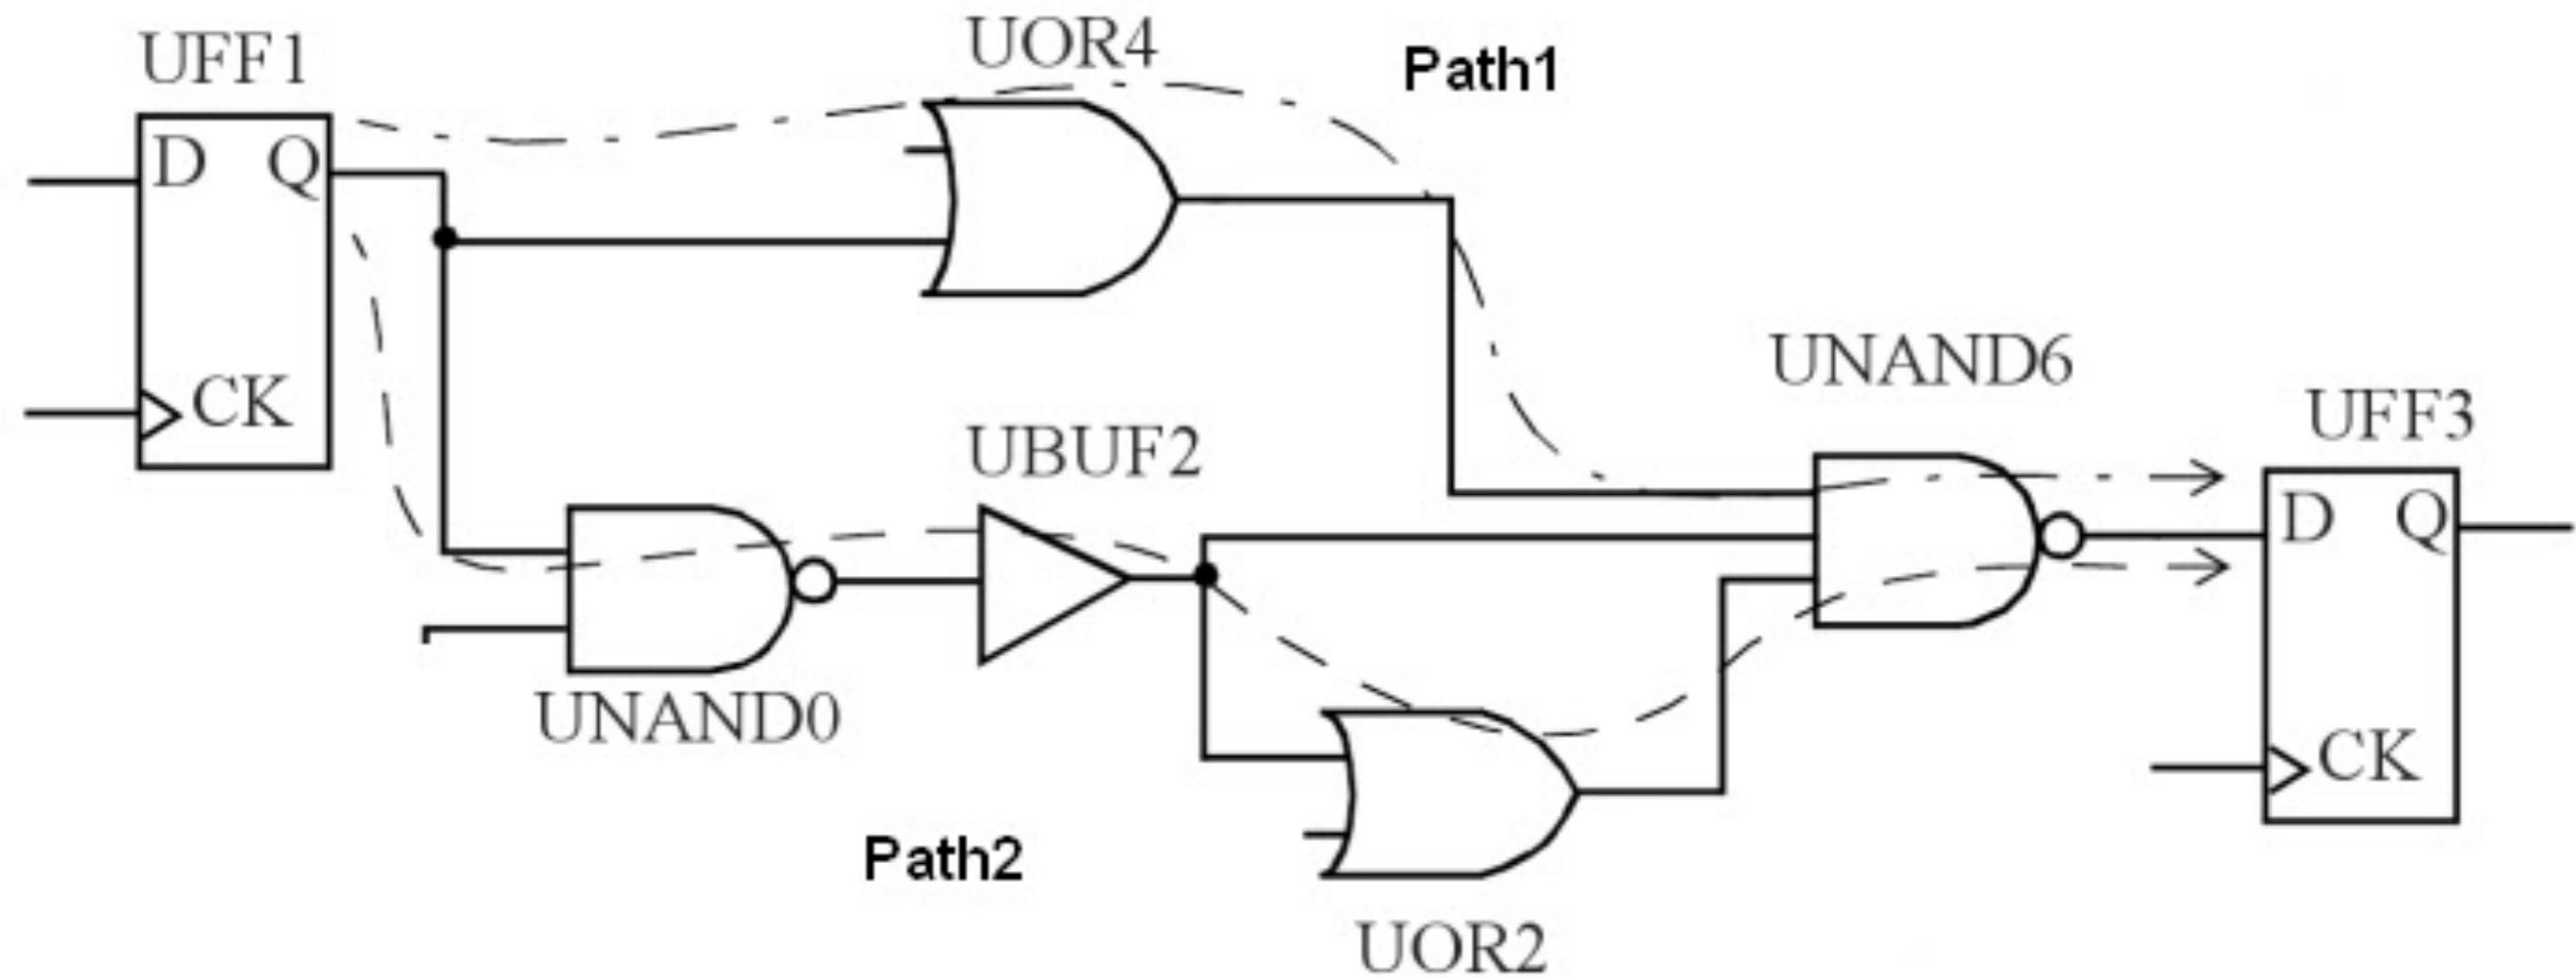
$T$  = pipeline cycle time

$$= t_l + L_{\max} + t_{\text{skew}}$$

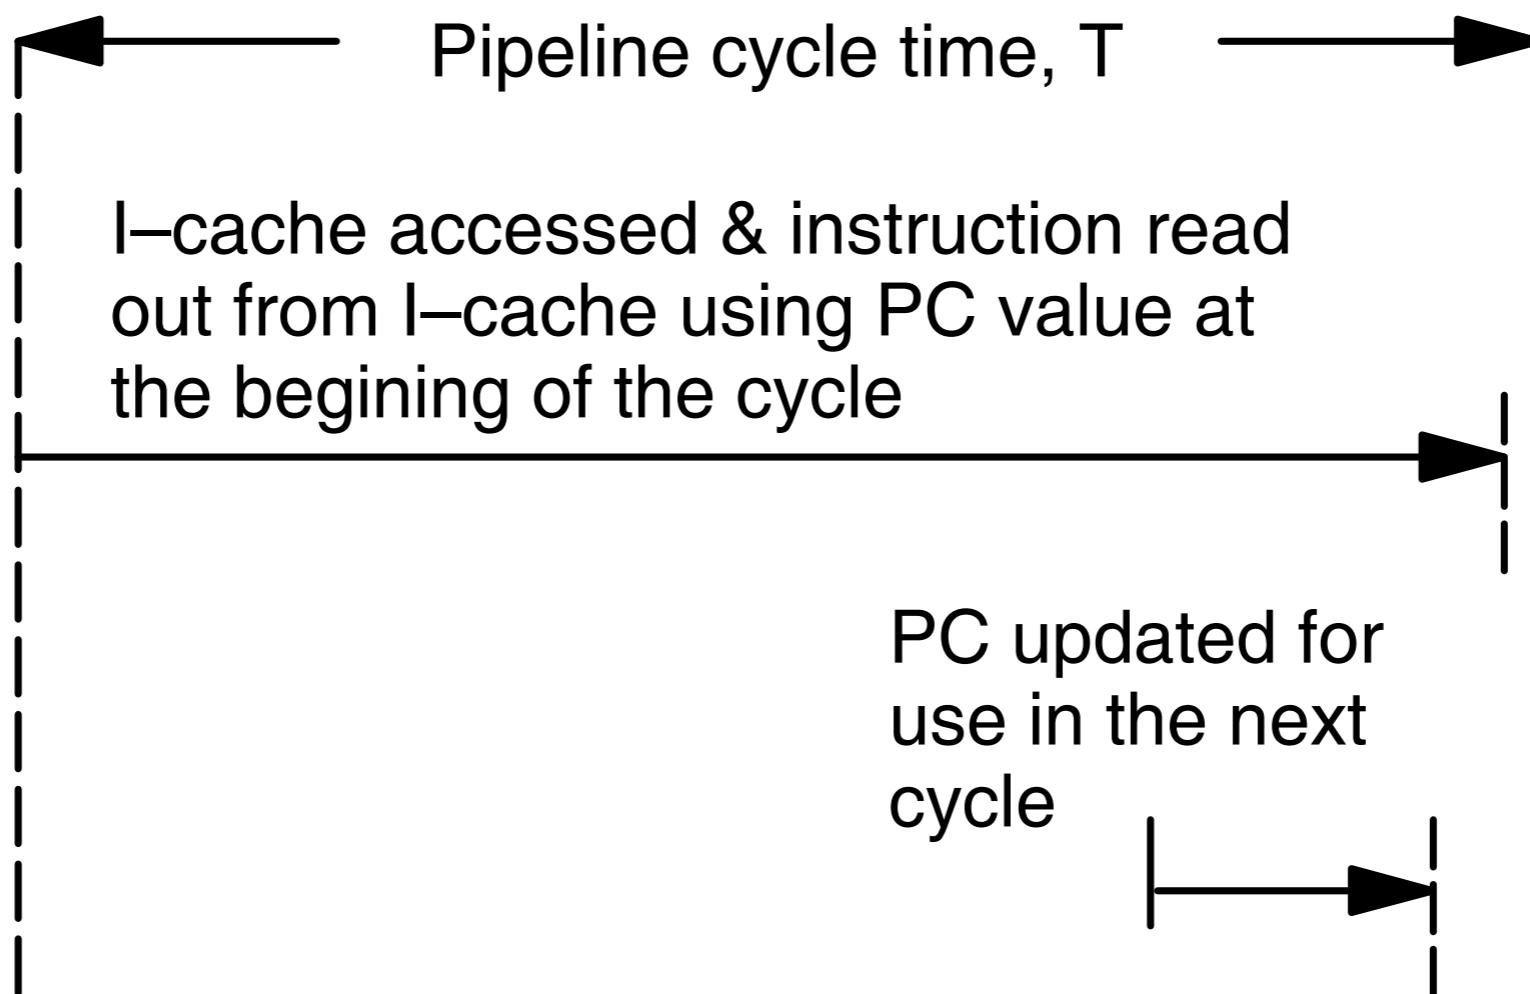
$t_{\text{skew}}$  = allowance for clock skew

$$\Delta = L_{\max} - L_i + t_{\text{skew}}$$

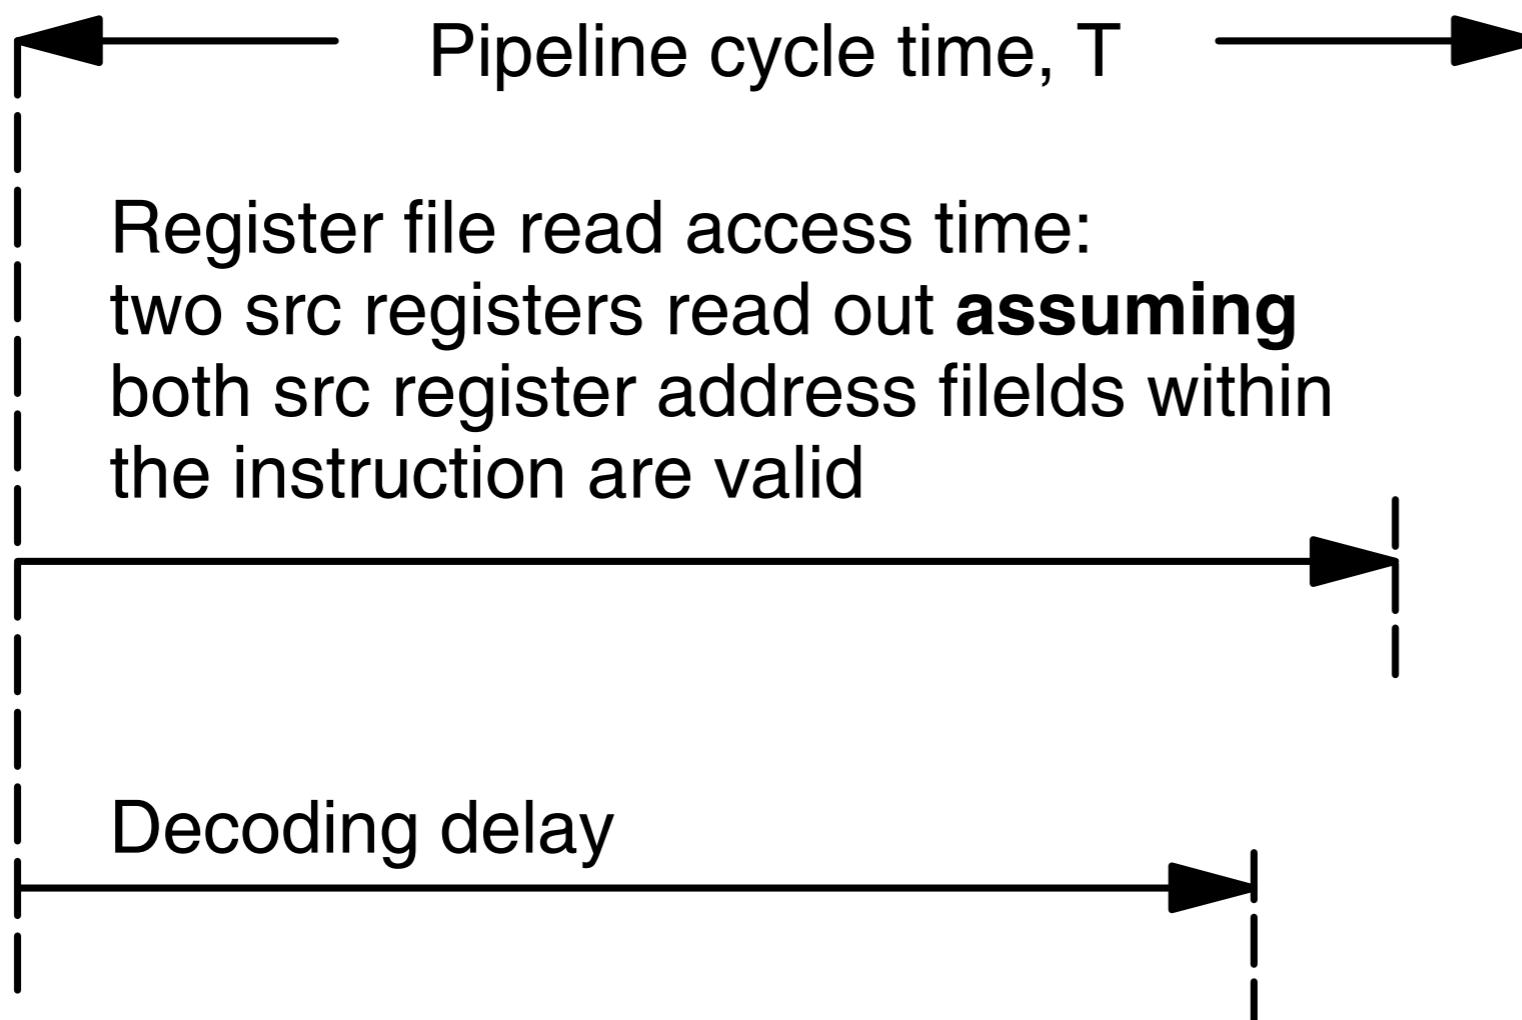
# Circuit Path



# Reading I-cache



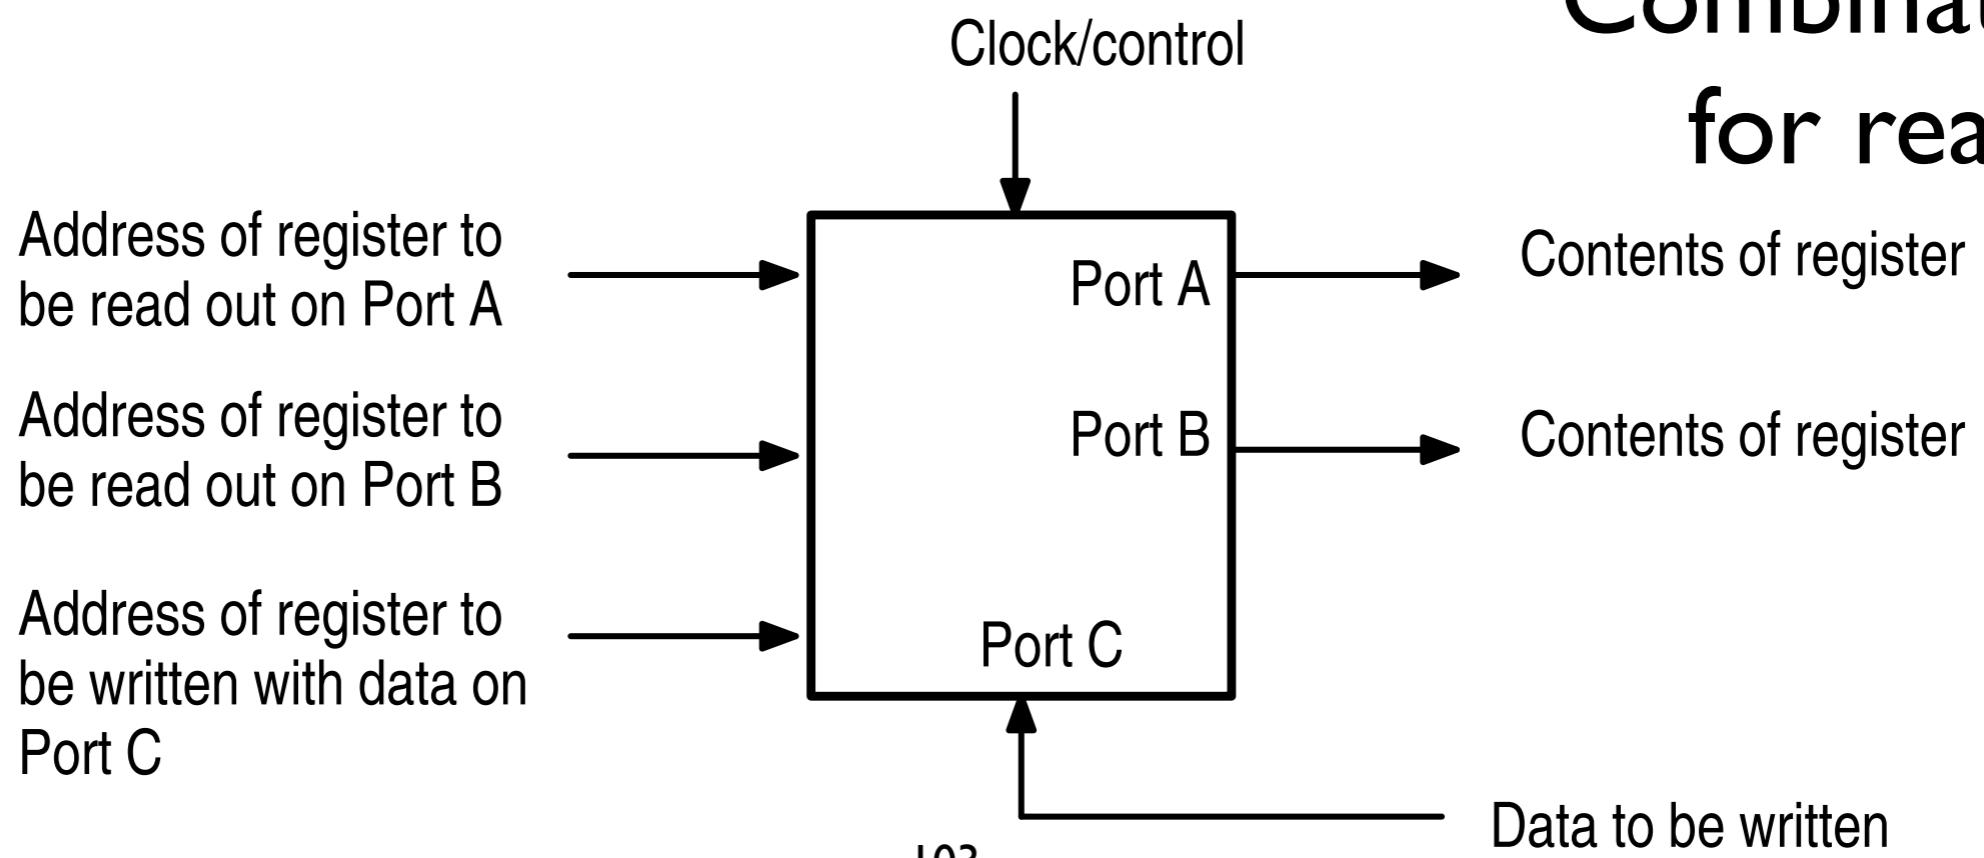
# Decoding and register file read



At this point, decoded info is available to discard data read from register(s) whose address field(s) was/were incorrectly assumed to be valid

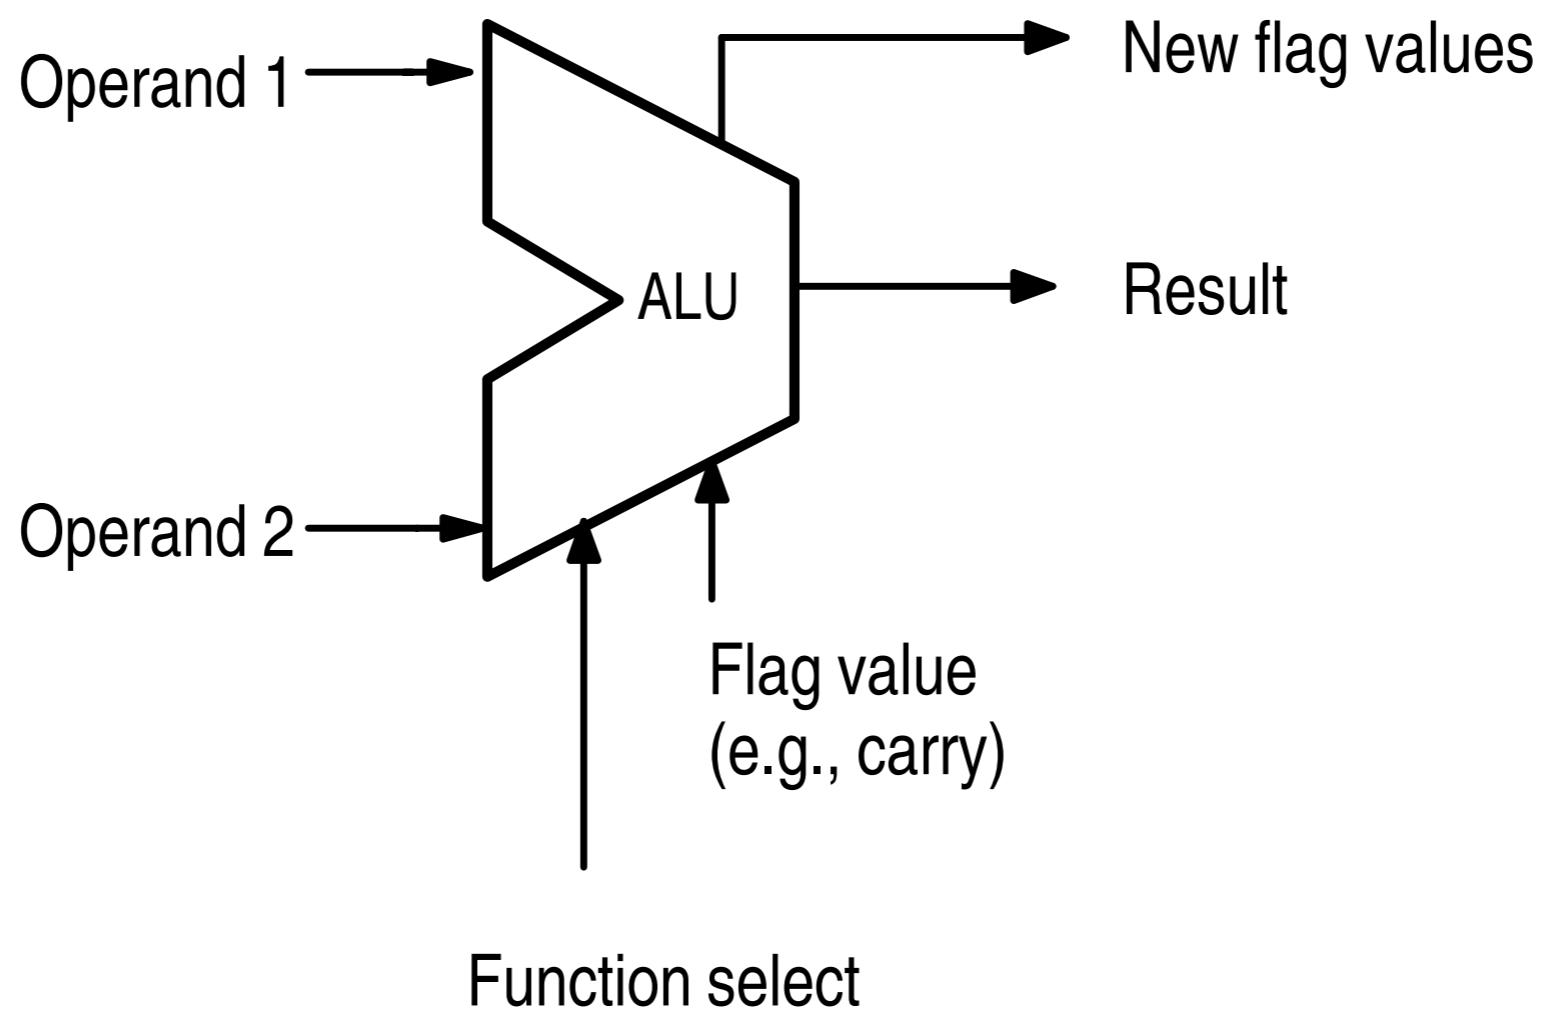
# Register file

- Small RAM
  - Multiple ports: 2 read, 1 write simultaneously
  - Each location is a register



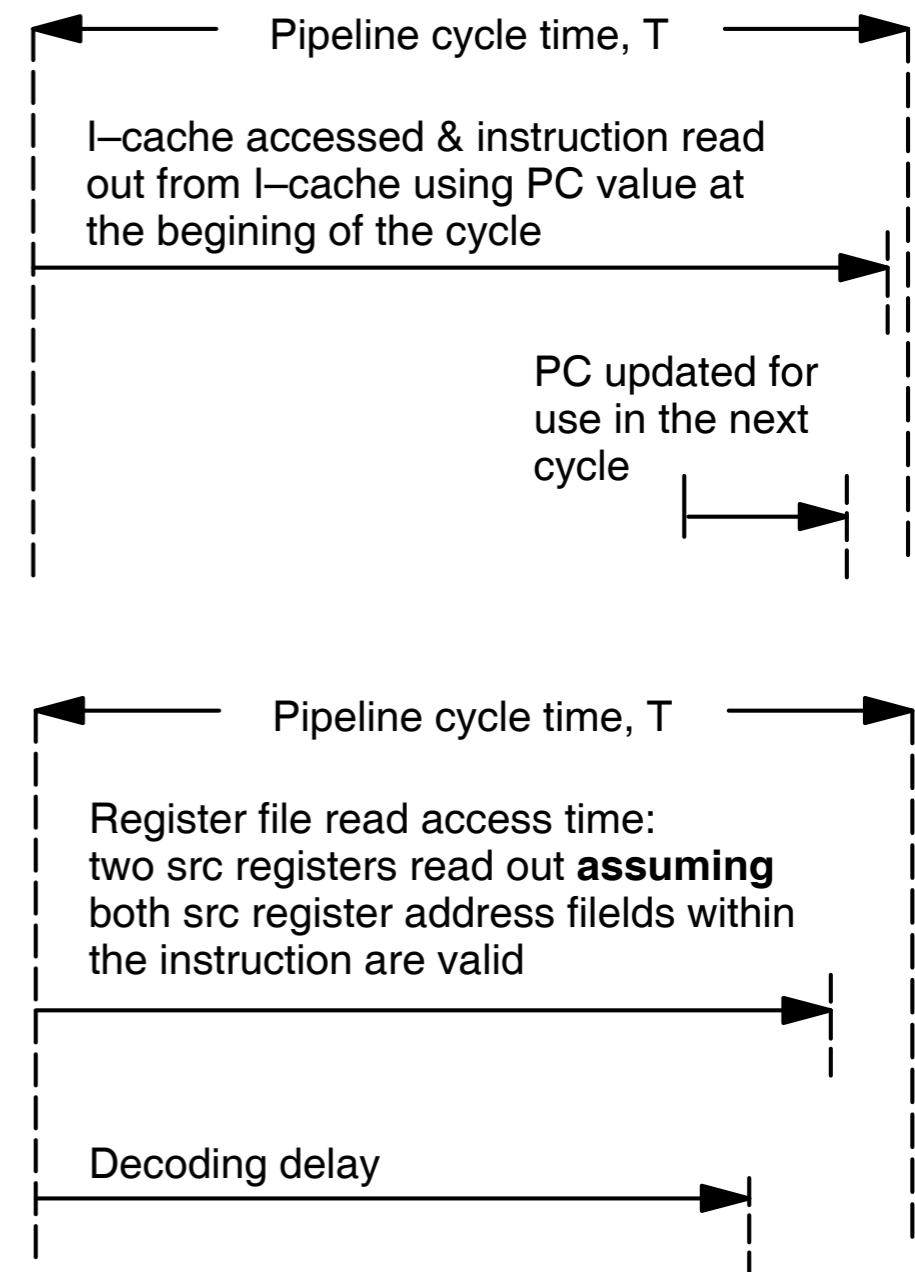
# ALU

- Combinational circuit for arithmetic and logic ops



# Logic utilization

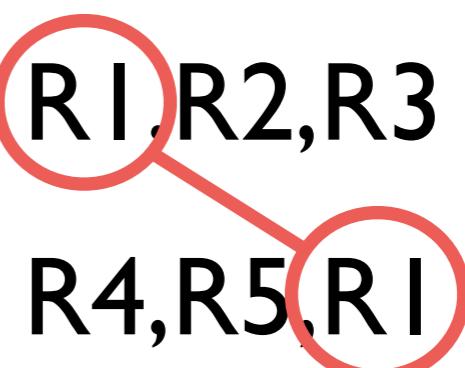
- Logic delays of stages are not uniform
- Fastest logic waits for slowest logic to catch up (utilization loss)
- Solution: Normalize logic delays across stages



# Dependencies and Parallelism

- ADD R1,R2,R3
- ADD R4,R5,R6
- ADD R1,R2,R3
- ADD R4,R5,R1

# Dependencies and Parallelism

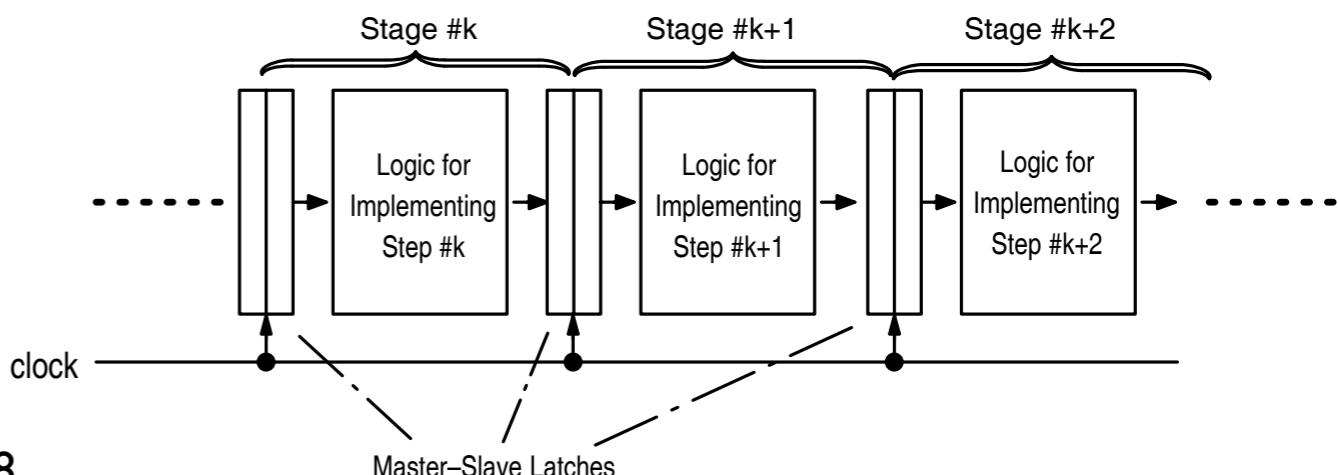
- ADD R1,R2,R3
  - ADD R4,R5,R6
- ADD R1,R2,R3
  - ADD R4,R5,R1
- 

# More on pipelining

- Pipelines exploit ILP -- Instruction-Level Parallelism
- Pipeline latency = number of cycles needed to process instruction



- Pipeline clock = common clock to drive interstage latches
- Clock rate =  $I/T$



# Clocking issues

- Clock period must be long enough for:
  - Critical paths
  - Latch overhead
  - Clock skew

# Pipeline efficiency

- $T_p$  = Total cycles when any stages are committed
- $K * T_p$  = Total resources committed over this duration (stage cycles)
- $K * N$  = Actual usage of stages, each stage used once by  $N$  instructions

# Pipeline efficiency

- Efficiency = Actual resources used/committed
  - $= (K * N) / (K * T_p) = N / (K + N - I)$
  - < 100% because some stages are unused during filling and flushing
  - < 100% because dependencies cause stalls

# Ideal CPI with pipelining

- $T_{exec} = T_p = K * T + (N - 1) * T$
- $= N * CPI * T$
- Best possible CPI approaches 1 while  $N \gg K$  and  $K > 1$
- Throughput =  $N / T_{exec}$
- IPC =  $1/CPI$

# Stalls

- If an instruction spends more than one cycle in a stage, prior stages *stall*.
  - Delays in fetching instruction, I-cache miss
  - Delays in accessing memory in MEM, D-cache miss
  - Waiting on dependencies
- Stalls cause “bubbles”

# Real CPI

- Higher than 1 due to:

- Data dependencies (interlocking)

```
LOAD  R1, R3, #50    /* R1 <- Mem [R3 + 50] */  
ADD   R2, R1, R6    /* R2 <- R1 + R6 */
```

- Branches

- Slow memory

- Resource contention

# Real CPI for pipelines

- Design goal: CPI as close to 1 as possible
- $f = \text{average number of cycles per instruction lost to bubbles}$ 
  - $\text{CPI} = (K - 1 + N * (1+f)) / N = \sim 1+f$
- Hardware and software solutions used to reduce CPI.

# **The APEX processor ISA**

# APEX ISA

load: LOAD <dest> <src1> <literal>  
store: STORE <src1> <src2> <literal>  
op: op <dest> <src1> <src2>  
opL: op <dest> <src1> <literal>

# APEX architecture

- Load-store architecture
- All instructions 32-bits
- 32 general-purpose registers, 0-31
- Registers and memory 32-bits
- Source operands: *rsrc*
- Destination operands: *rdest*
- *Mem[X]* is memory address at addr X.

# APEX pipeline



The APEX Pipeline

## Processing Step

Fetch Instruction

Decode and fetch register operands

Execute register-to-register operation

Perform memory operation, if any

Write results back

## Associated Stage Name

Fetch stage

Decode, register fetch

Execute

Memory

Writeback

## Abbreviation

F

D/RF

EX

MEM

WB

# Instruction flow

- Instruction, operands, and addresses flow down pipeline as necessary.
- Program counter also: PC-relative addressing, return address for subroutines and interrupts

# Example instructions

ADD rdest rsrc1 rsrc2 :       $rdest \leftarrow rsrc1 + rsrc2$

ADDL rdest rsrc1 literal :  $rdest \leftarrow rsrc1 + \text{literal}$

LDI rdest rsrc1 literal :  $rdest \leftarrow \text{Mem}[rsrc1 + \text{literal}]$

LDIR rdest rsrc1 rsrc2 :       $rdest \leftarrow \text{Mem}[rsrc1 + rsrc2]$

STI rsrc1 rsrc2 literal :  $\text{Mem}[rsrc2 + \text{literal}] \leftarrow rsrc1$

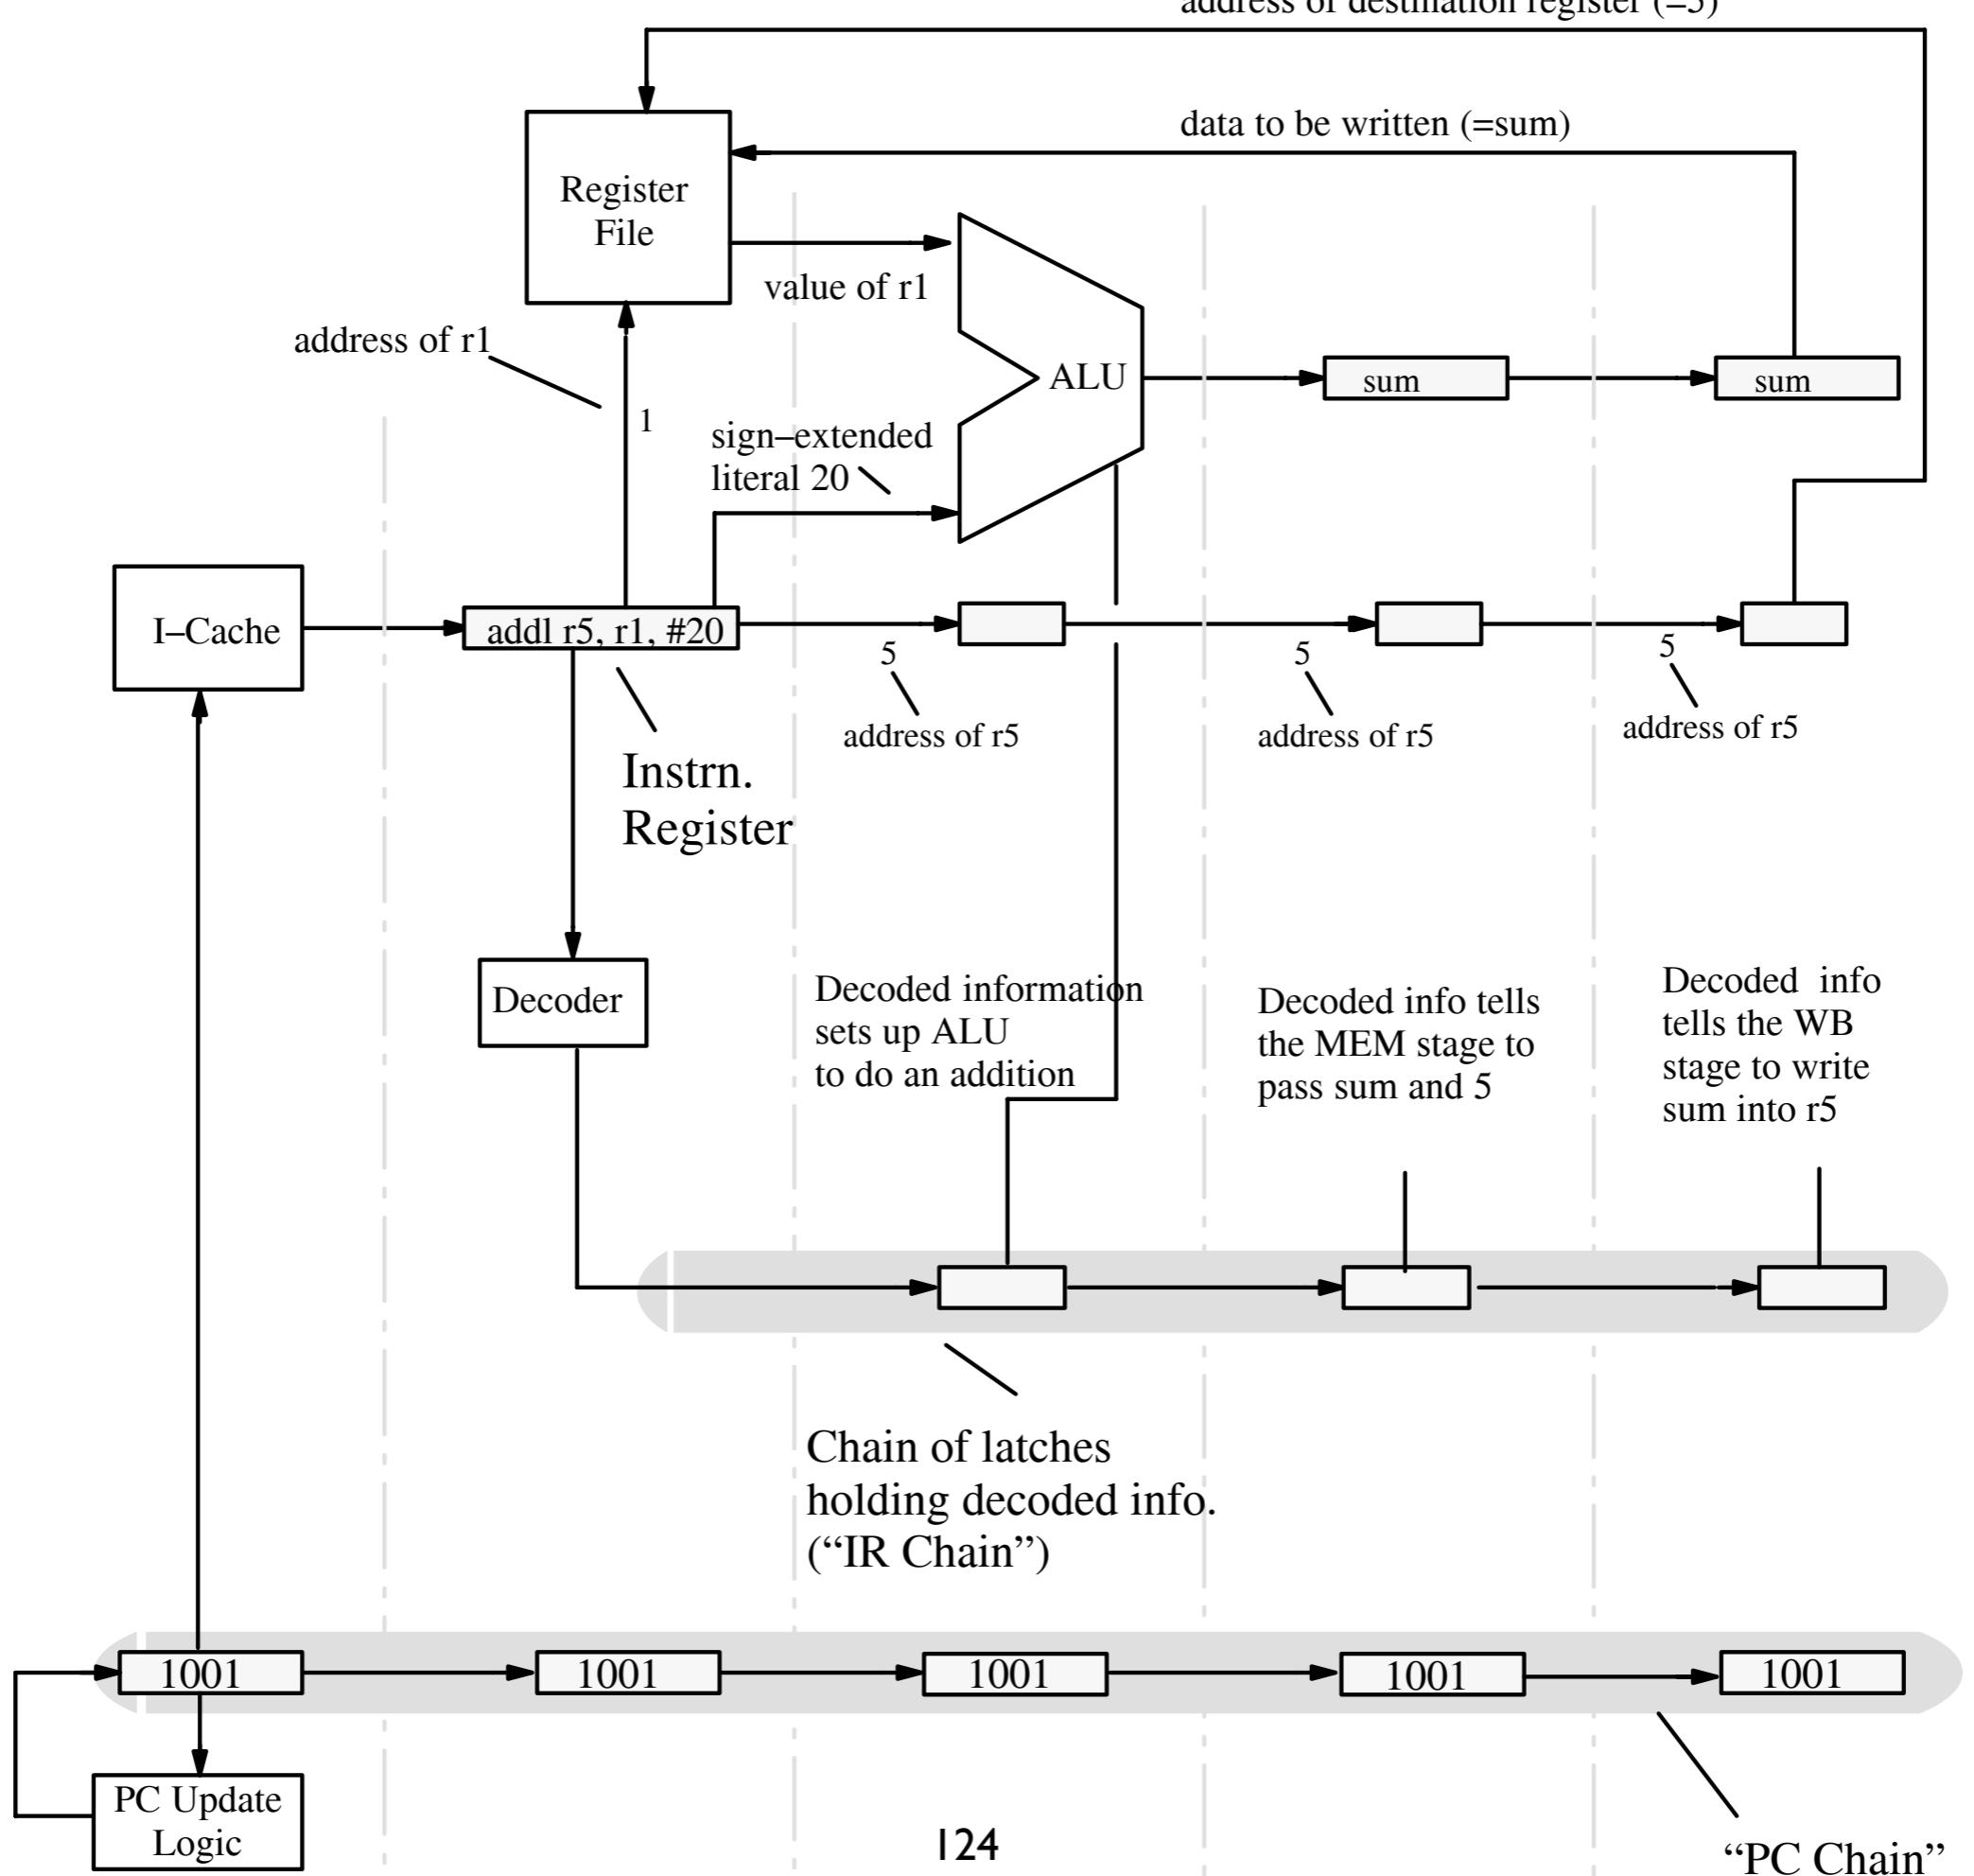
INSTRUCTION \ Stage	F stage	D/RF stage	EX stage	MEM stage	WB stage
Reg-to-reg: ADD rdest, rsrc1, rsrc2	Read instruction from memory	Decode instruction; Read rsrc1 and rsrc2 from the RF	Add contents of rsrc1 and rsrc2; Set condition code flags	Hold on to the result produced in the EX stage in the previous cycle;	Write result to register rdest
Load indexed–literal offset: LDI rdest, rsrc1, <literal>	Read instruction from memory	Decode instruction; Read rsrc1; sign extend <literal>	Add contents of rsrc1 and sign extended literal; Set condition code flags;	Read the contents of the memory location whose address was computed in the previous cycle in the EX stage	Write the data read out from memory to the register rdest
Load indexed–reg. offset: LDIR rdest, rsrc1, rsrc2	Read instruction from memory	Decode instruction; Read rsrc1 and rsrc2 from the RF	Add contents of rsrc1 and rsrc2; Set condition code flags	Read the contents of the memory location whose address was computed in the previous cycle in the EX stage	Write the data read out from memory to the register rdest
Store indexed–literal offset: STI rsrc1, rsrc2, <literal>	Read instruction from memory	Decode instruction; Read rsrc1 and rsrc2; sign extend <literal>	Add contents of rsrc2 and sign extended literal; Set condition code flags	Write the contents of rsrc1 (read out in the D/RF stage) to the memory location whose address was computed in the EX stage in the previous cycle	No action

Use of the stages of APEX by register-to-register, load and store instructions

INSTRUCTION \ Between Stages	F and D/RF	D/RF and EX	EX and MEM	MEM and WB
Reg-to-reg: ADD rdest, rsrc1, rsrc2	<ul style="list-style-type: none"> <li>– The instruction read from memory</li> <li>– The PC value</li> </ul>	<ul style="list-style-type: none"> <li>– Instrn. info;</li> <li>– The PC value</li> <li>– Contents of rsrc1 and rsrc2;</li> <li>– Address of rdest</li> </ul>	<ul style="list-style-type: none"> <li>– Instrn. info;</li> <li>– The PC value</li> <li>– Address of rdest;</li> <li>– Result of operation in the EX stage</li> </ul>	<ul style="list-style-type: none"> <li>– Instrn. info;</li> <li>– The PC value</li> <li>– Address of rdest;</li> <li>– Result of operation in the E stage</li> </ul>
Load indexed–literal offset: LDI rdest, rsrc1, <literal>	<ul style="list-style-type: none"> <li>– The instruction read from memory</li> <li>– The PC value</li> </ul>	<ul style="list-style-type: none"> <li>– Instrn. info;</li> <li>– The PC value</li> <li>– Contents of rsrc1;</li> <li>– Sign extended literal</li> <li>– Address of rdest</li> </ul>	<ul style="list-style-type: none"> <li>– Instrn. info;</li> <li>– The PC value</li> <li>– Address of rdest;</li> <li>– Result of operation in the EX stage</li> </ul>	<ul style="list-style-type: none"> <li>– Instrn. info;</li> <li>– The PC value</li> <li>– Address of rdest;</li> <li>(Data being loaded comes from memory)</li> </ul>
Load indexed–reg. offset: LDIR rdest, rsrc1, rsrc2	<ul style="list-style-type: none"> <li>– The instruction read from memory</li> <li>– The PC value</li> </ul>	<ul style="list-style-type: none"> <li>– Instrn. info;</li> <li>– The PC value</li> <li>– Contents of rsrc1 and rsrc2;</li> <li>– Address of rdest</li> </ul>	<ul style="list-style-type: none"> <li>– Instrn. info;</li> <li>– The PC value</li> <li>– Address of rdest;</li> <li>– Result of operation in the EX stage</li> </ul>	<ul style="list-style-type: none"> <li>– Instrn. info;</li> <li>– The PC value</li> <li>– Address of rdest;</li> <li>(Data being loaded comes from memory)</li> </ul>
Store indexed–literal offset: STI rsrc1, rsrc2, <literal>	<ul style="list-style-type: none"> <li>– The instruction read from memory</li> <li>– The PC value</li> </ul>	<ul style="list-style-type: none"> <li>– Instrn. info;</li> <li>– The PC value</li> <li>– Contents of rsrc1 and rsrc2;</li> <li>– Sign extended literal</li> </ul>	<ul style="list-style-type: none"> <li>– Instrn. info;</li> <li>– The PC value</li> <li>– Contents of rsrc1;</li> <li>– Result of operation in the EX stage</li> </ul>	None

\* “Instrn. Info” indicates information about the instruction, such as the instruction itself or its decoded form – see text

Information flow among the stages of APEX for register-to-register, load and store instructions



# Features of APEX Pipeline

- Not all instructions use all stages -- data forwarded
  - Reg-to-reg: No processing in MEM
  - Store: WB not used.
- Registers read only in D/RF
  - Minimize number of ports
- Registers written only in WB
  - Even if result is available early
- EX state implements ALU and effective address calculation for LOAD and STORE

# Dependencies in Pipelines, long pipelines, multifunction pipelines

# Functional Pipelines

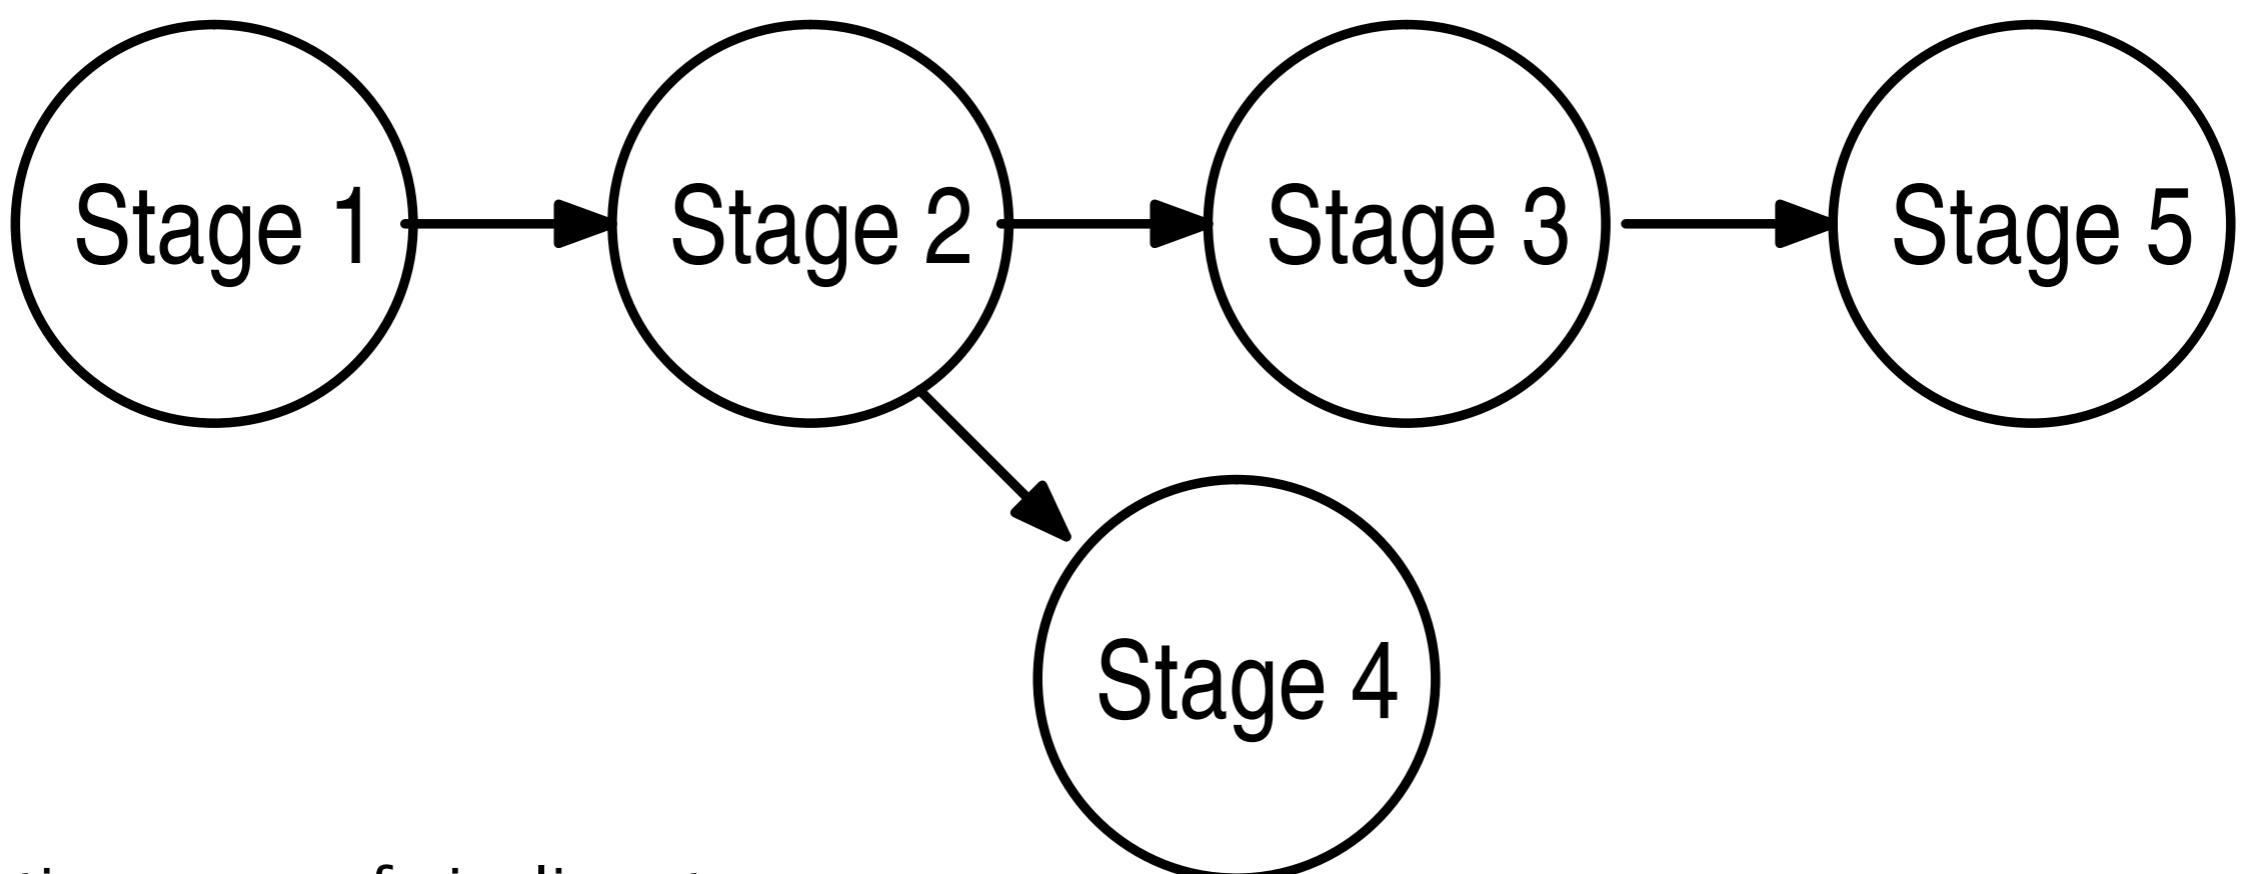
- Why pipeline?
- Pipelining to divide work among functional units
- Complex operations can be pipelined to increase throughput.

# Optimizing number of pipeline stages

- Calculating Speedup ( $S$ )
- $S = N * K * T / (K * T + (N - I) * T)$
- $S = T_{np} / T_p$
- $S_{max} = K$  ( $N \gg K, K > I$ )

# Precedence Graphs

- nodes = stages
- directed arcs = input-output order



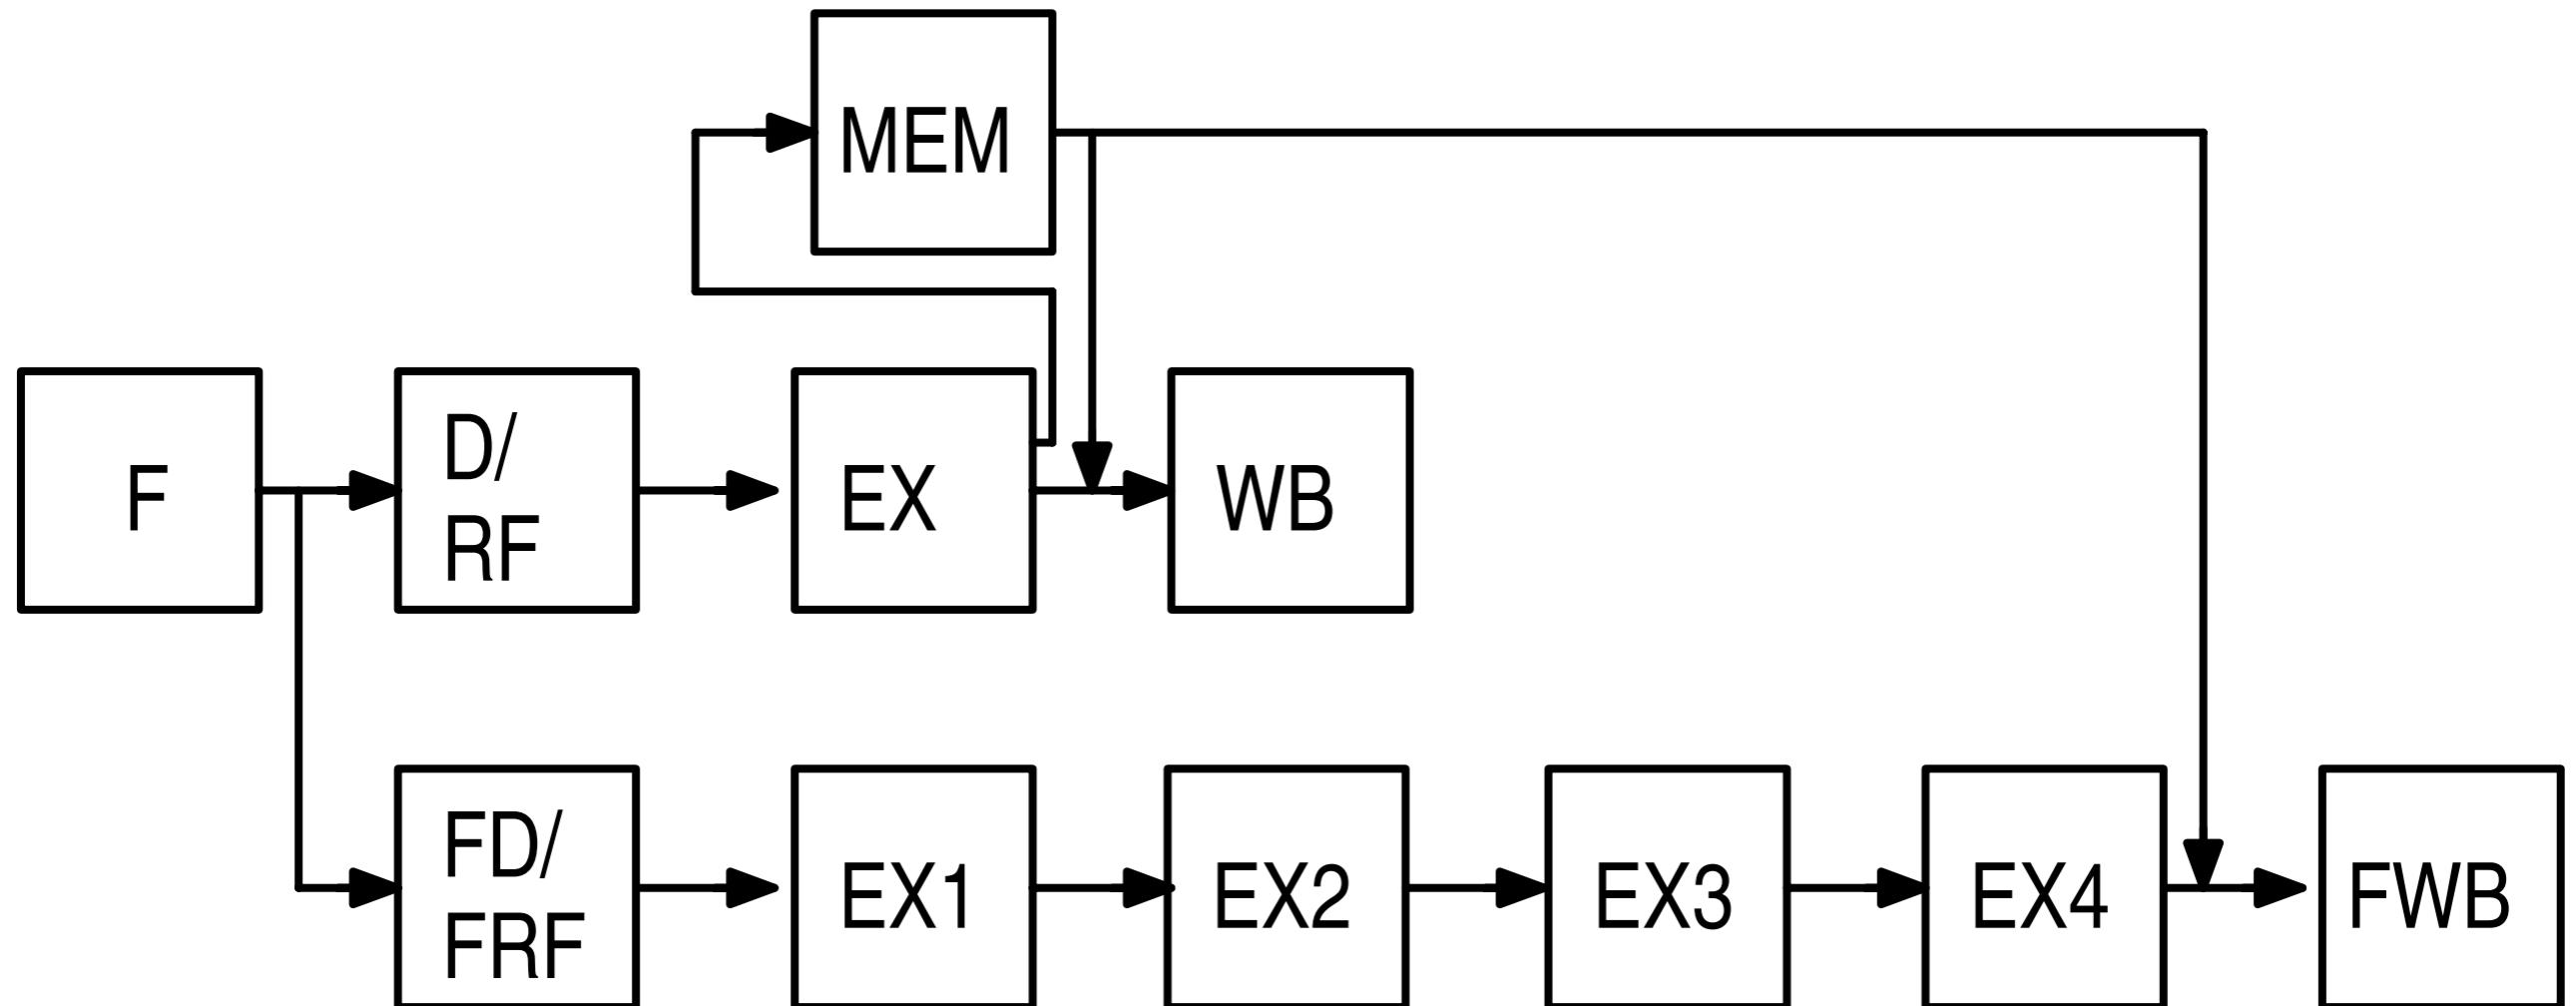
Depicting usage of pipeline stages

# Reservation tables

- One row per stage
- One column per use by instruction (time slot)
- Total columns = pipeline latency in cycles
- X marks usage of stage S in slot m.

		column number	1	2	3	4	5
			F	D/RF	EX	MEM	WB
Stages	0	X					
	1		X				
	2			X			
	3				X		
	4					X	
← Pipeline Latency →							

# More realistic pipeline



# Features of realistic pipeline

- Separate registers for floating point and integer data.
- Separate execution units for
  - Integer (IU)
  - Floating point (FPU)
  - Load/store unit (LSU)
- Floating point execution units more deeply pipelined.
- FPU and MEM may be further pipelined
- Potential resource conflict in using WB and FWB to write RF and FRF.

# Pipeline Classifications

- Linear vs. nonlinear (feedback)
  - For nonlinear, some rows in reservation table can have two or more X marks
- Synchronous vs. asynchronous
  - Synchronous: Common clock
  - Asynchronous: Clockless, handshake signals

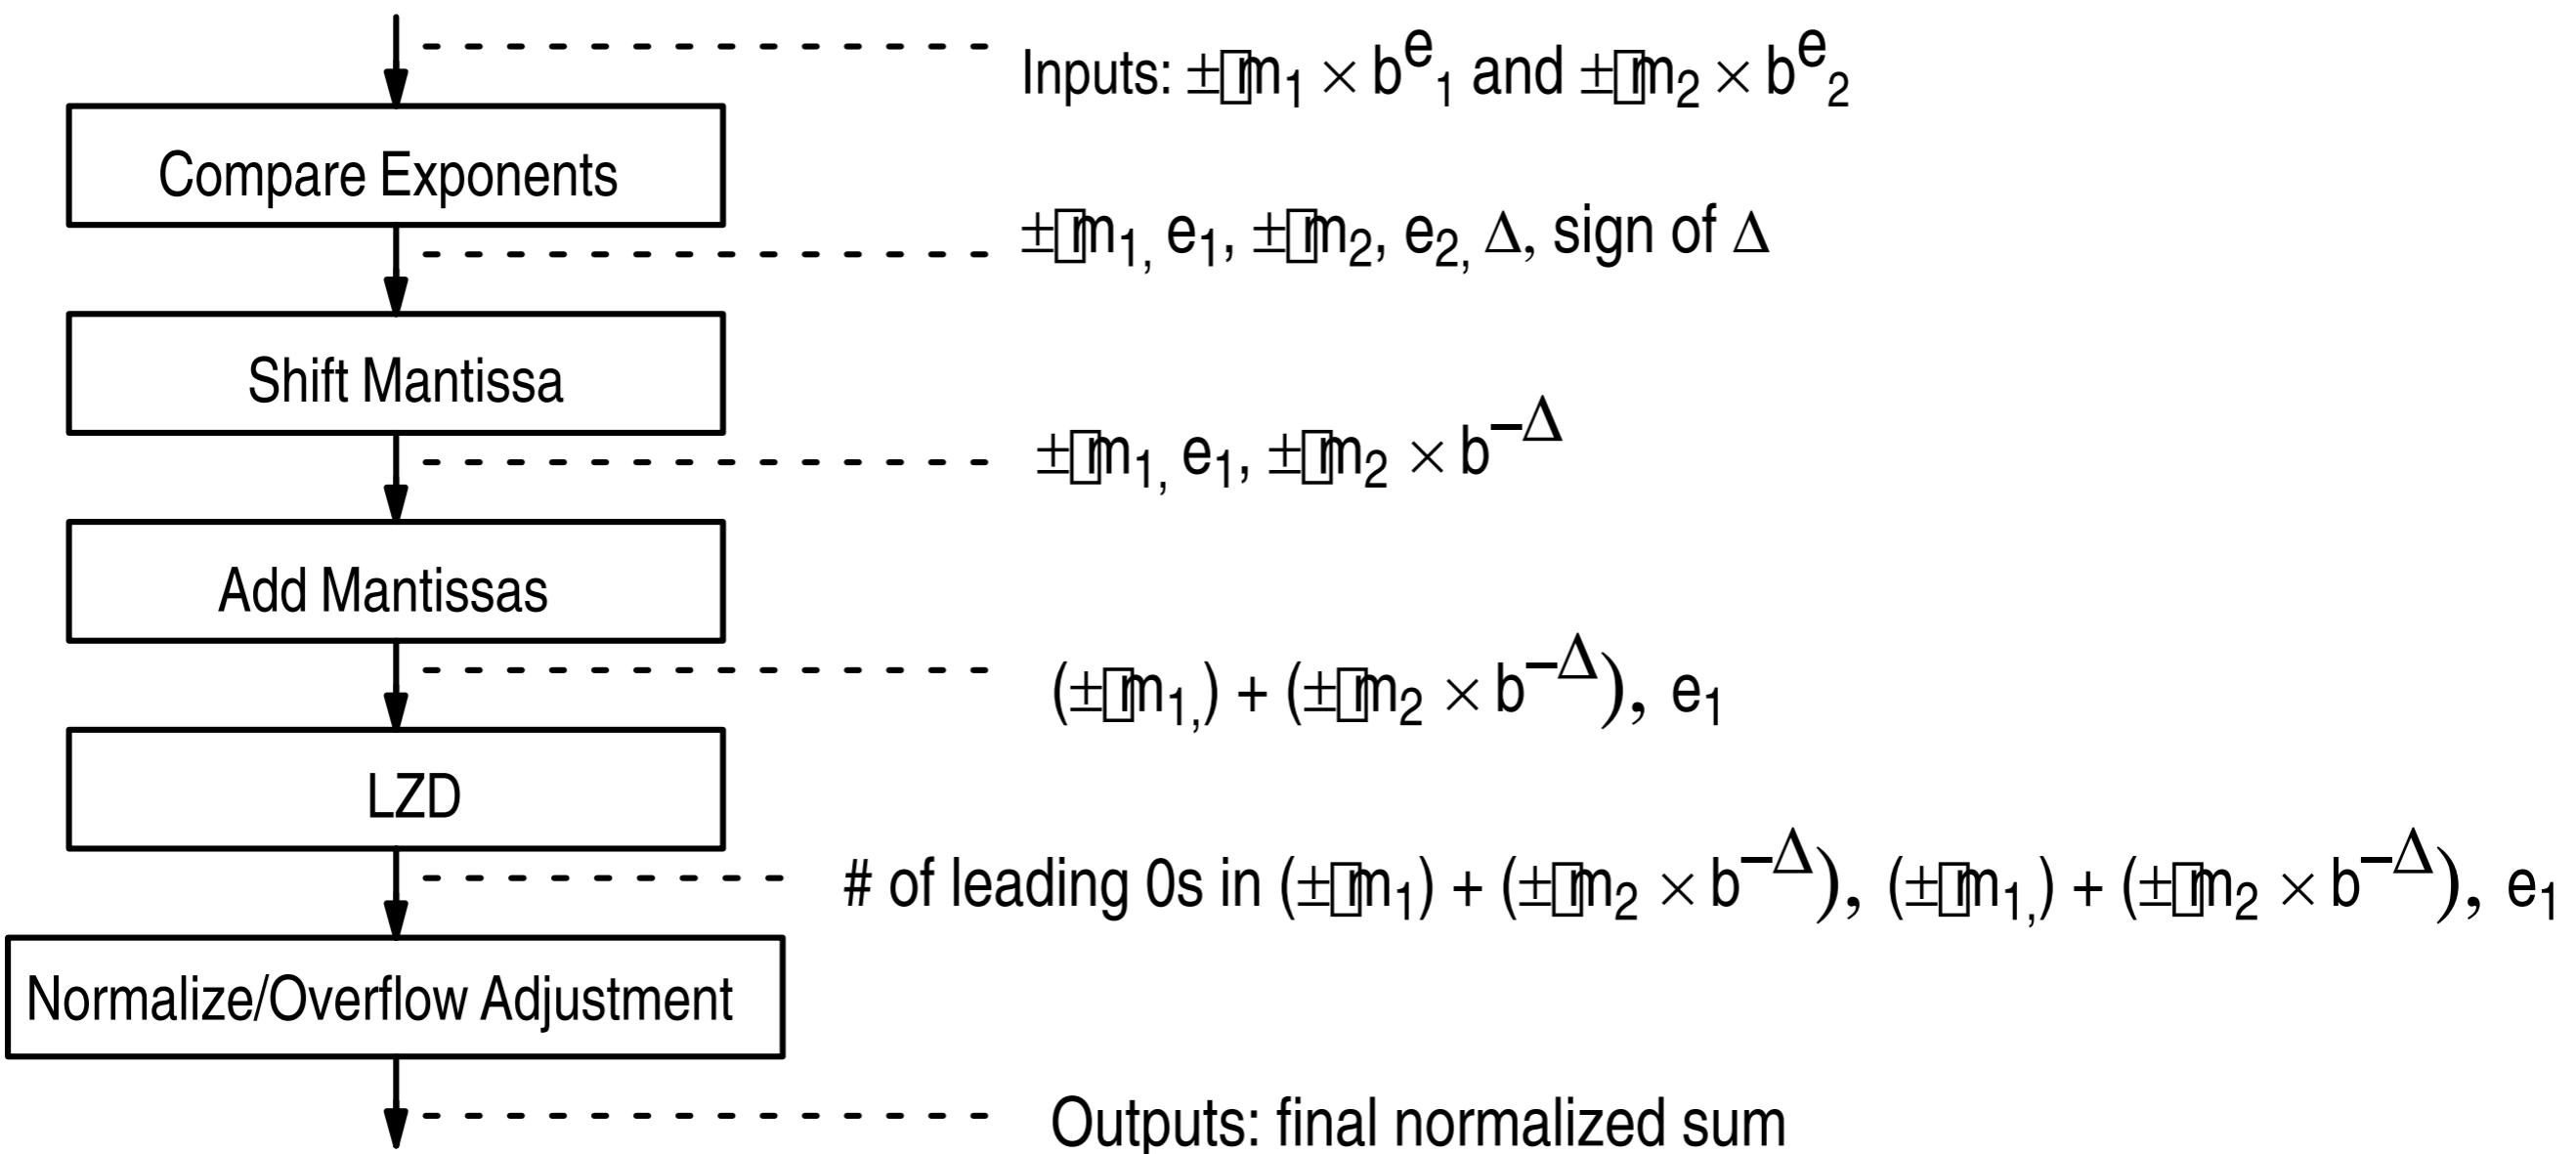
# Floating Point addition

- Input operands  
 $m_1 \times b^{e_1}$  and  $m_2 \times b^{e_2}$
- Base  $b$  usually power of 2, assume 2
- Normalized format
  - Leading 1 of mantissa needn't be stored
- Logic blocks:
  - Adders, comparators, barrel shifters, count leading zeros

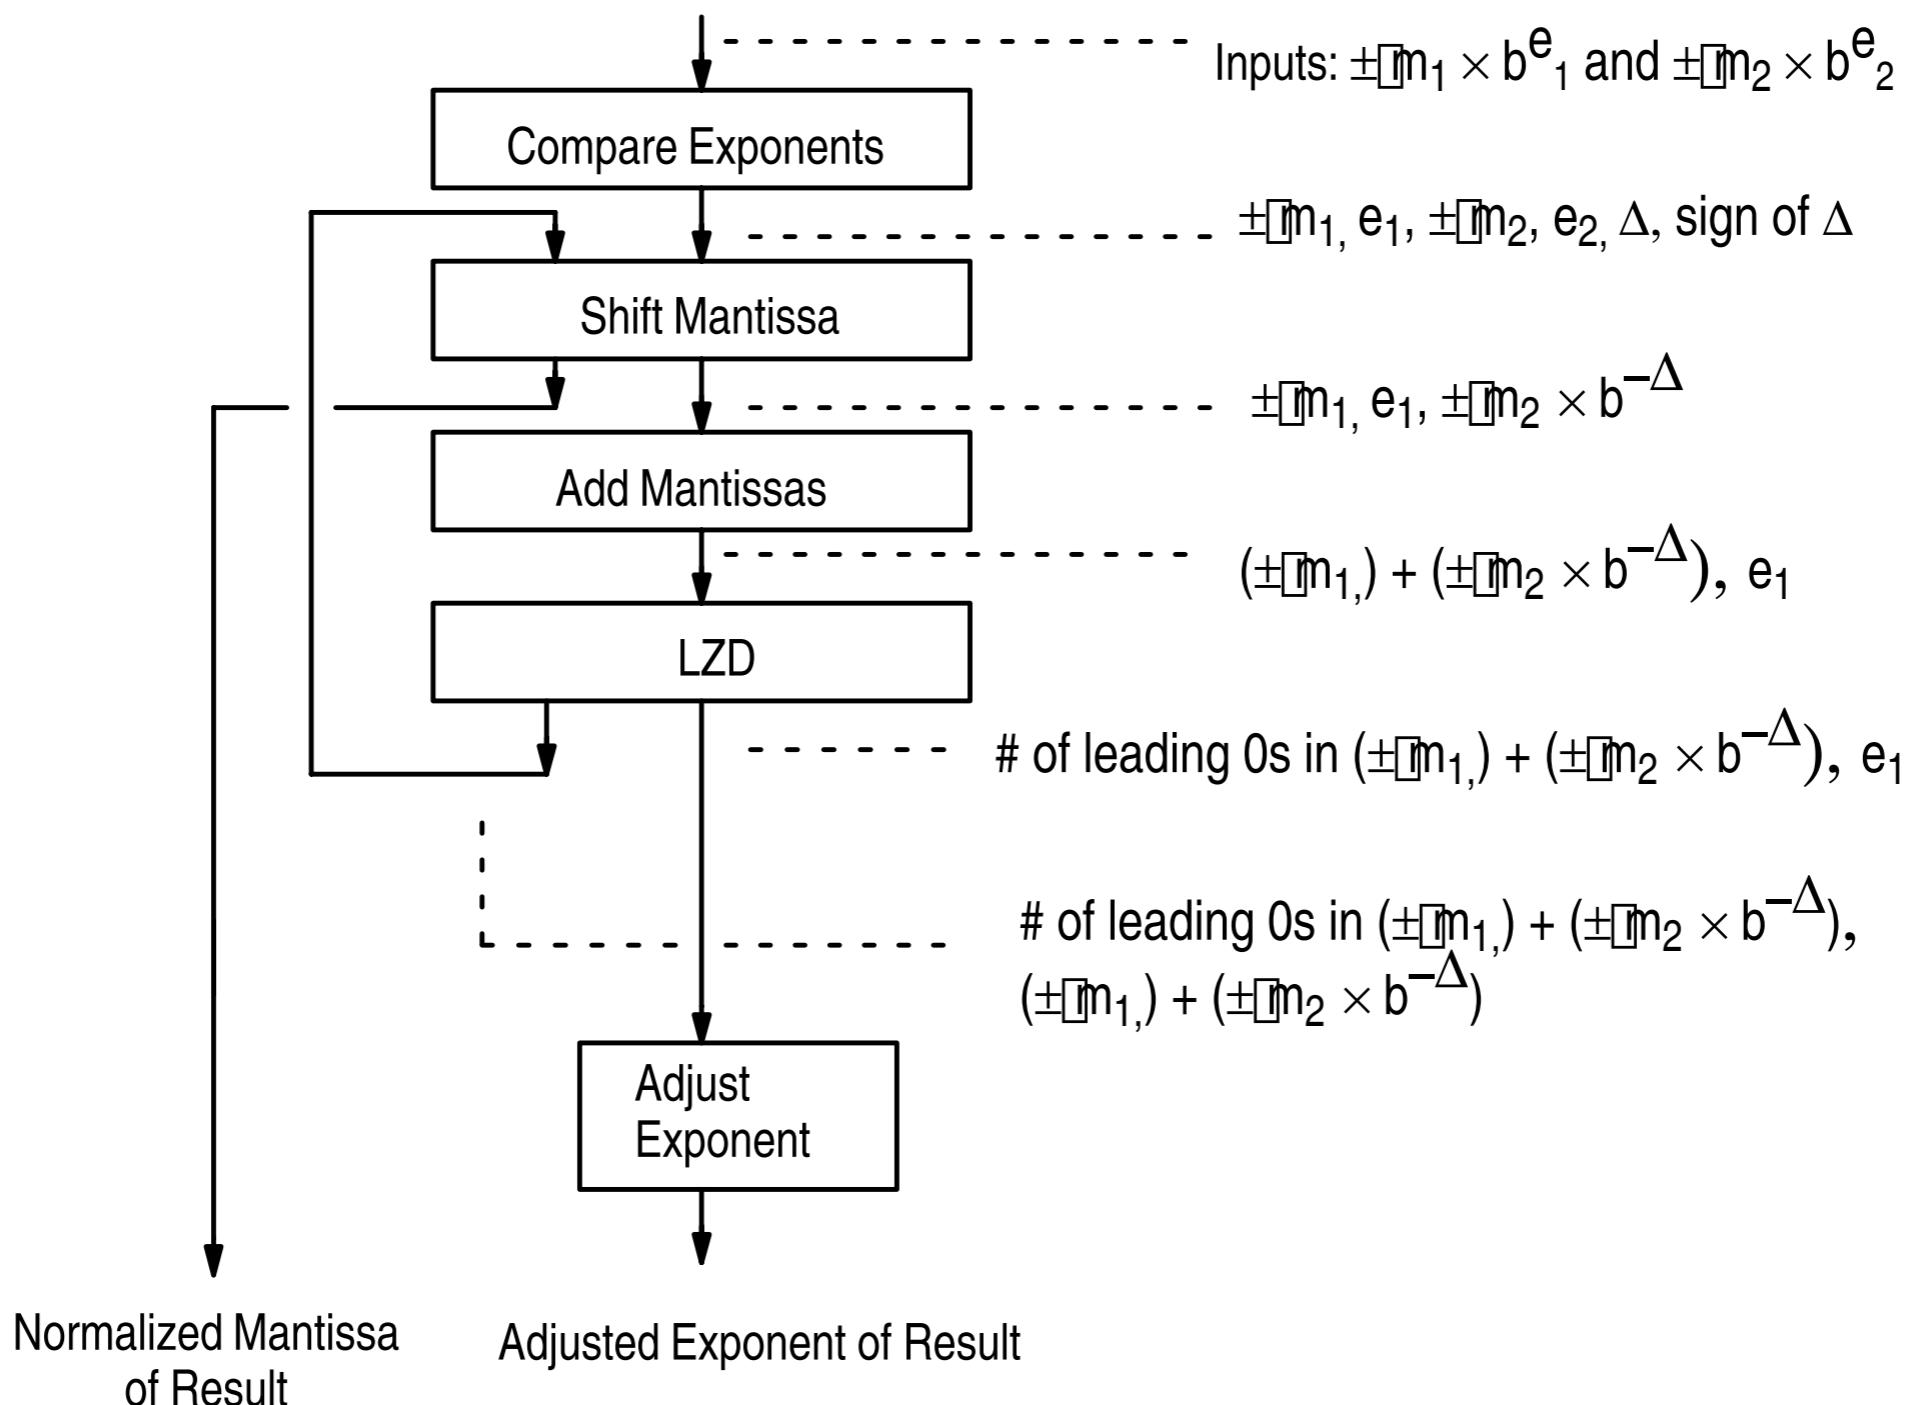
# FP ADD Algorithm

- 1: Compare exponents
  - $\Delta = | e_1 - e_2 |$
- 2: Shift  $m_1$  right to so  $e_1 = e_2$
- 3: Add  $m_3 = m_1 + m_2$ 
  - new exponent  $e_3 = e_1$
- 4: Normalize  $m_3$ 
  - Shift  $m_3$  left and dec  $e_3$  until MSB is 1

# FP Adder



# FP Adder (alternate)



# FP Add reservation tables

column number →

	1	2	3	4	5
Compare Exponents	X				
Shift Mantissa		X			
Add Mantissas			X		
LZD				X	
Normalize/ Overflow Adjustment					X

Reservation table for the floating point adder of Figure 2.7(a)

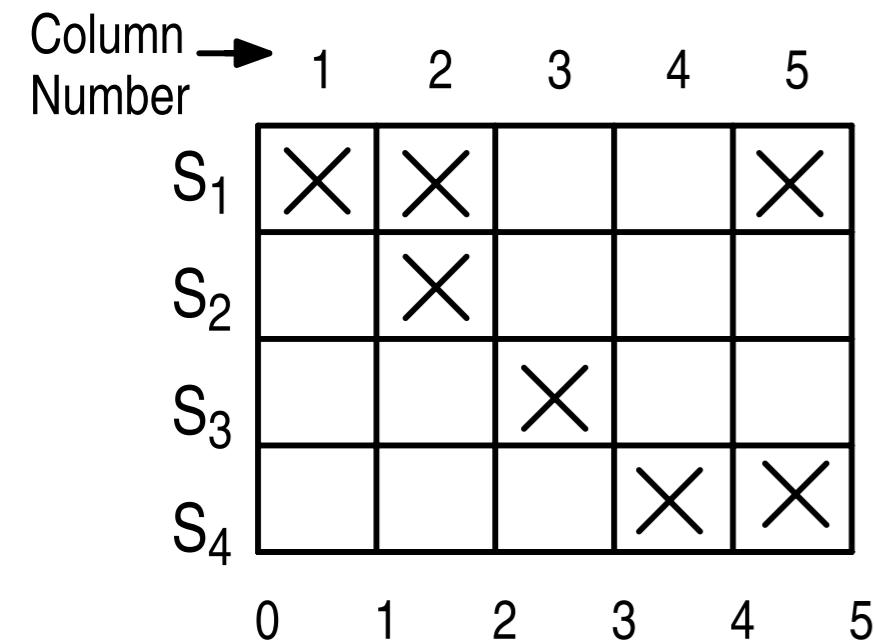
	1	2	3	4	5
Compare Exponents	X				
Shift Mantissa		X			X
Add Mantissas			X		
LZD				X	
Adjust Exponent					X

Reservation table for the floating point adder of Figure 2.7(b)

# Nonlinear pipeline throughput

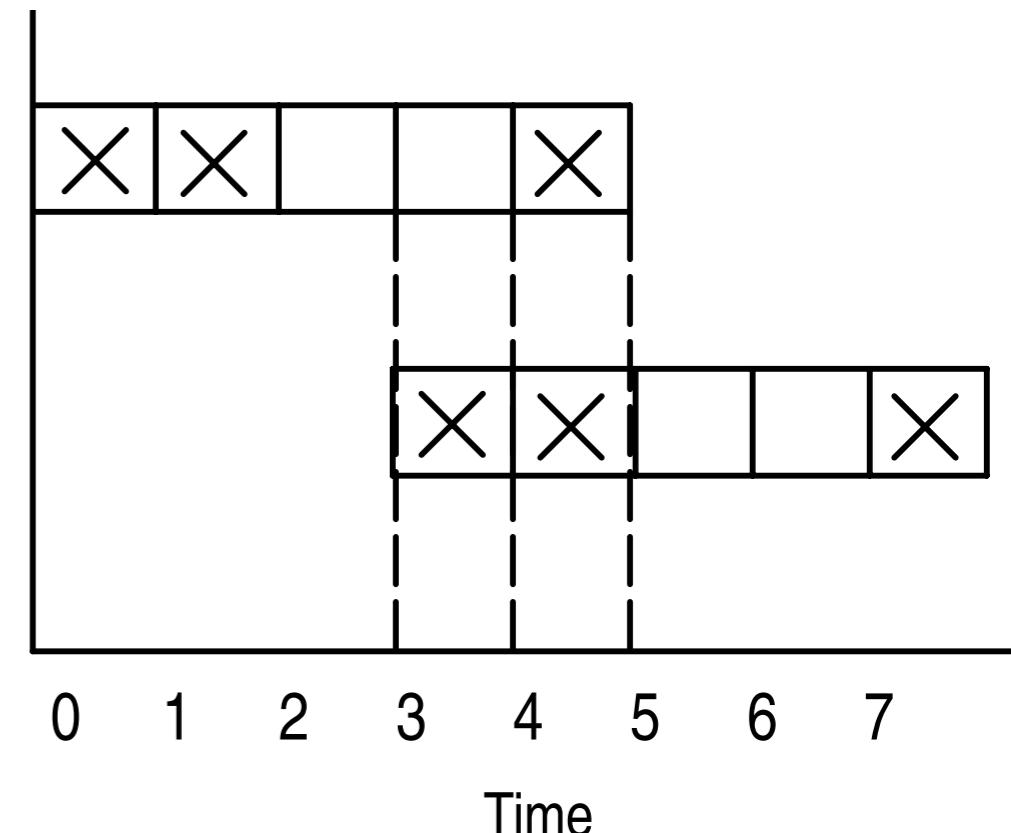
- Two or more X's in a row of res. table
  - Stage used more than once by same op.
  - Already started op may contend for use of this stage with an op started later.
    - Collision
    - Cannot initiate new op every cycle

# Nonlinear pipeline throughput



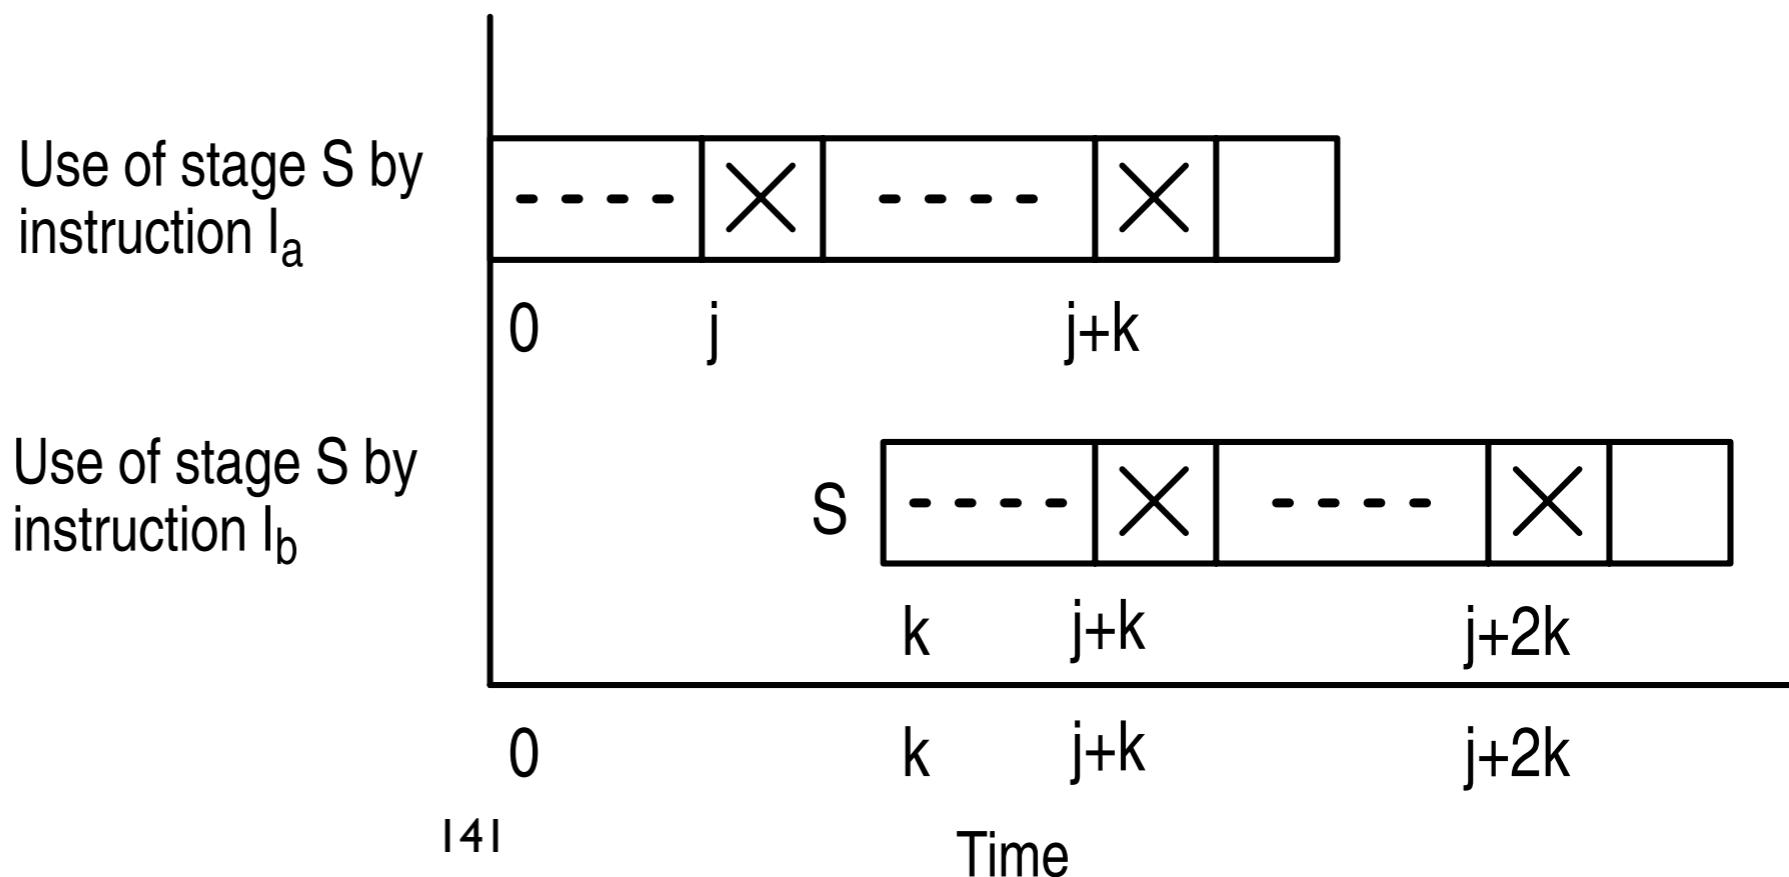
Use of stage S<sub>1</sub> by instruction I<sub>a</sub>

Use of stage S<sub>1</sub> by instruction I<sub>b</sub> initiated 3 cycles after I<sub>a</sub> was initiated



# Nonlinear pipeline throughput

- In general:
  - If the res. table has a row where any two X marks are separated by distance  $k$ , the two ops initiated at an interval  $k$  will collide.



# Collision avoidance:

## Static

- Scheduling interval not equal to  $k$
- Can be done by compiler
- Trivial:
  - Any distance higher than pipeline latency
  - Any distance higher than max of all distances between any two X's in same row.
- Better:
  - More complex schedules that produce best throughput while avoiding collisions.

# Collision avoidance: Dynamic

- Use hardware to avoid collisions -- delay initiation of new operation to avoid collision with already-issued ops.

# Terminology

- ***Initiation***: Starting of new op in pipeline
- ***Initiation latency***: Time between two successive initiations
- ***Initiation cycle***: Repetitive sequence of initiations
- ***Forbidden latency***: Initiation latency that causes collisions
- ***Collision-free initiation***: Initiation that does not collide with already-initiated ops.

# Consequence of feedback pipeline

- $M$  = number of X's in row with most X's
- $T_p$  = Runtime on feedback architecture
- $N$  = Number of instructions
- Busiest stage uses  $N*M$  cycles.
- Efficiency:  $\epsilon = N * M / T_p \leq 1$
- Throughput:  $\tau = N / T_p$
- Implication:  $\tau \leq 1 / M$

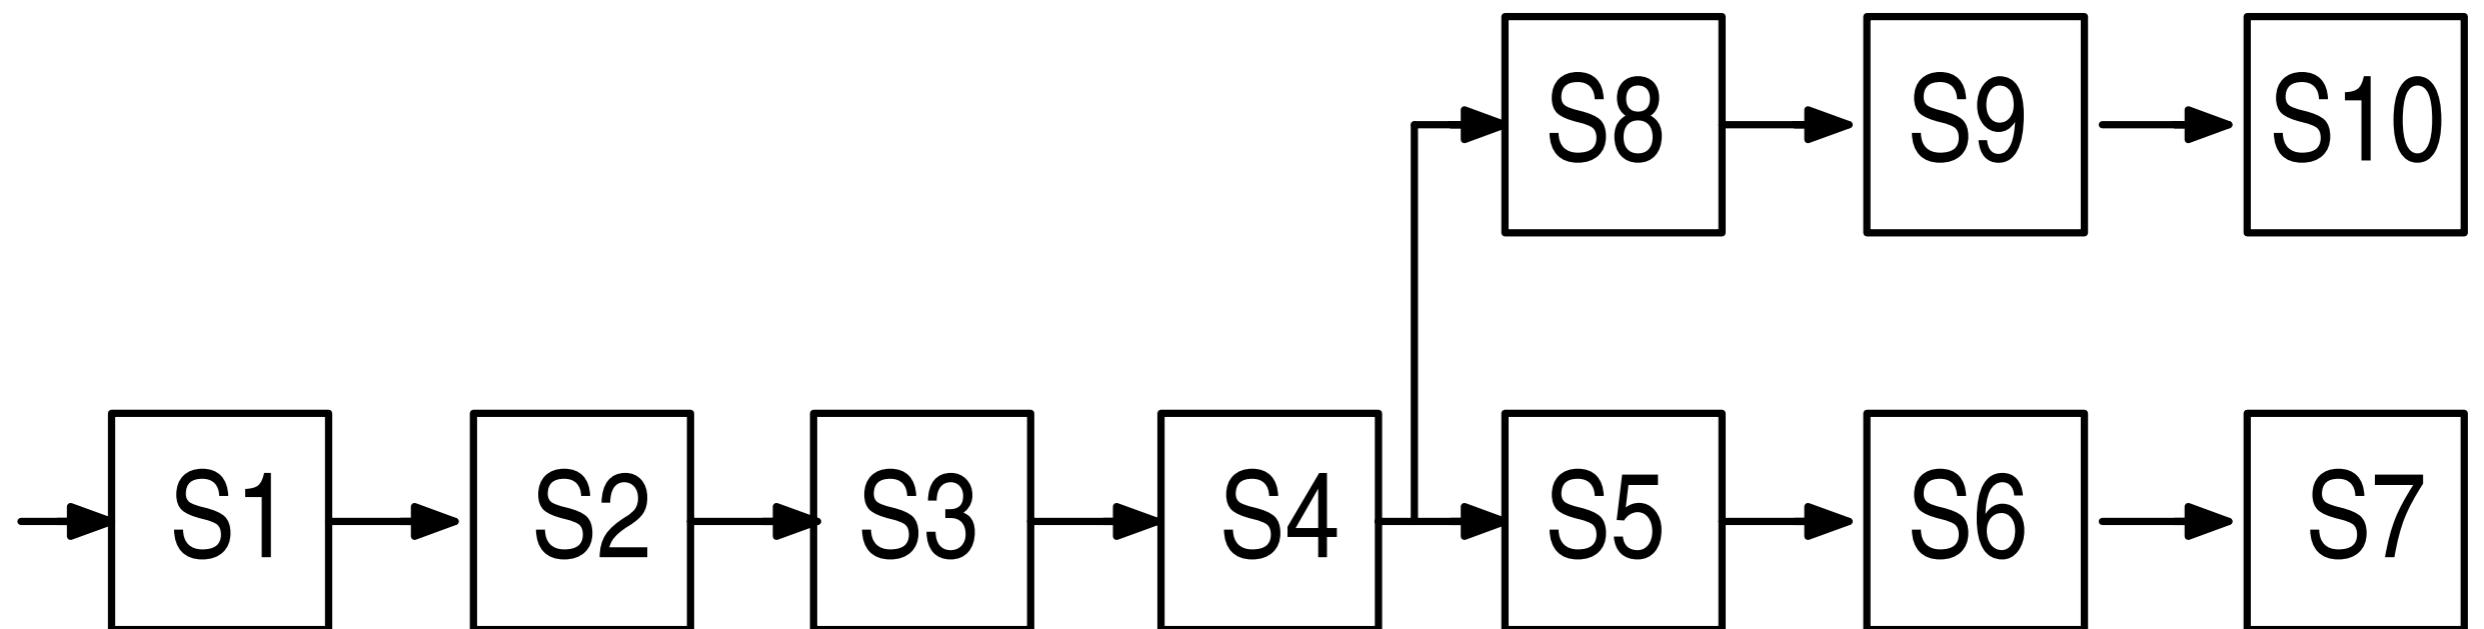
$$\text{MAT} \leq \text{MPRT} = 1/M$$

# Limitations on pipelining

- Factors
  - Clock skew
  - Latch overhead
  - Pipeline bubbles, data dependencies
  - State info to save/restore on interrupt
- Consequences
  - Limits pipelines to 4 to 8 stages
  - Longer pipelines are *decoupled*, using queues.

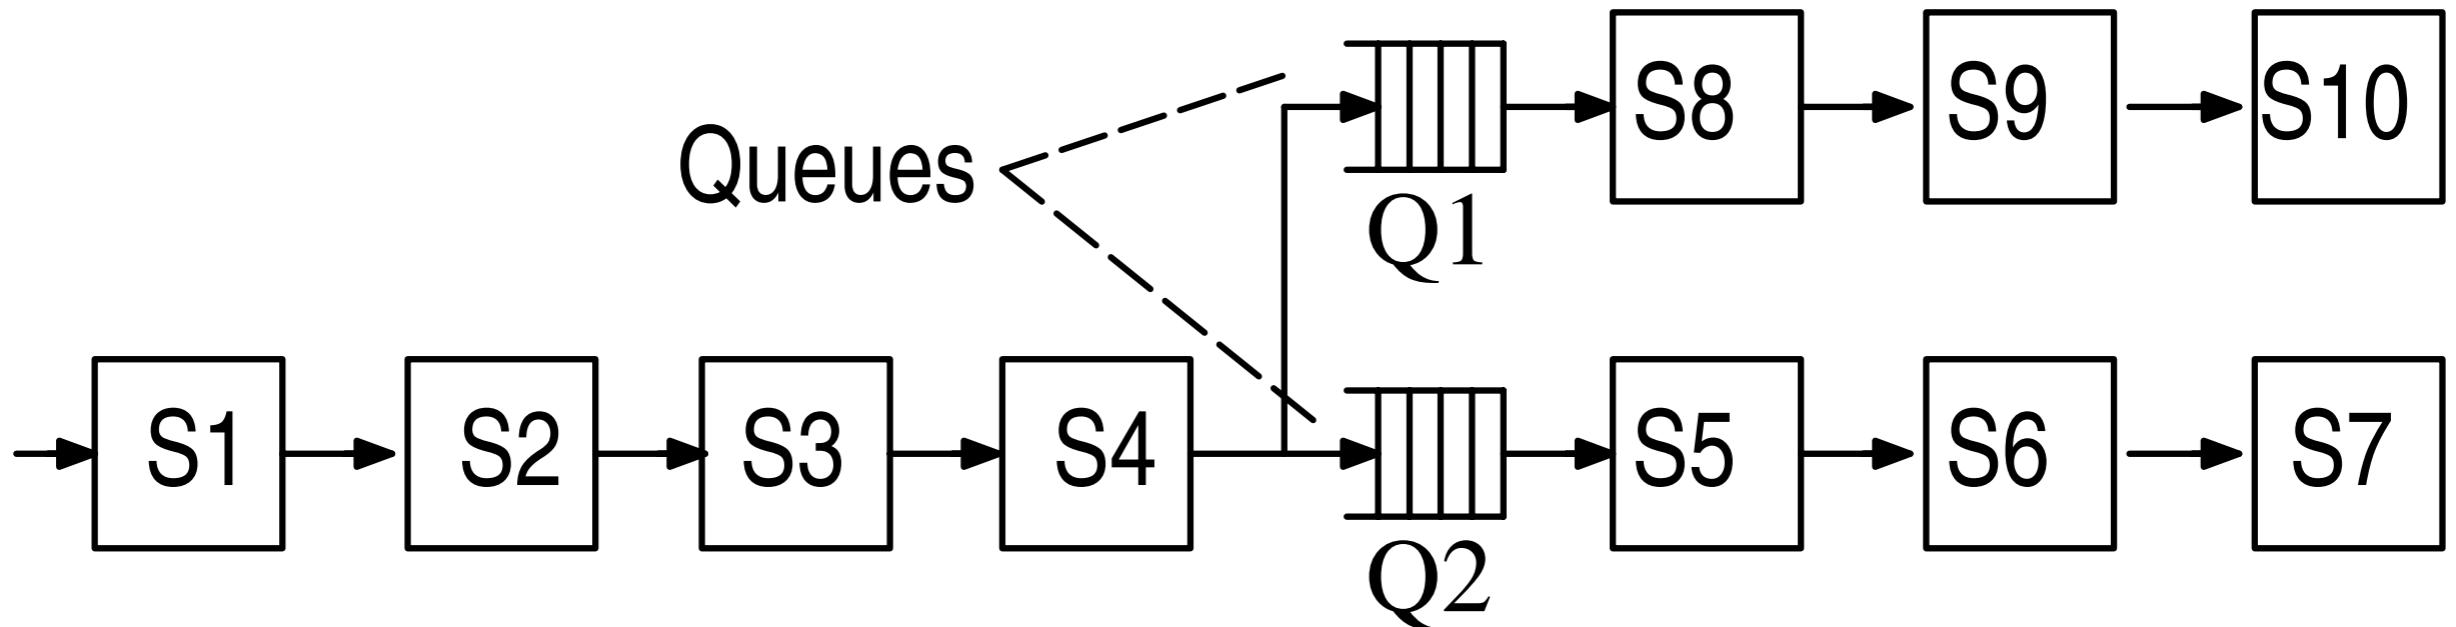
# Non-decoupled pipeline

- Clock rate accounts for larger clock skew
- Blocking stage can stall whole pipeline



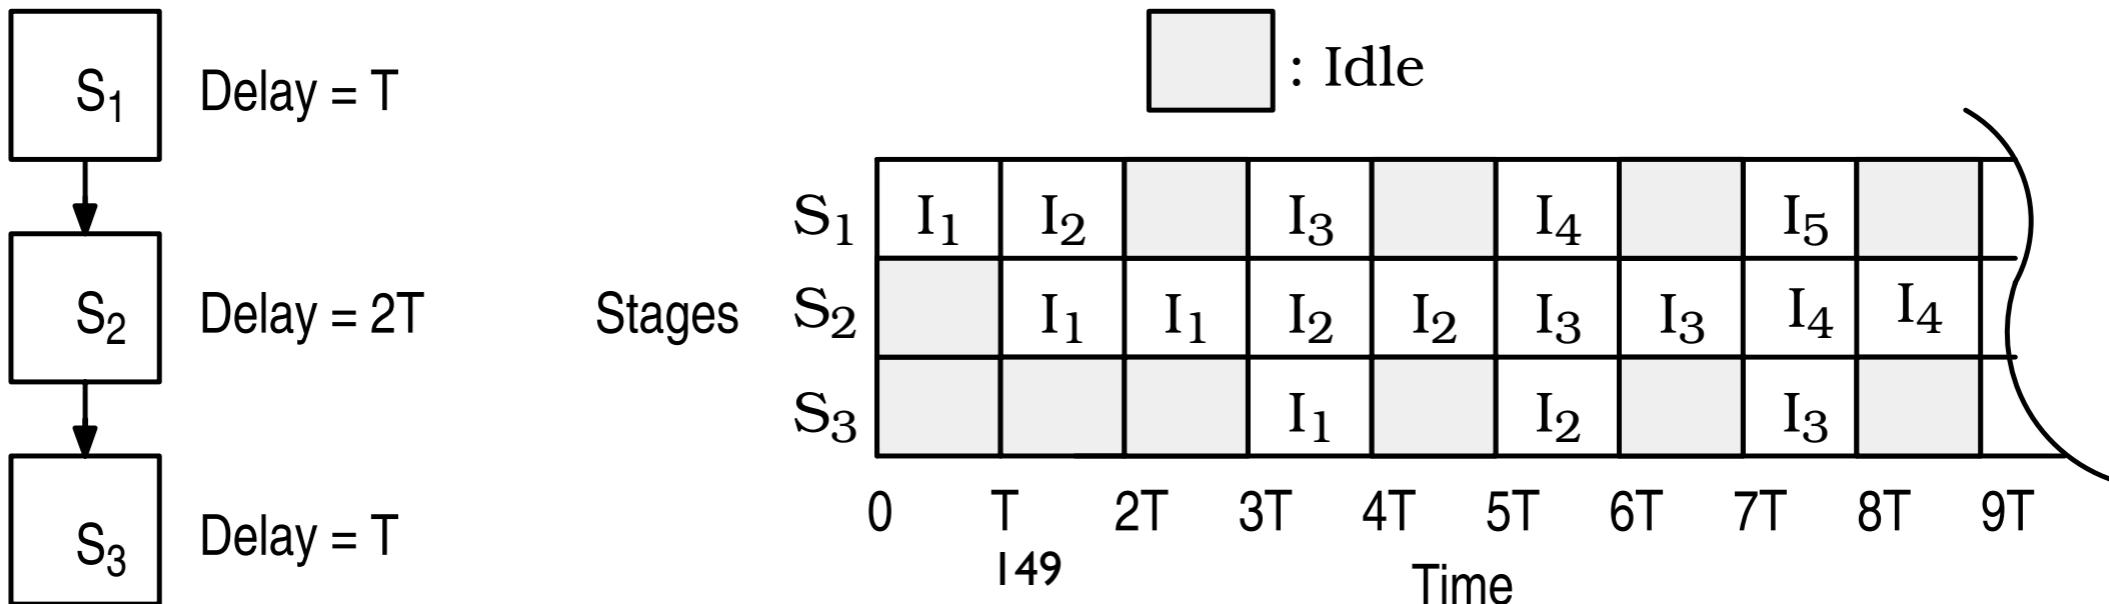
# Decoupled pipeline

- Three independent pipeline segments
- Independently clocked (same frequency)
- Queues help hide stalls



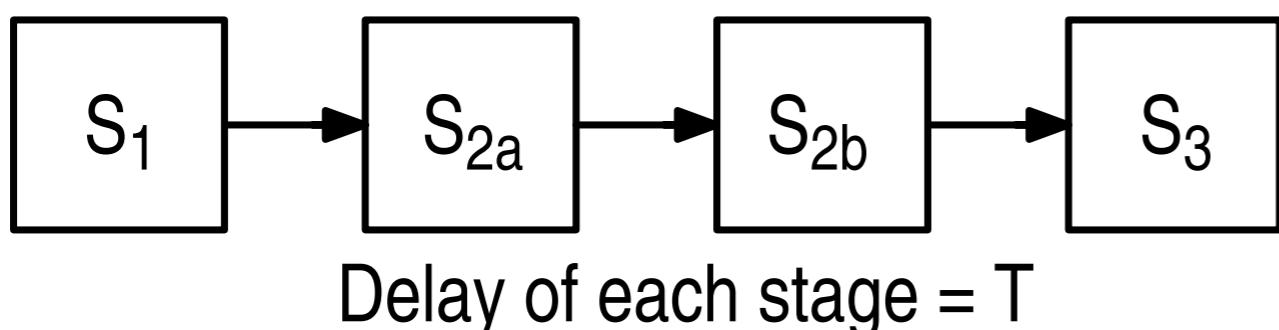
# Impact of Pipeline Bottlenecks

- K-stage pipeline with non-uniform stage delays:  $T_p = \sum_i T_i + (N-1).T_s$
- $T_s = \max \{ T_i \}$
- Instructions complete one every  $T_s$
- Slowest stage is performance bottleneck.
- Stage with delay  $T_i$  unutilized for  $(T_s - T_i)$



# Segmentation

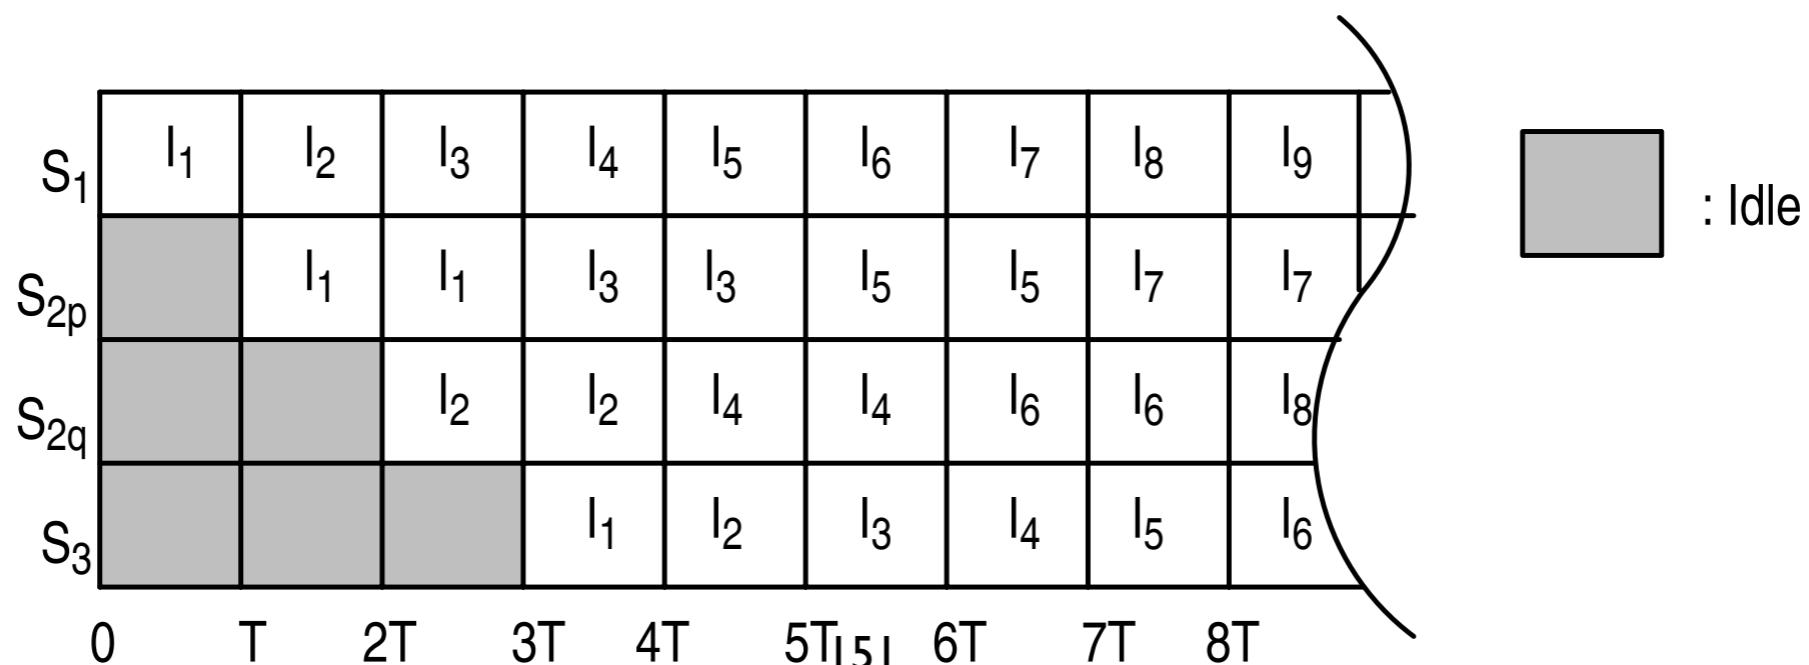
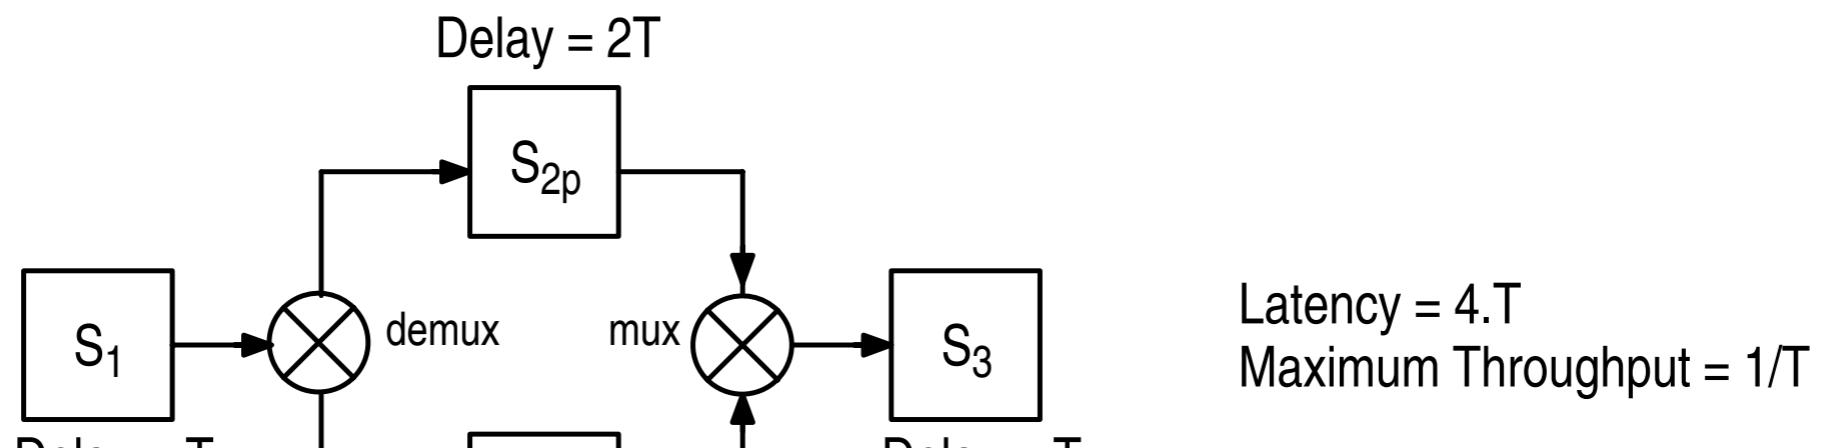
- Split slower stages into sequences of simpler stages.



Latency =  $4.T$   
Maximum Throughput =  $1/T$

# Replication

- Replicate slower stages and alternate

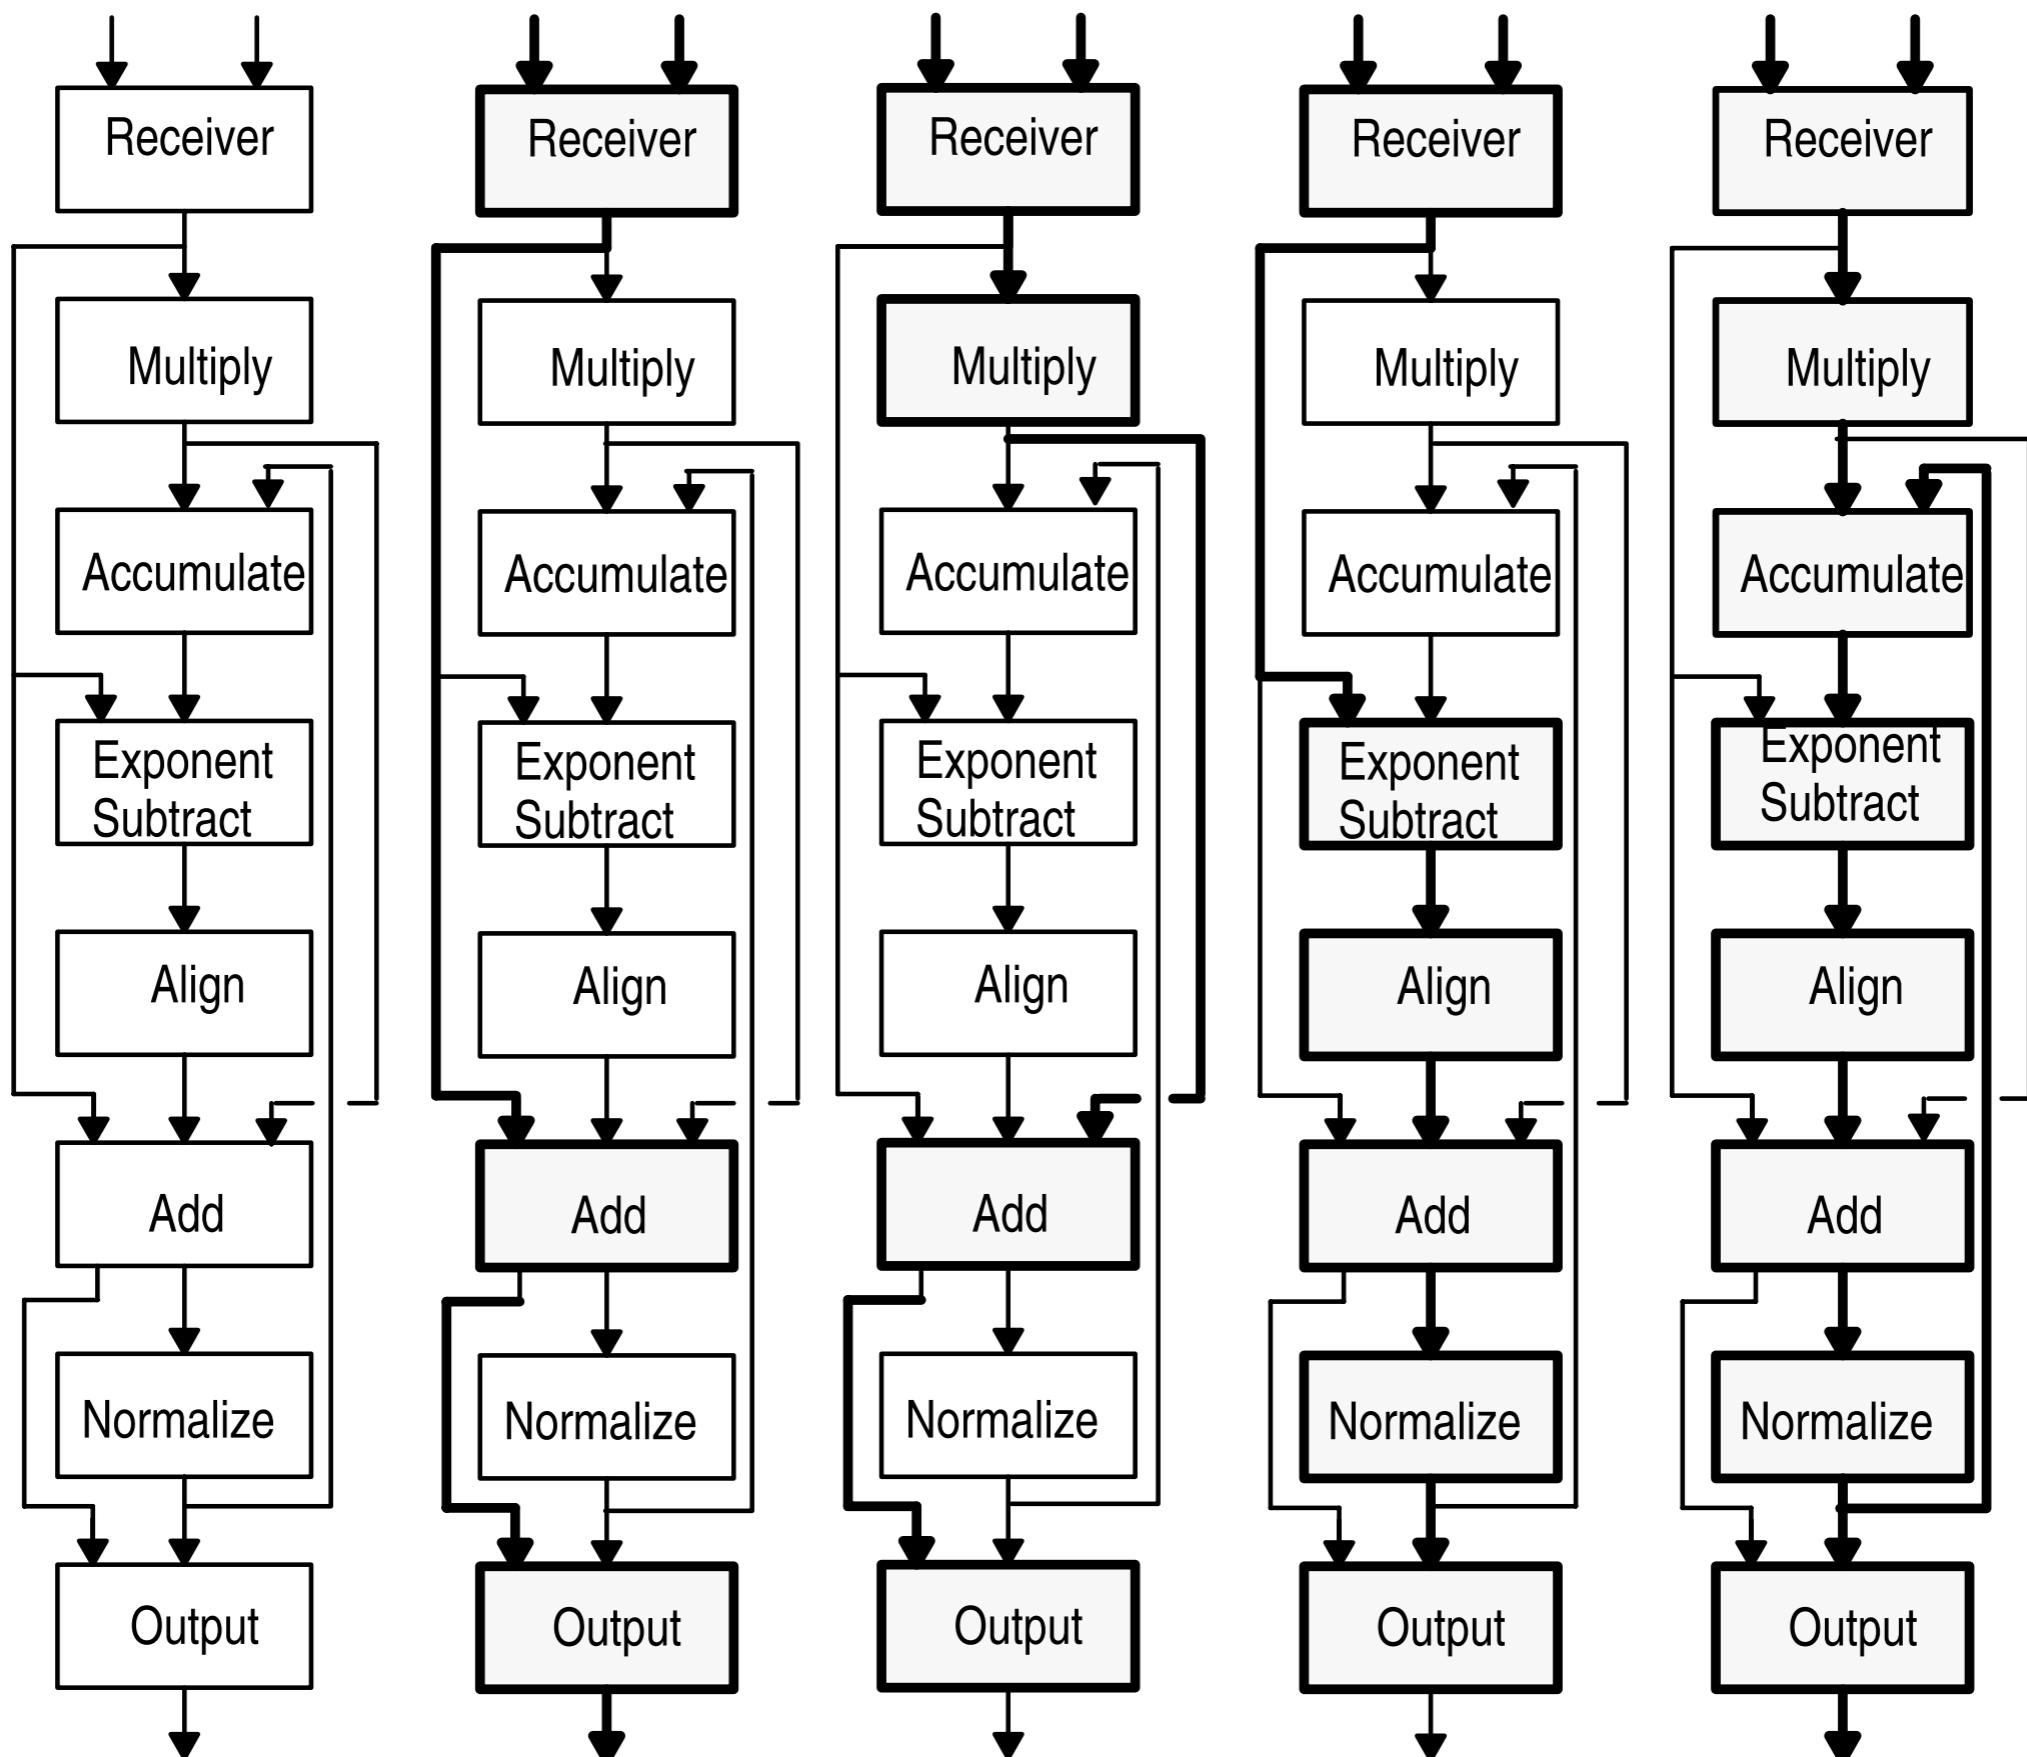


# Pipeline Classifications

- Unifunction vs. multifunction
  - Unifunction: Pipeline implements one function
  - Multifunction: Common pipeline for multiple operations
- Static vs. Dynamic
  - Static: Fixed function
  - Dynamic: Reconfigurable on the fly

# Multifunction pipelines

- Example: The arithmetic unit pipeline of TI ASC (Advanced Scientific Computer)
- All reconfiguration done by microcode.
- Cycle time is 60ns.
- Inputs and outputs are vectors.



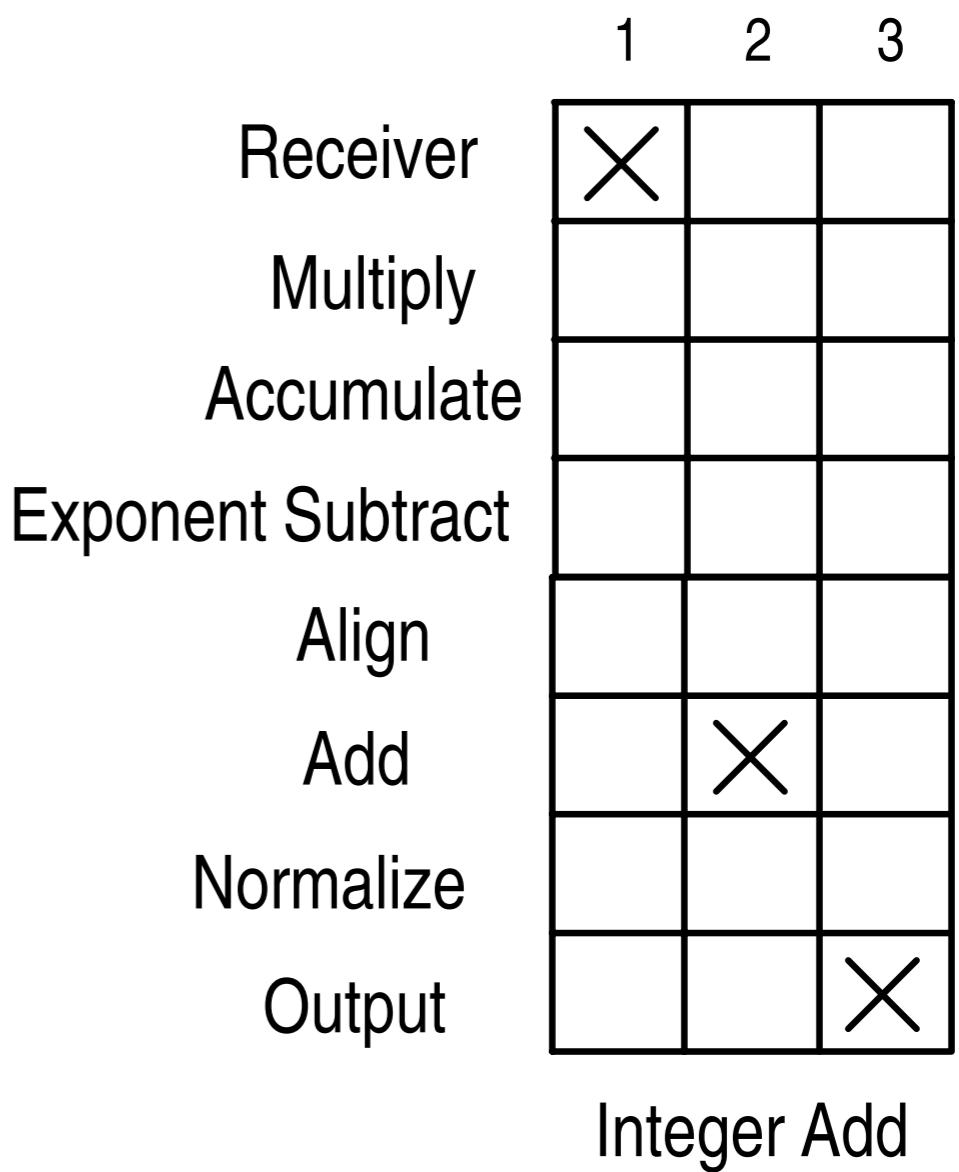
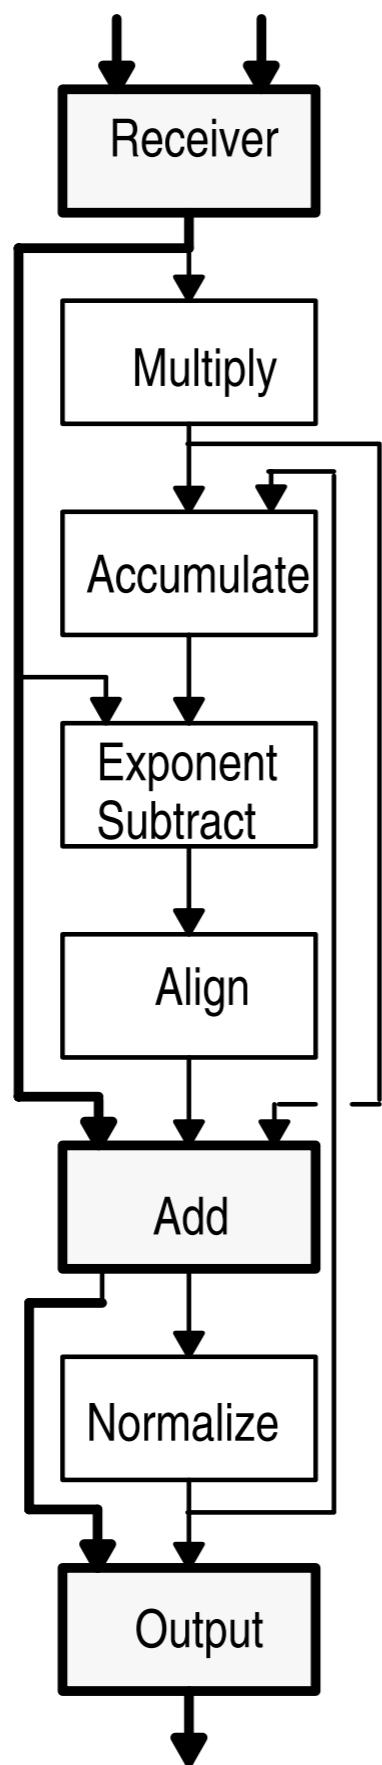
Basic Pipeline

Integer Add

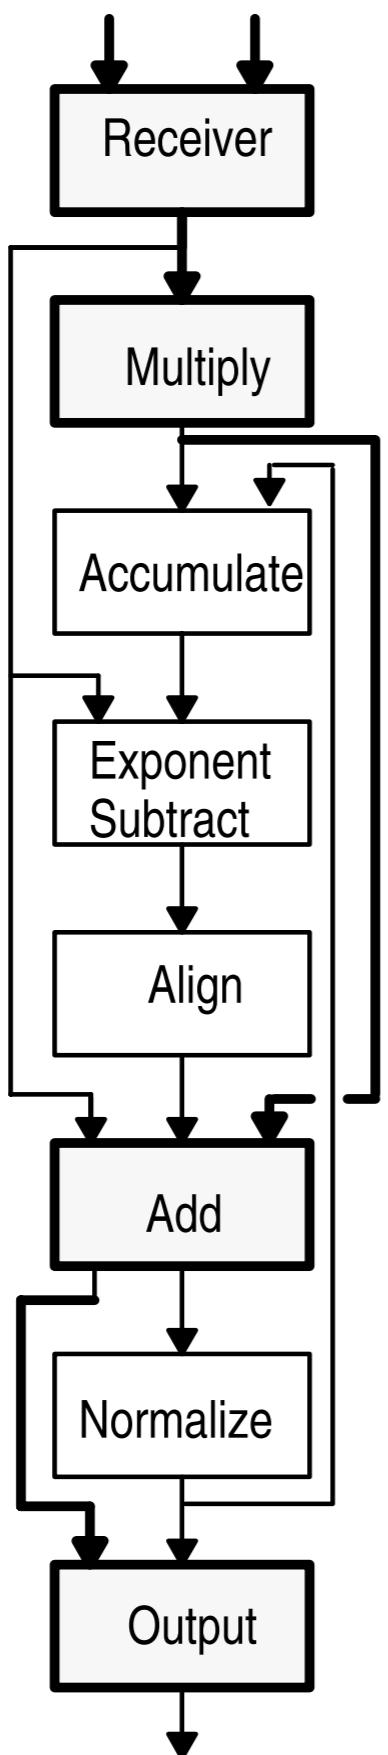
Integer Multiply

Floating  
Point  
Add

Floating Point  
Vector Dot Product



Integer Add

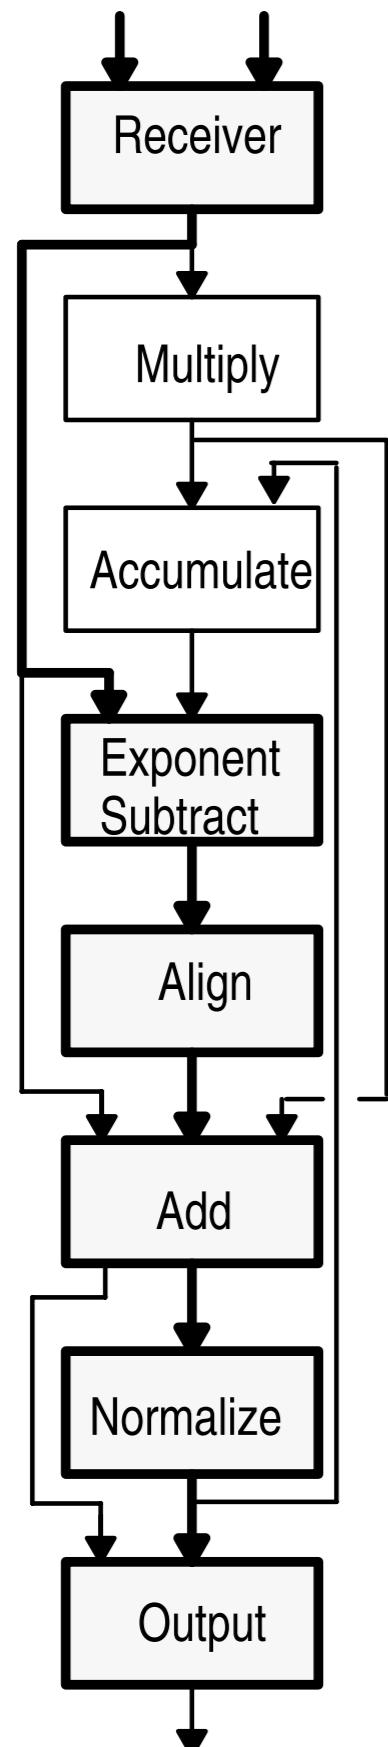


	1	2	3	4
Receiver	X			
Multiply		X		
Accumulate				
Exponent Subtract				
Align				
Add			X	
Normalize				X
Output				

Integer Multiply

Adds a zero

Integer Multiply



Floating Point Add

1      2      3      4      5      6

Receiver

Multiply

Accumulate

Exponent Subtract

Align

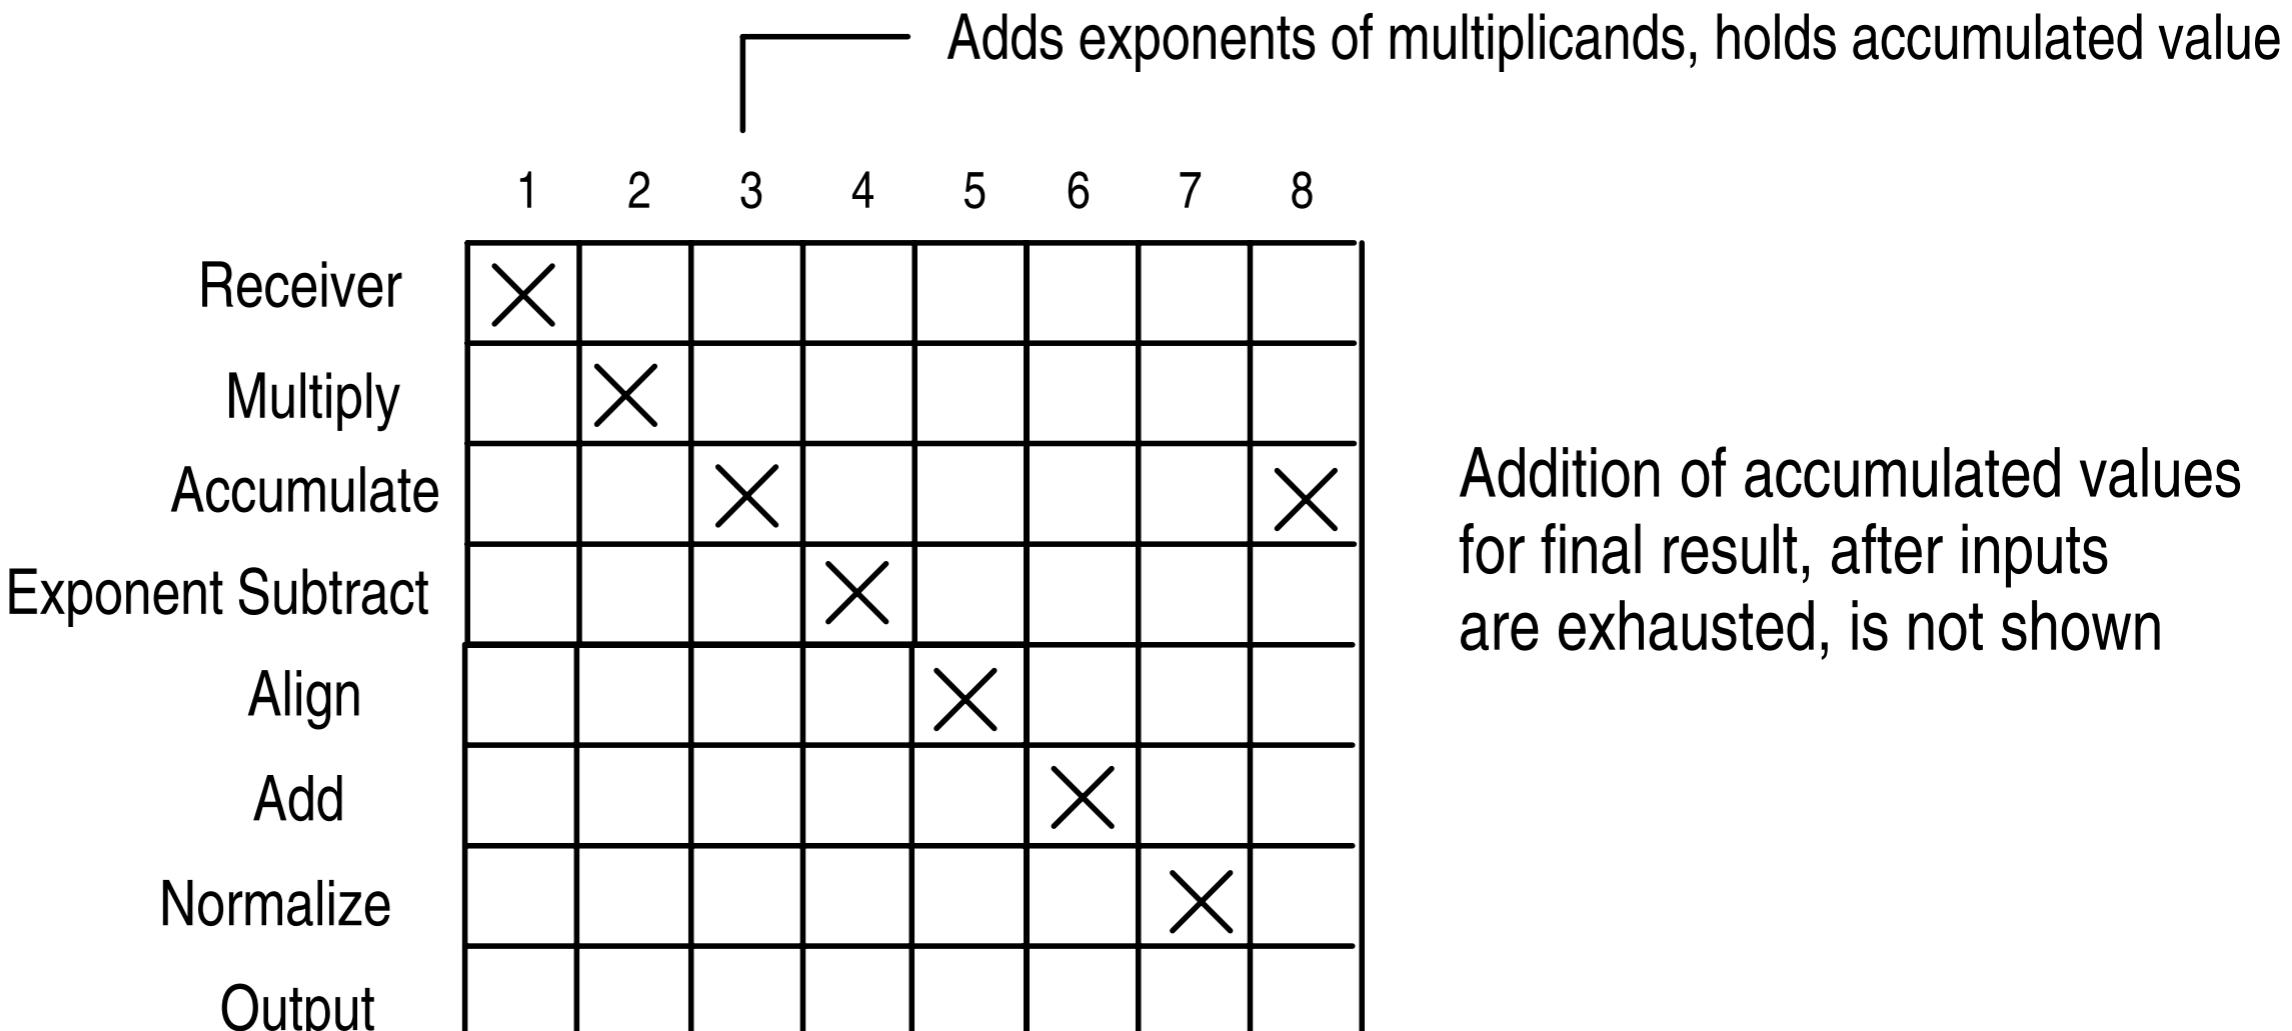
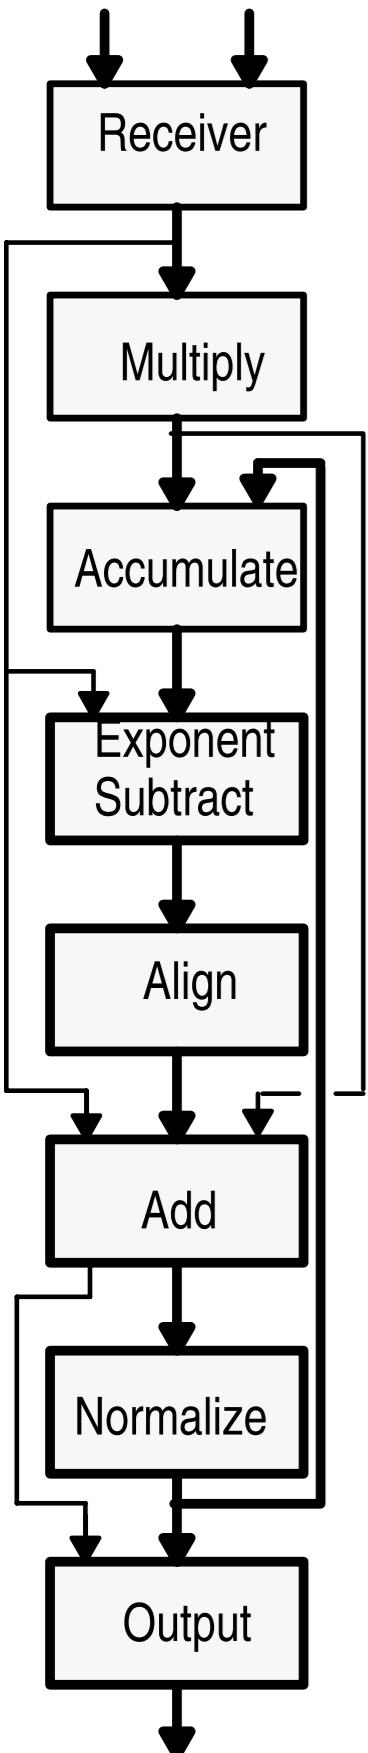
Add

Normalize

Output

1	2	3	4	5	6
X					
	X				
			X		
				X	
					X
					X

Floating Point Add



## Vector Inner Product

# Multifunction pipelines: Software controlled

- Fixed multiply & add  
 $\text{operand\_1} * \text{operand\_2} + \text{operand\_3}$
- Three functions:
  - Multiply and add
  - Multiply: Force  $\text{operand\_3}$  to be zero
  - Addition: Force  $\text{operand\_1}$  to be one
- Reservation table same for all three