

Speculative Execution

Main Idea

- General: Execute instructions that may not be needed, before it is known whether or not they will be needed.
- Optimistic: Execute instructions whose correctness is uncertain.
- If execution turns out to be incorrect or unnecessary, discard or undo their results.
- Examples:
 - Exceptions
 - Branch prediction
 - Memory reordering (hoisting)

Speculative Execution for Branches

- Use prediction to guess branch direction.
- Execute along speculative path.
- Architectural state is not committed until after branch direction is resolved.
- Major advancement over delayed branch and squashing

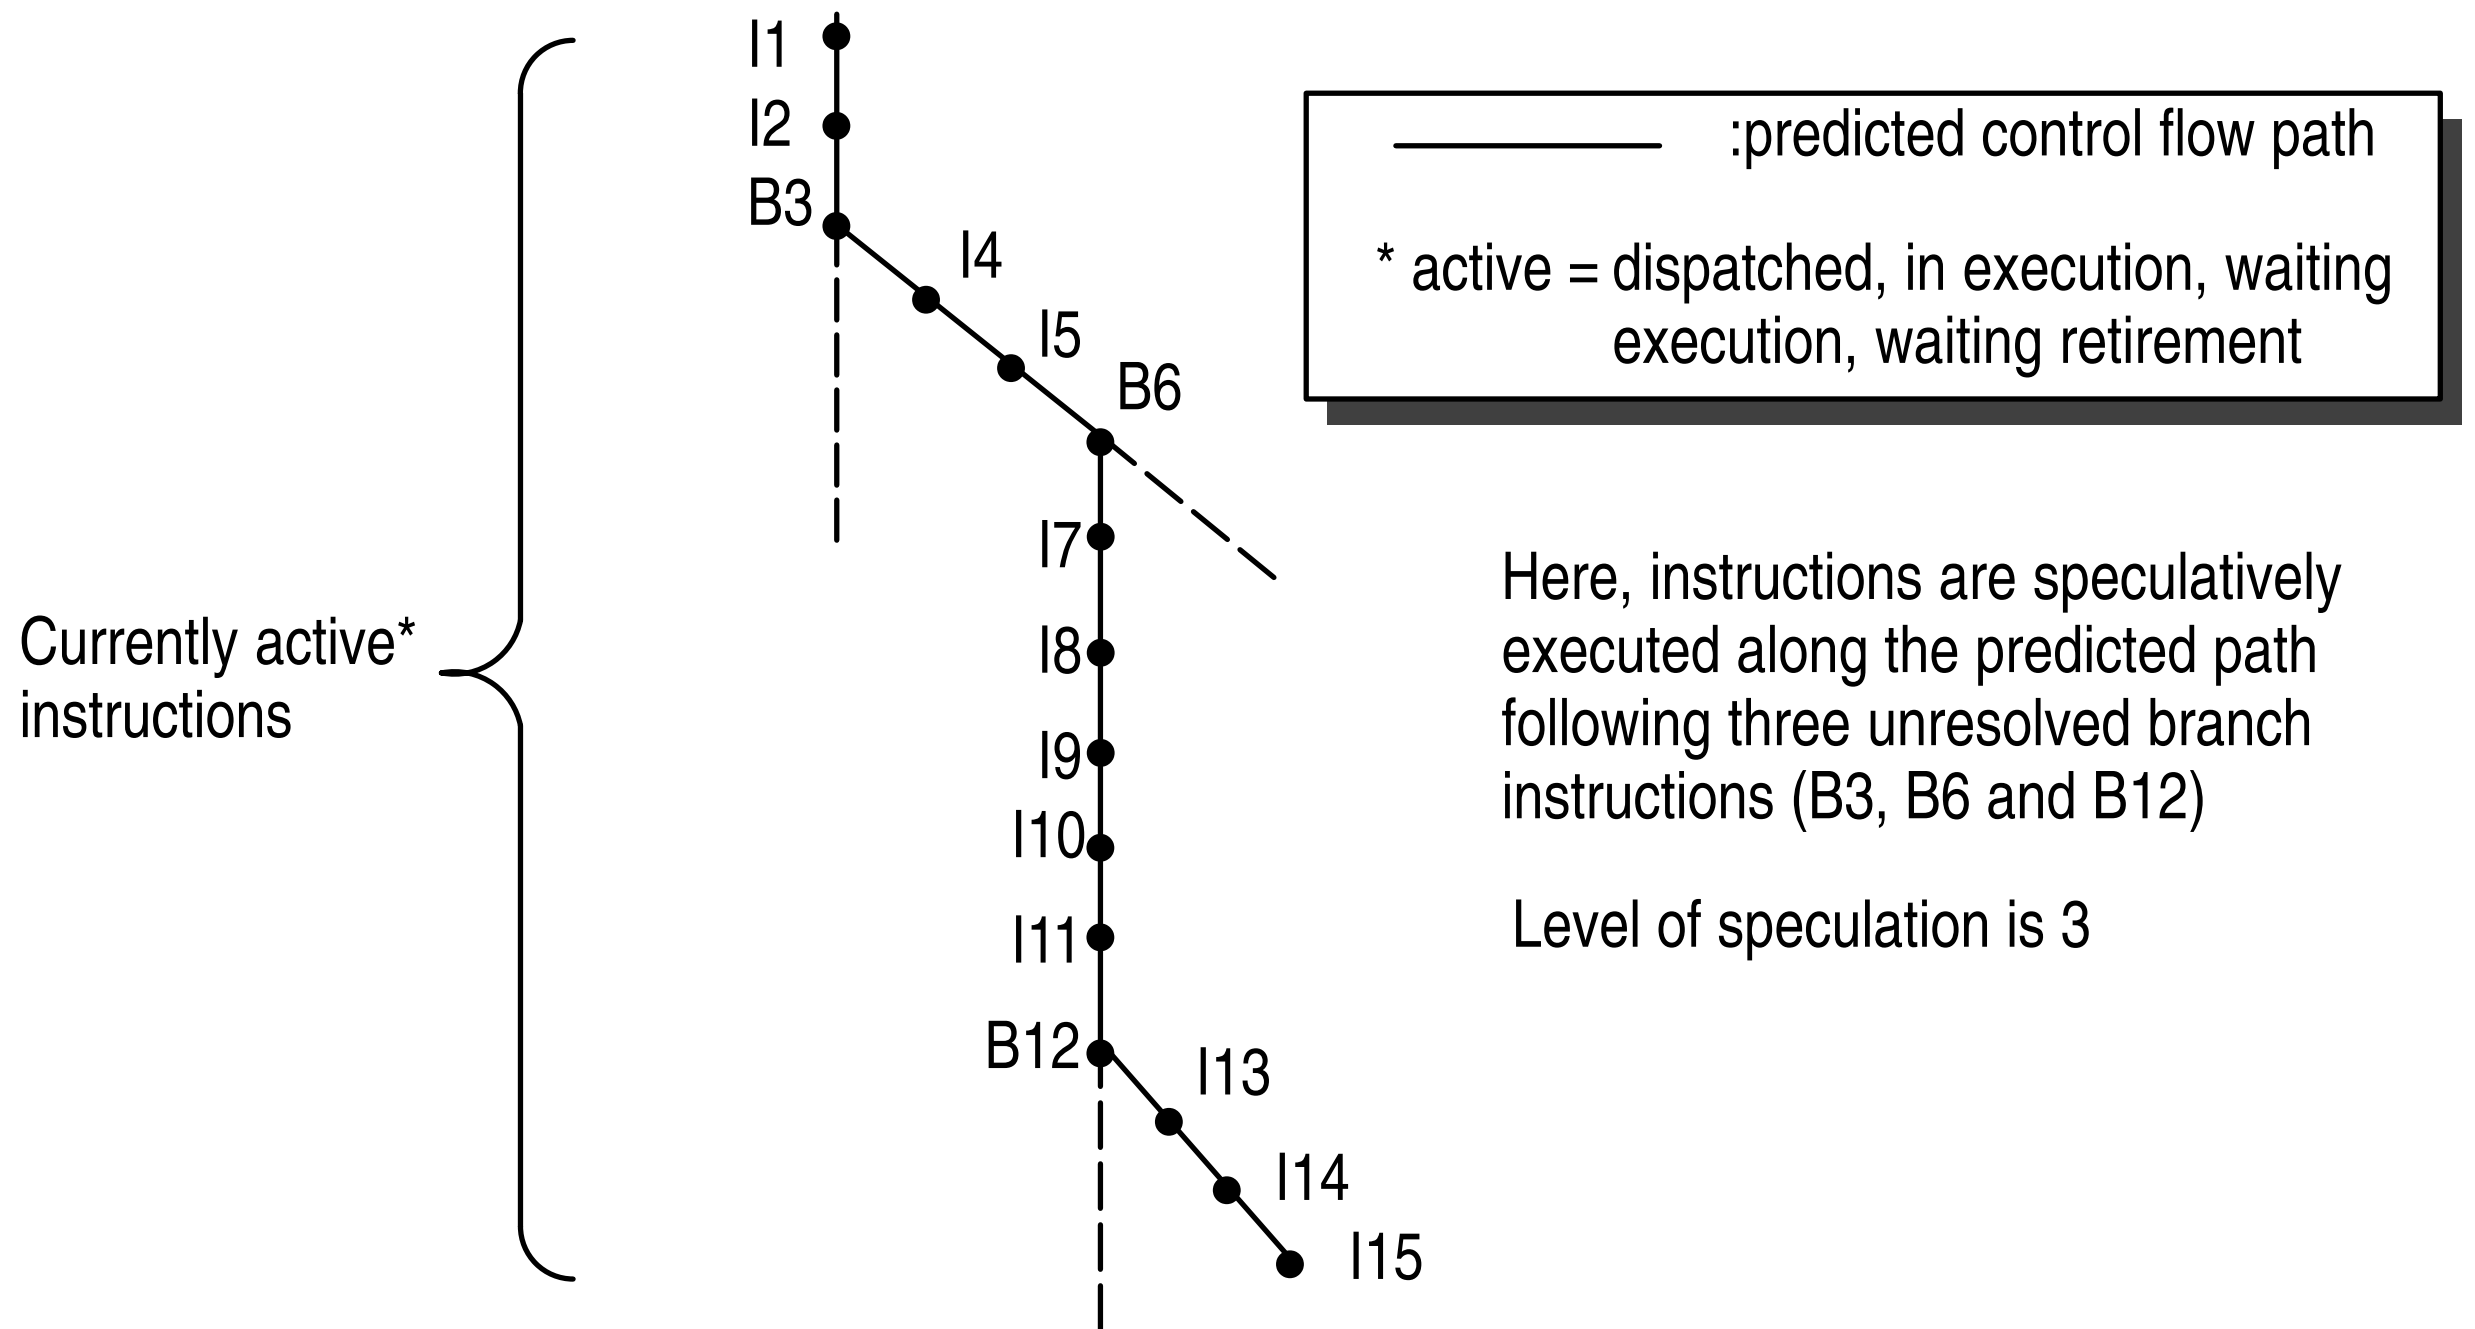
Degrees of Speculation

- Just Fetch?
- Fetch and Decode?
- Fetch, Decode, and Dispatch?
- Fetch, Decode, Dispatch, Complete with forwarding, but don't update architectural state. [Most common solution]

Levels of Speculation

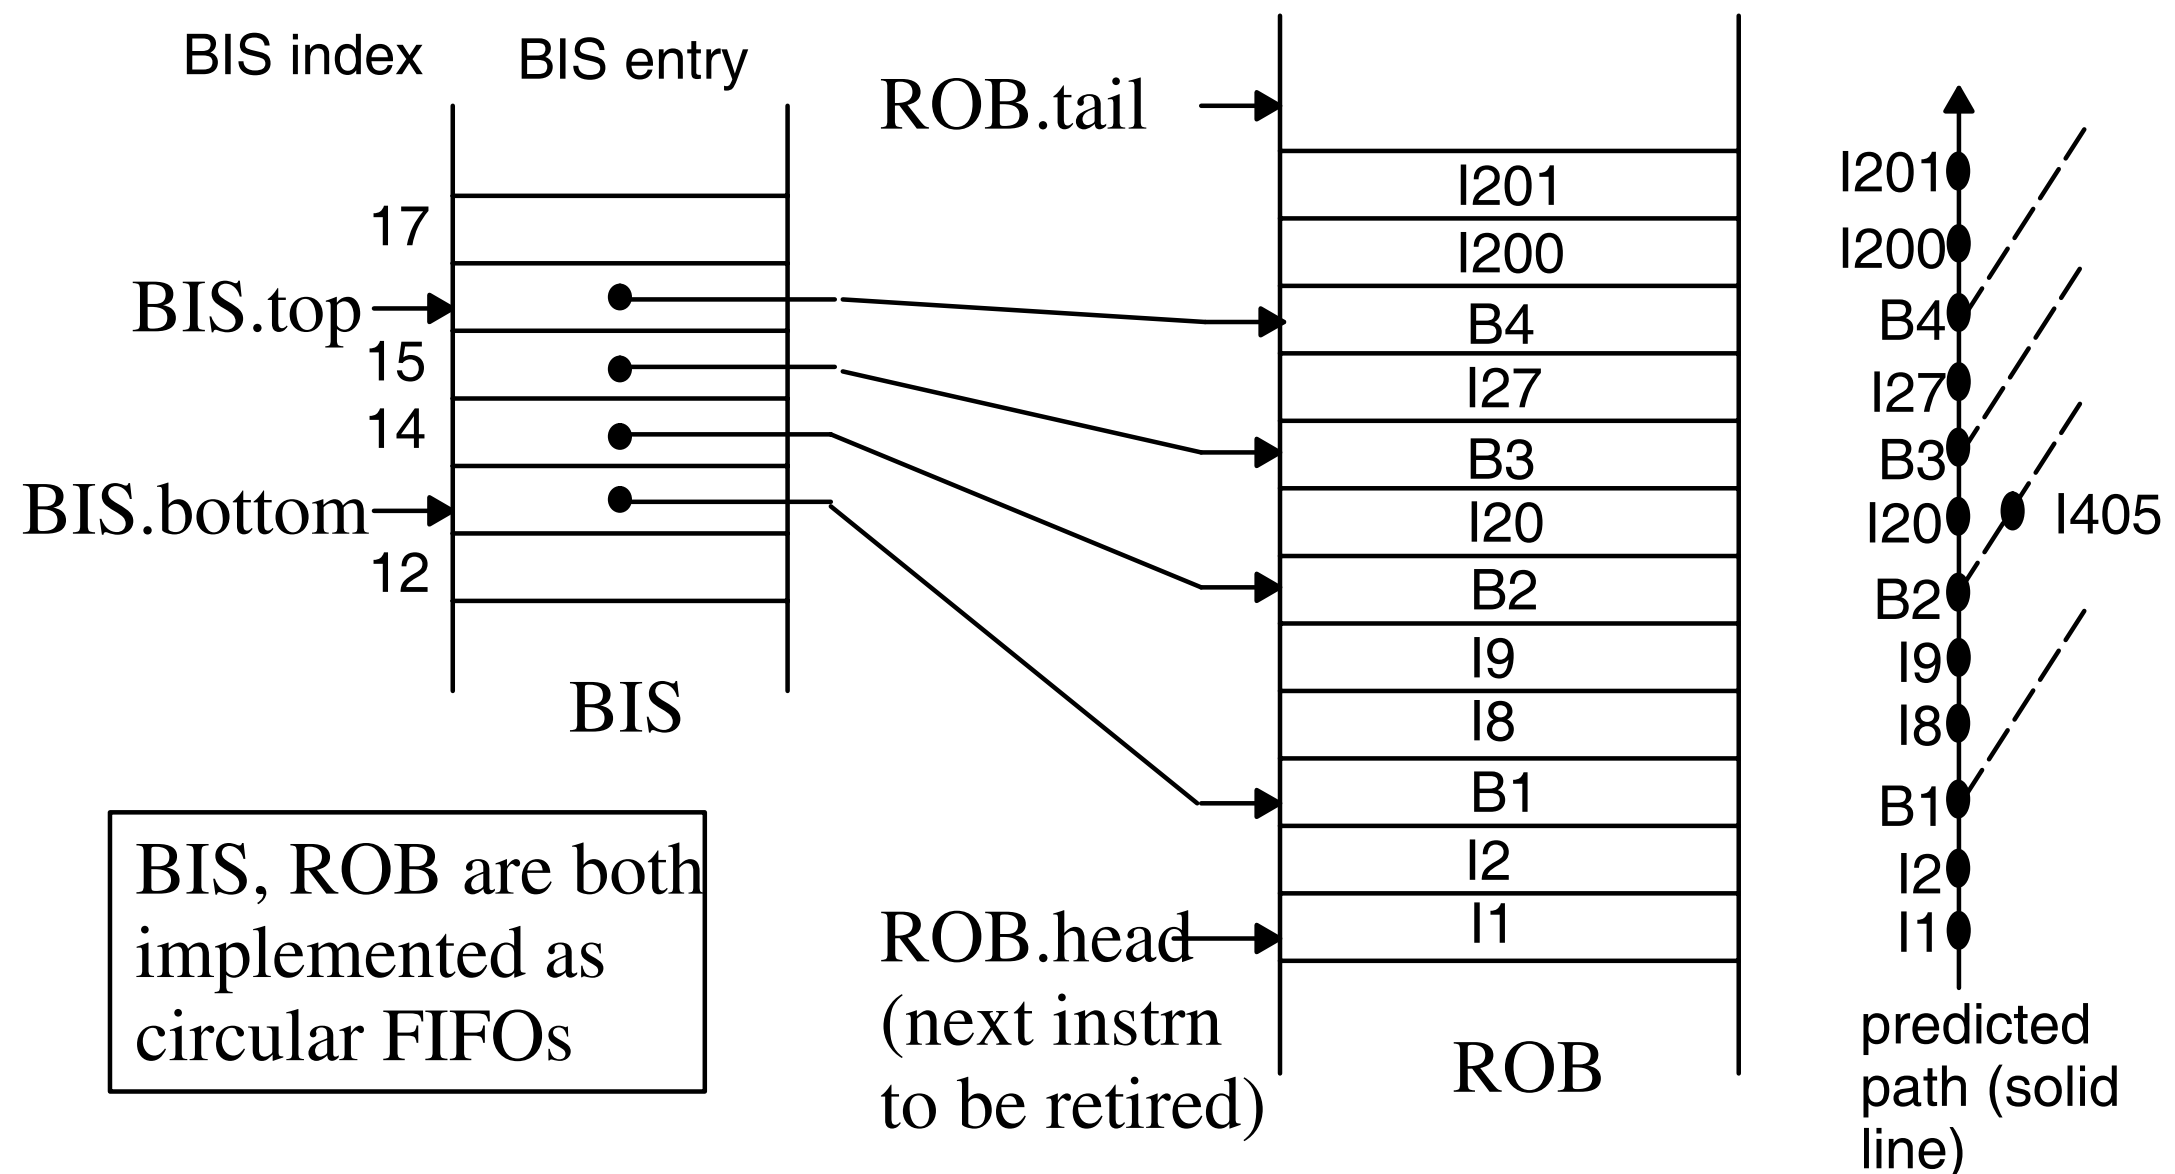
- Typically, a CPU can speculate through multiple predicted but unresolved branches.
- The number of outstanding branches is finite.

- Speculative execution along predicted branch direction:



Handling Speculative Execution with a ROB

- Additional facilities needed:
- ▶ A branch instruction stack, BIS, implemented as a circular FIFO queue, with pointers for the stack top and stack bottom. Both pointers move, as seen below:



Branch Tags

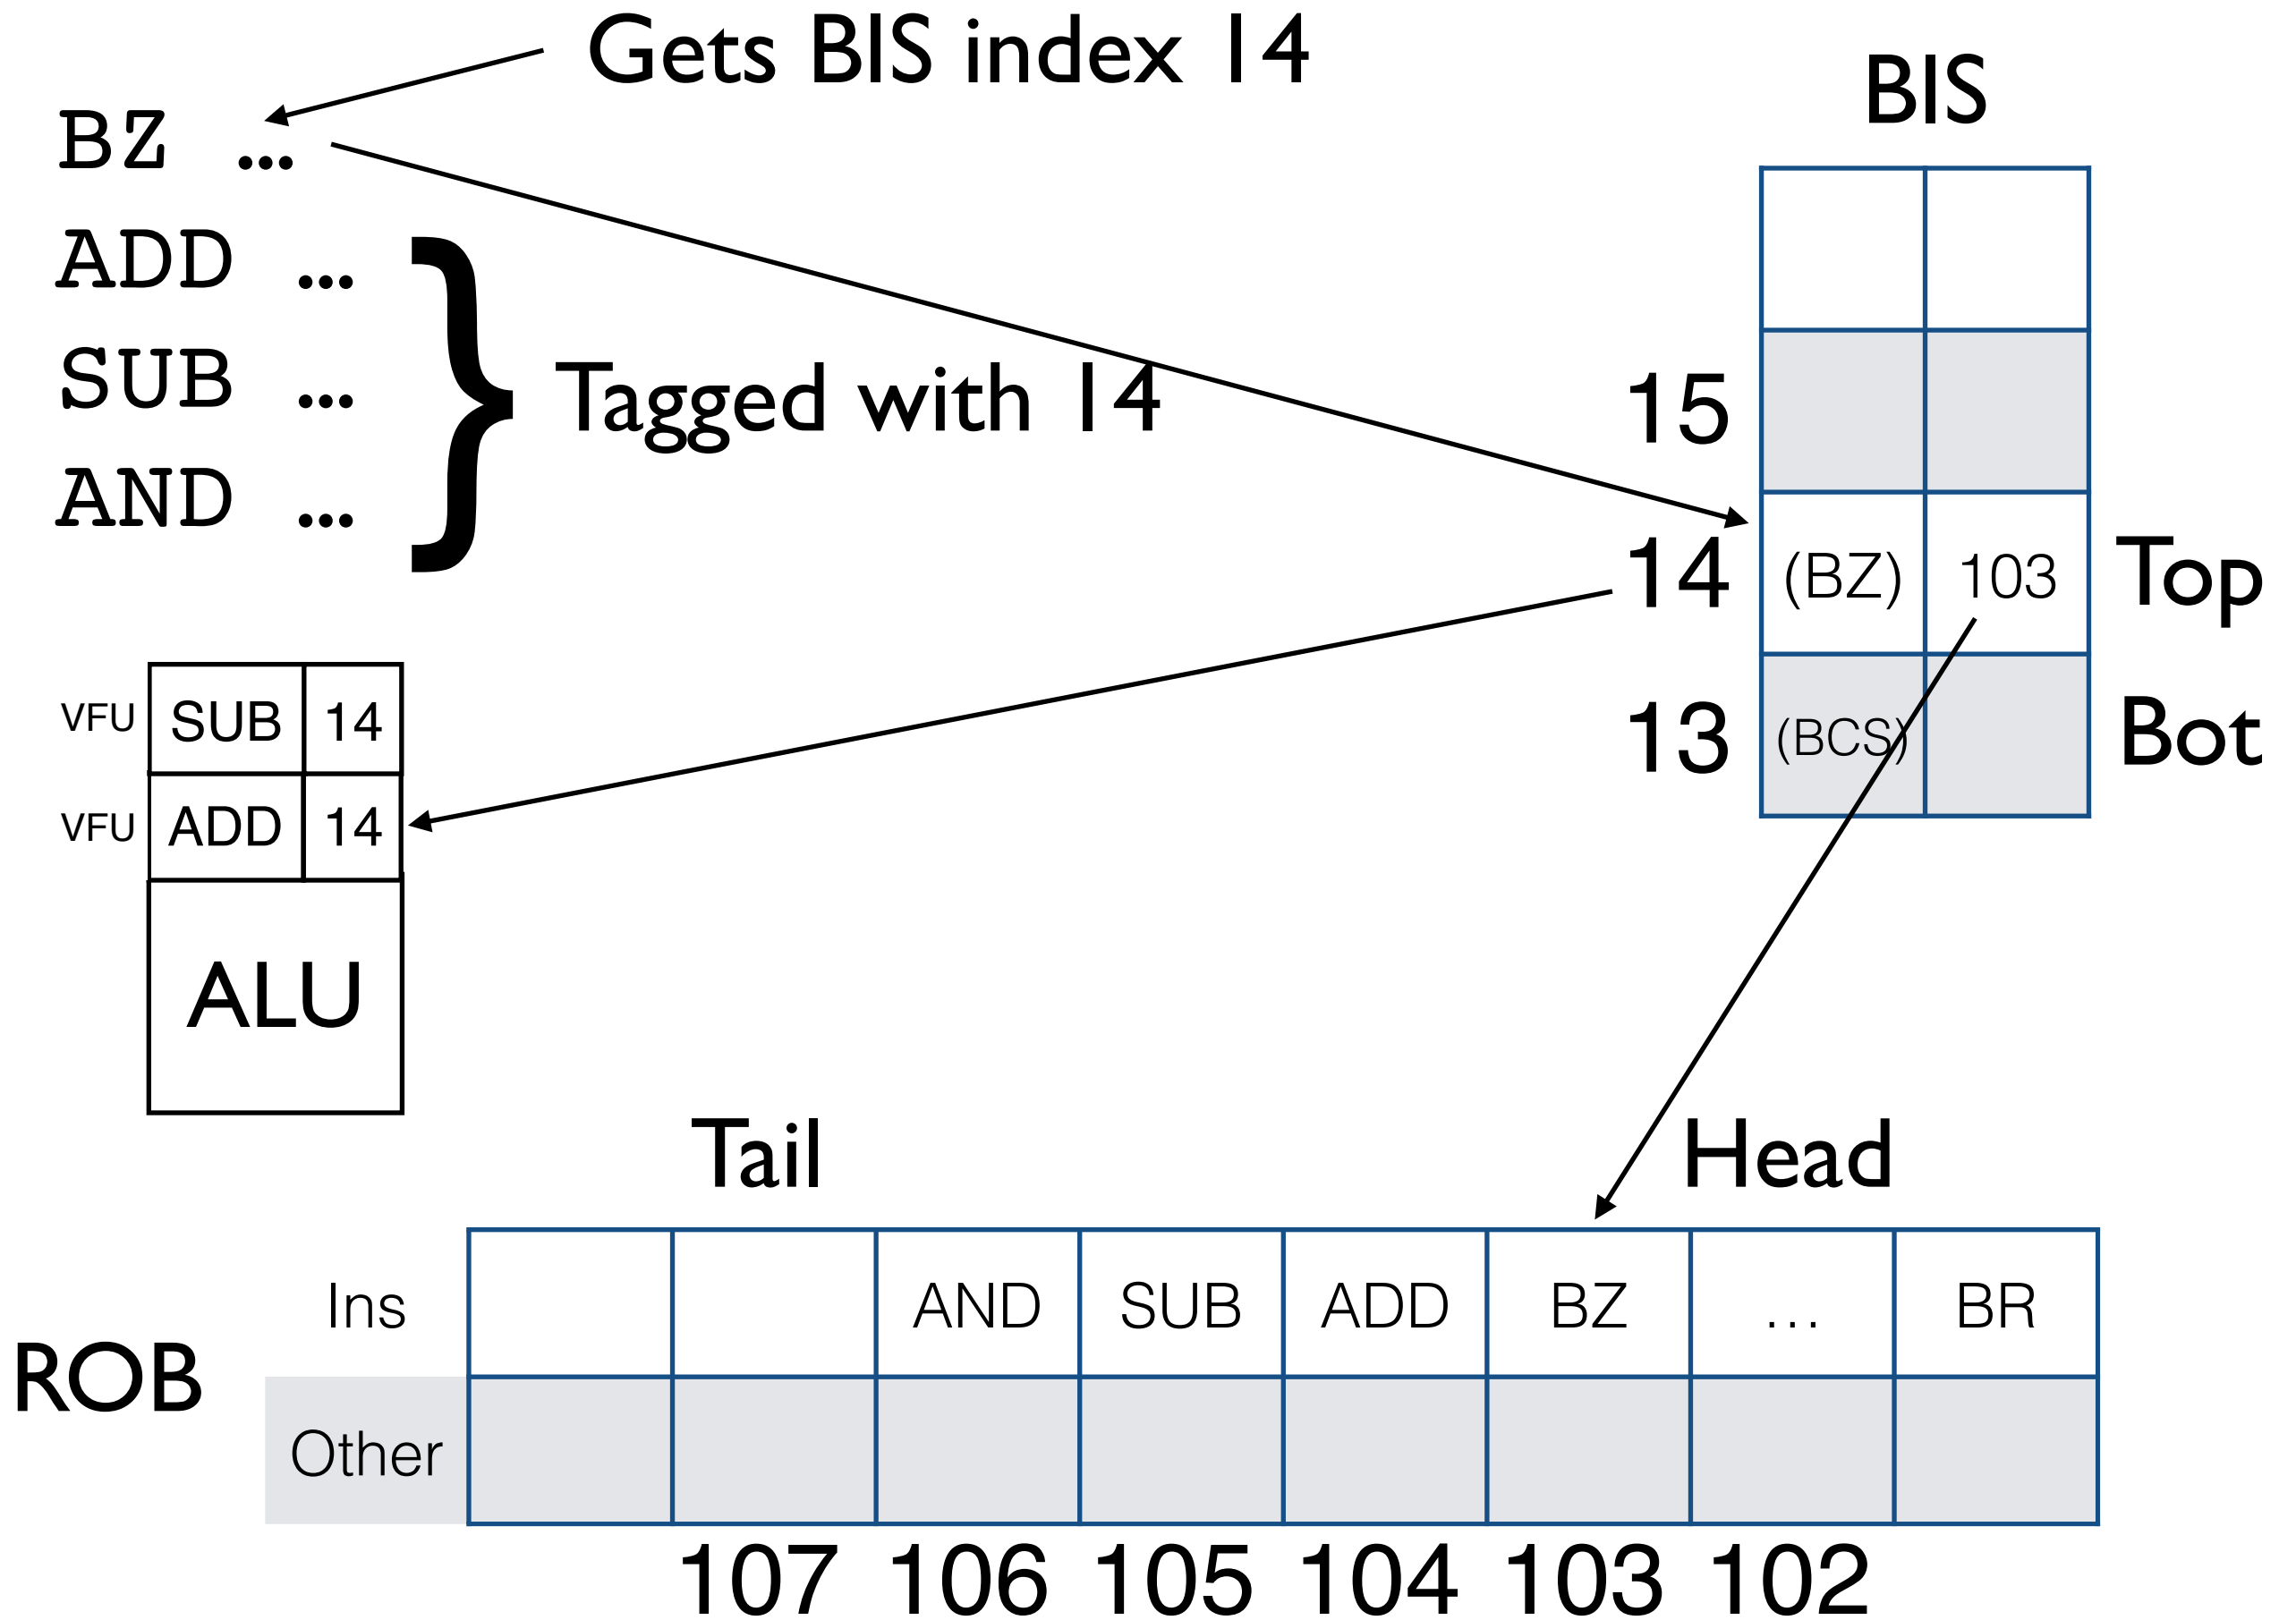
- Dispatched instructions are tagged with the BLS index of the closest preceding branch.
- Tag follows execution through to the RoB.
- Flush speculated instructions from RoB based on tag.

Dispatching Branches

- Allocate entries in: PRF, IQ, RoB, and BLS
- To allocate in BLS if not full:
 - Increment BLS.top to point to next free entry
 - BLS[top].rob = RoB entry for branch ins.

Dispatching Non-branch Instructions

- Allocate entries in: PRF, IQ, RoB
- Tag instructions in IQ with BLS.top



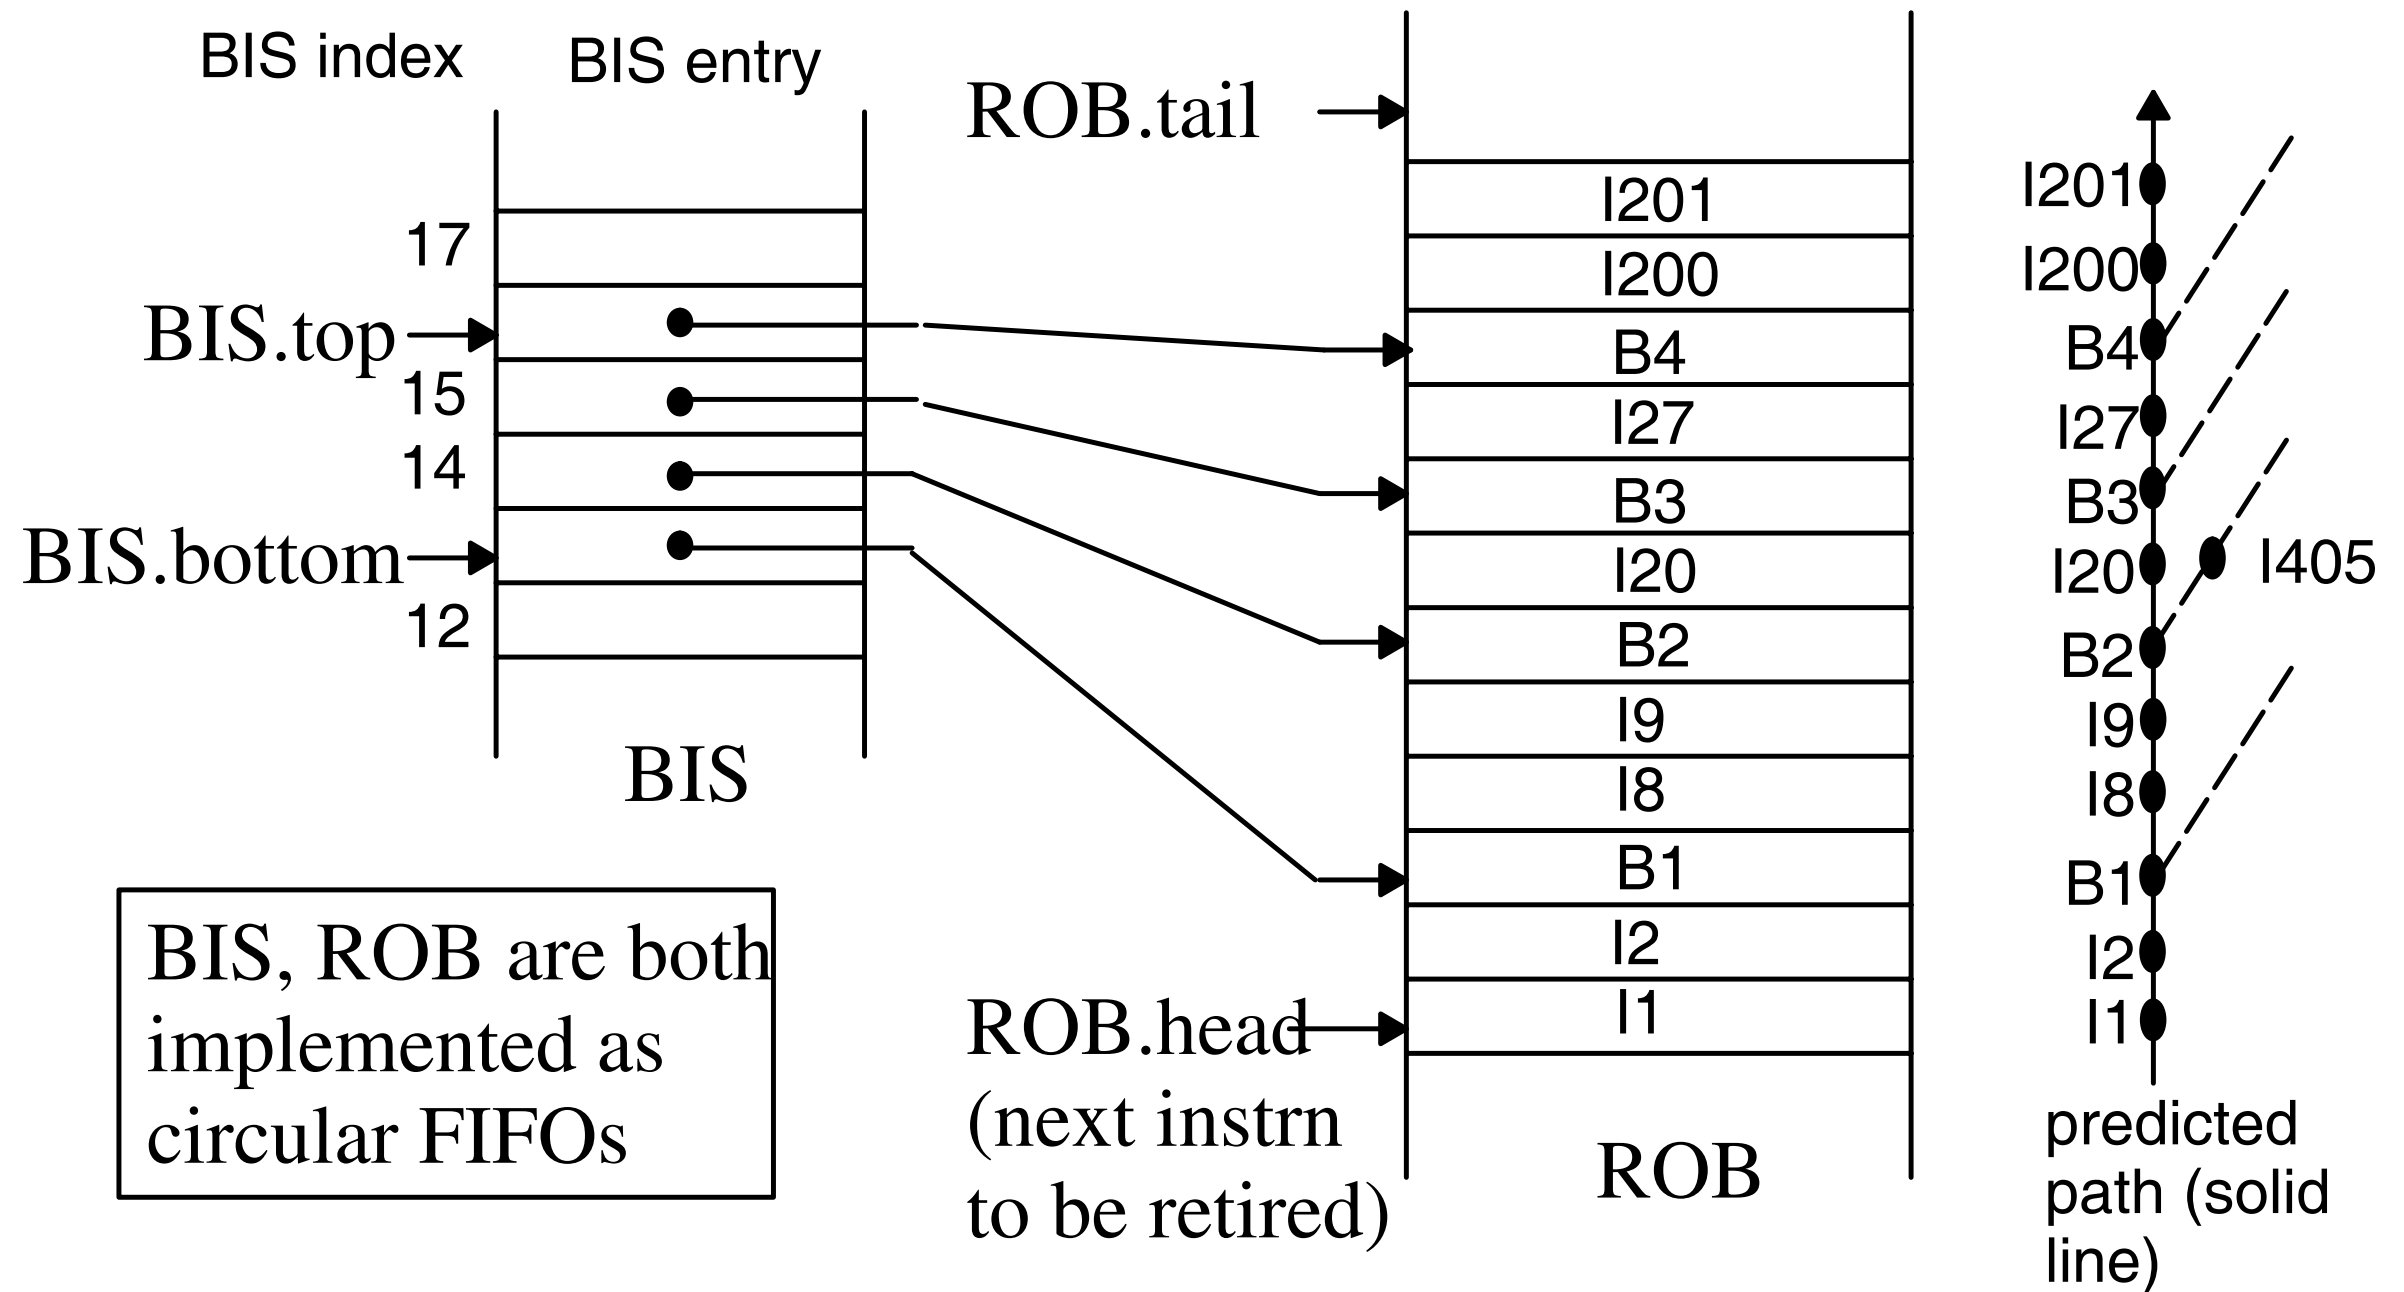
Handle Correctly- predicted Branch

- No special action on resolving branch.
- When committing RoB entry, increment BIS.bottom to free BIS entry.

Handle Mispredicted Branches

- Use BLS entry of mispredicted branch to locate branch in RoB.
- Flush from RoB all instructions following this branch.
 - Point RoB.tail to entry following branch.
- Flush from IQ and EX by broadcasting tags.
 - Broadcast branch tags from BLS
 - Decrement BLS.top until BLS.top points to mispredicted branch.

Example



Example Facts

- Flushes everything following mispredicted branch in program order
- Much in common with exception handling with precise interrupts
- Needn't wait for retirement to undo mispredicted branch.

Register Renaming & Speculation

- Speculated instructions allocated physical registers: Deallocate them
- Restore rename table to time of branch

Deallocating physical registers

- Integrated RoB and PRF:
 - Automatic with moving ROB.tail
- Separate RoB and PRF:
 - A: Walk RoB in reverse, freeing registers
 - B: Take checkpoint for each branch

Objective

- Situation:
 - Discover that branch in the **middle** of the instruction window was mispredicted.
- Objective:
 - Immediately reset front end of processor to time of mis-speculated instruction.
 - Only erase from RoB and IQ instructions **after** mispeculated one.

Restoring Rename Table

- Integrated RoB and PRF with future RAT and retirement RAT:
 - On discovering mispredict, stop dispatch
 - **Wait** for instructions prior to branch to commit, establishing precise architectural state at time of branch
 - Copy retirement RAT to future RAT
 - ▶ Treats mispredict like exception
 - ▶ Can be very slow


Restoring Rename Table Using Checkpoints

- Checkpoint RAT at time of branch
- Point BLS entry at checkpoint
- On mispredict, restore RAT from checkpoint
- ▶ Requires a lot of storage

Restoring Rename Table

“Walking Back”


- On mispredict:
 - Copy retirement RAT to future RAT
 - Fast forward from RoB.head to branch, updating RAT with arch to phys mappings found in RoB
 - “Walking backwards” towards the tail




Restoring Rename Table

“Walking Forward”

- Keep “undo” queue for RAT modifications
 - On dispatch, store old value of RAT[dest] in RAT history buffer.
- On mispredict:
 - Rewind ROB.tail back to branch, undoing all RAT changes made after branch
 - “Walking forward” towards ROB.head



Dispatching
ADD RI, ...



Old New

Put this into
undo buffer

Example

ADD R1, ...

MUL R0, ...

AND R1, ...

BZ ...

ADDC R2, ...

SUB R1, ...



ADD P4, ...

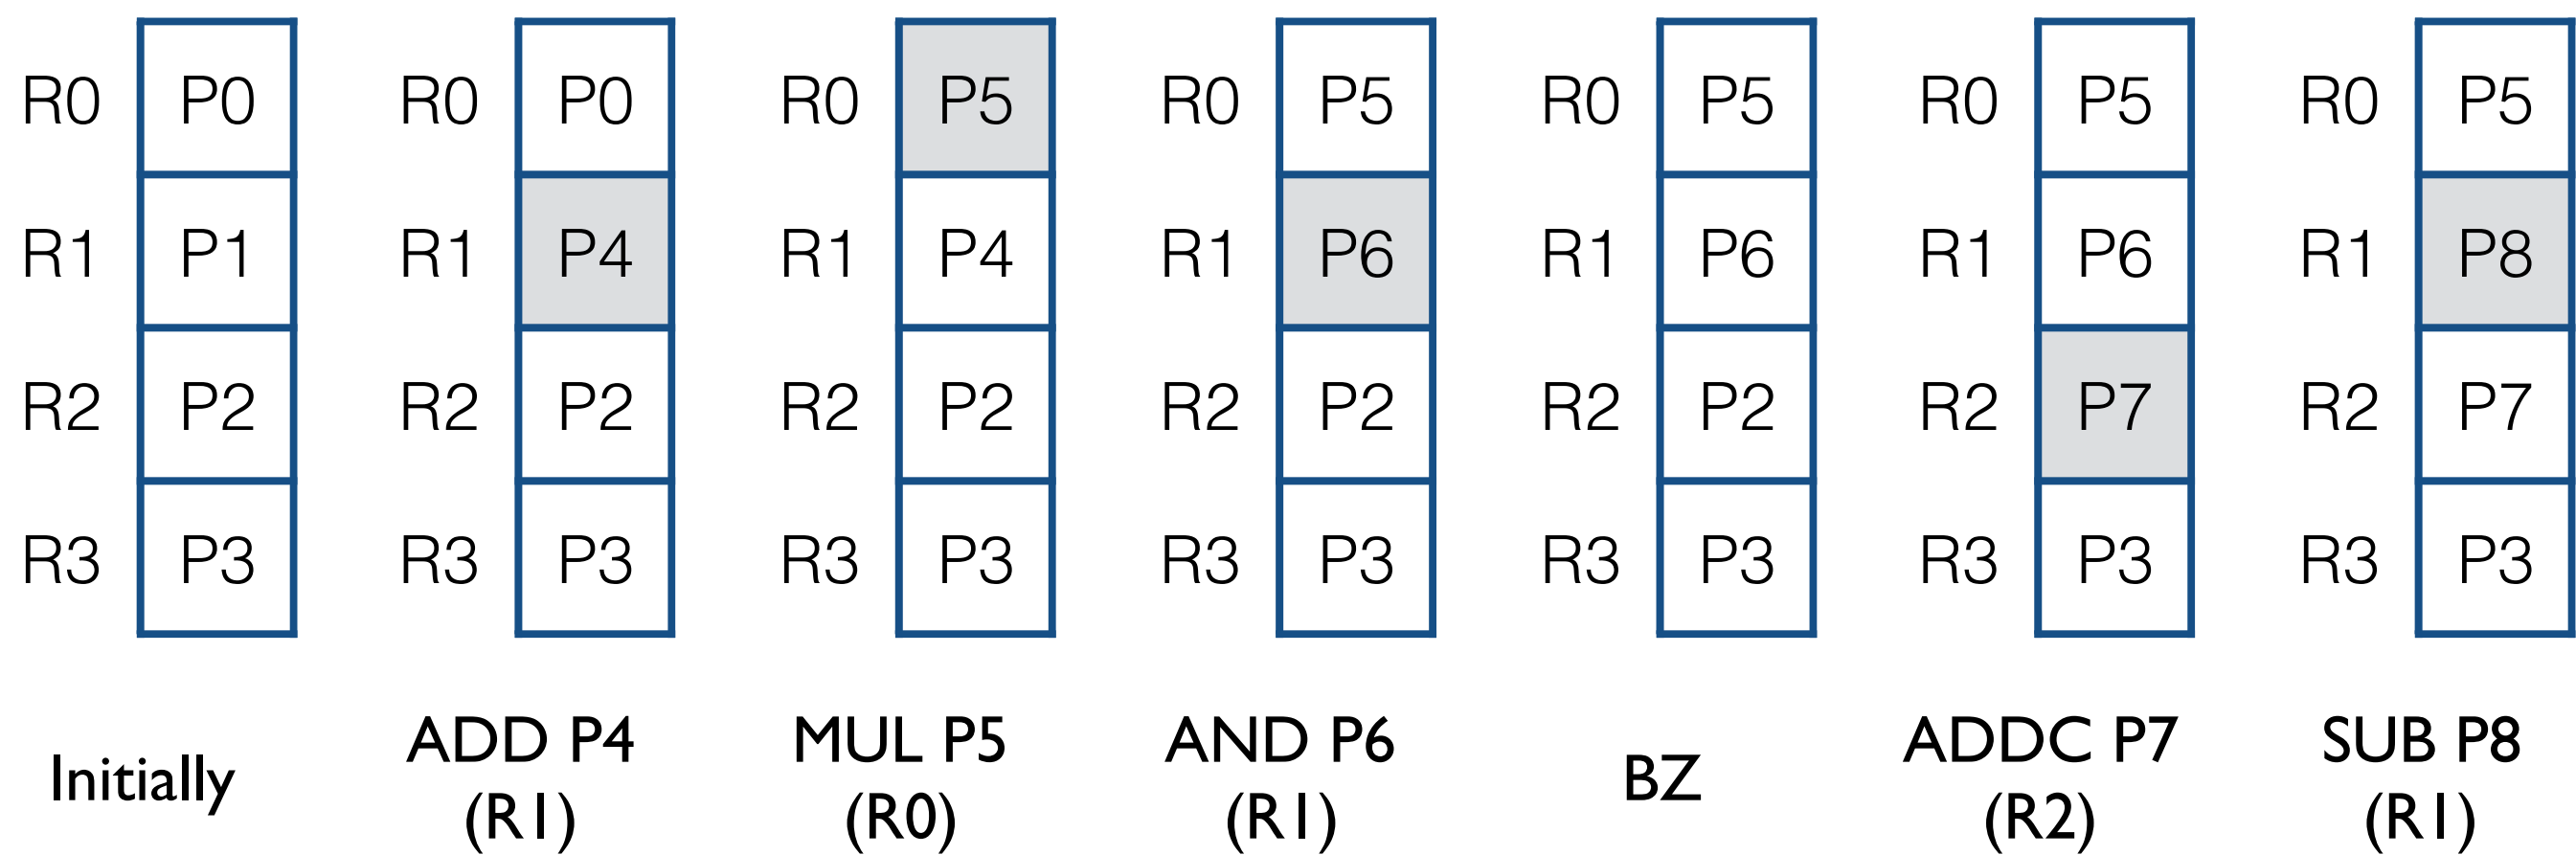
MUL P5, ...

AND P6, ...

BZ ...

ADDC P7, ...

SUB P8, ...

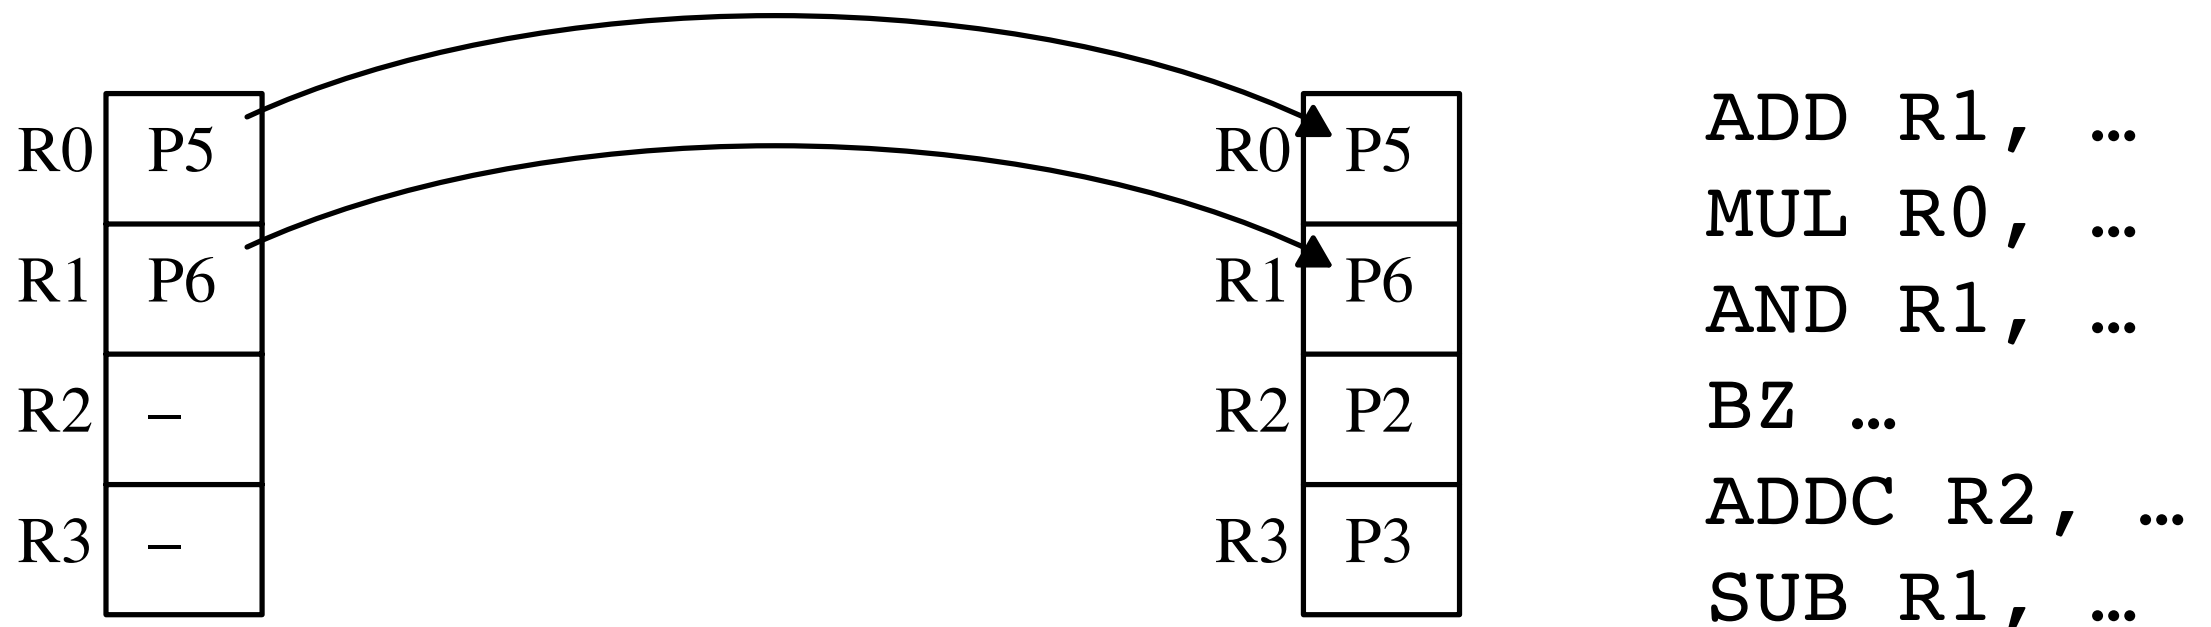


	Tail				Head			
Ins		SUB	ADDC	BZ	AND	MUL	ADD	
Arch		R1	R2		R1	R0	R1	
Phys		P8	P7		P6	P5	P4	
Other								

- **Alternative 1:** Treat mispredicted branch as an exception. Restore rename table by copying valid entries from the retirement RAT to the front end RAT (FE-RAT) after all instructions ahead of BZ are committed:

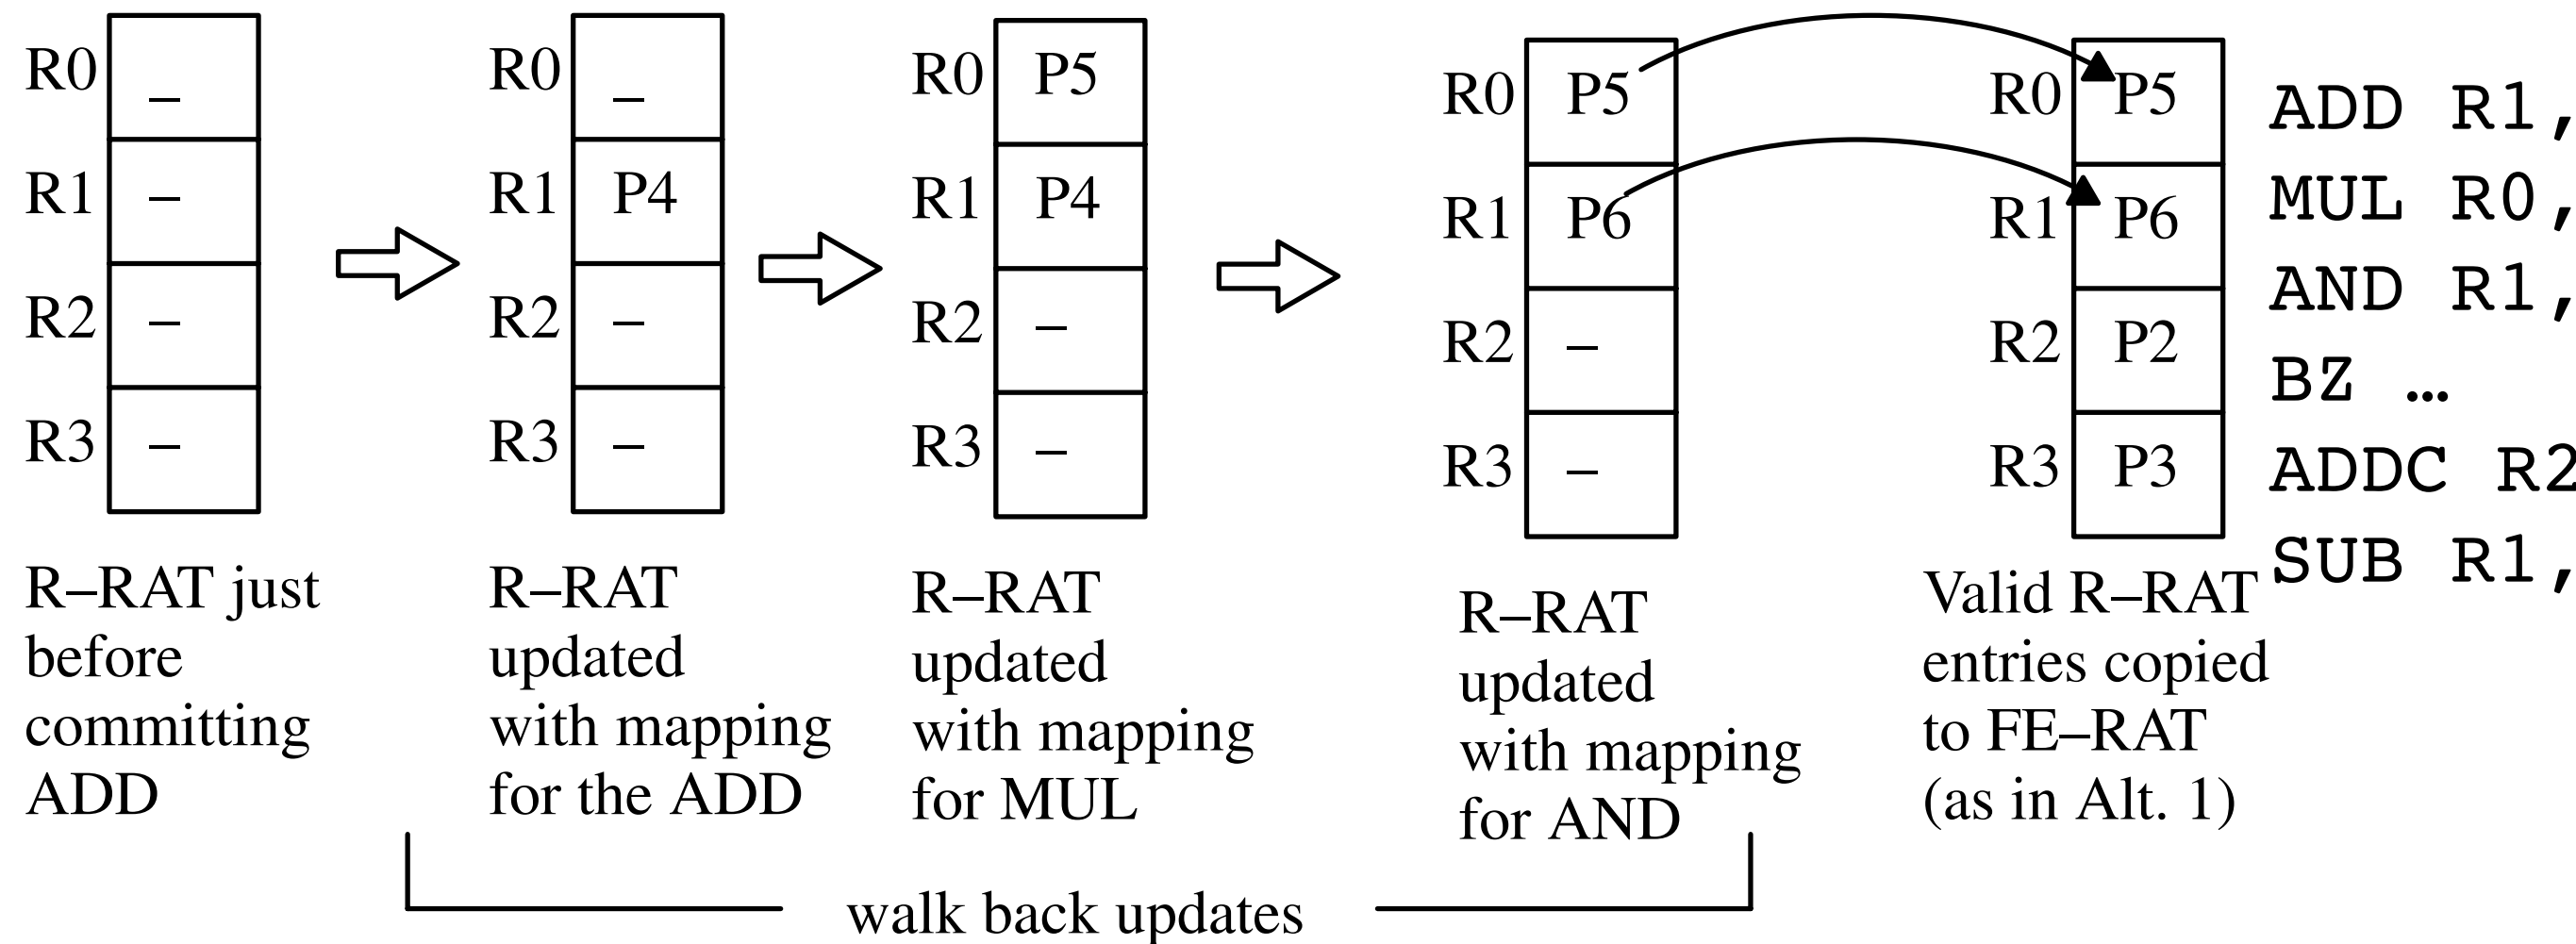
Retirement RAT (R-RAT) after committing instructions preceding BZ:

Restored front-end RAT: valid entries from R-RAT copied to FE-RAT



R-RAT entries for R2 and R3 are not copied to the FE-RAT, as they are not initialized

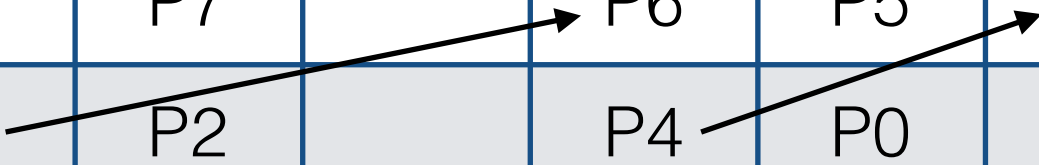
- Alternative 2:** Checkpoint the rename table – works with all datapath variations – here rename table at the point of dispatching BZ is saved.
- Alternative 3:** Improvement to Alternative 1 – use retirement RAT and **walk back** from the commit end of the ROB to the entry for the mispredicted branch:




ADD R1, ...
MUL R0, ...
AND R1, ...
BZ ...
ADDC R2, ...
SUB R1, ...

Alternative 4:

	Tail				Head			
Ins		SUB	ADDC	BZ	AND	MUL	ADD	
Arch		R1	R2		R1	R0	R1	
Phys		P8	P7		P6	P5	P4	
Old		P6	P2		P4	P0	P1	
Other								



Rename table (RAT) updates on walking forward:



RAT just
after
dispatching
SUB

RAT updated
with old
mapping
overwritten
by SUB

RAT updated
with old
mapping
overwritten
by ADDC

No up-
dates
for BZ:
this is the
restored
RAT

Free P8


Free P7

walk forward updates

- Checkpointing is the fastest scheme – hardware implementation used should avoid entry-by-entry copying.

Example: fast, parallel restoration similar to that used in the AMD K6:

- One entry in the rename table: permits four checkpoints, “a” through “d”; each entry uses 7 bits (bits 6 through 0) to name a destination



- For each bit, there are 4 values – bits “a” through “d”, corresponding to the 4 checkpoints. The bits holding these 4 values are set up logically as a circular shift register, as indicated by the arrows

Restoring Rename Table Analysis

- Walking techniques sequential and slow
 - Dependent on position of branch in RoB
- RAT restoration technique influences how physical registers are freed.
- Hardware requirements:
 - Most: Checkpointing
 - Walking Forward
 - Least: Walking Back or waiting